

Structured Query Language

SQL is everywhere

- just about every computer and every person on planet eventually touches something running SQL
- incredibly successful and solid technology
- runs universities (System Informatics iNtegrated University), banks, hospitals, governments, small businesses, large ones
- all Android Phones and iPhones have easy access to SQL database called SQLite
 - many applications on your phone use it directly

learning SQL ?

- weird obtuse kind of "non-language"
 - that most programmers can't stand
- based on solid mathematically built theory of operation (relational theory (extension of set theory))
- actually learn important theoretical concepts that apply to nearly every data storage system past and present

Non Procedural, Declarative Programming Language

- (procedural) programming languages - variables and data structures, using conditional logic (i.e., if-then-else) and looping constructs (i.e., do while ... end),
- procedural language defines both desired results and mechanism, or process, by which results are generated
- Nonprocedural languages define desired results, but process by which results are generated is left to an external agent
- manner in which statement is executed is left to component of your database engine known as *query optimizer*
- integrate SQL with programming language

SQL

- pronounce SQL "sequel"
- but you can also say "ESS-QUEUE-ELL" if you want
- stands for Structured Query Language
- language for interacting with data in (relational) database
- matches theory established many years ago defining properties of well structured data

SQL operations

- Create
 - putting data into tables - INSERT
- Read
 - query data out of tables - SELECT
- Update
 - change data already in table - UPDATE
- Delete
 - remove data from table – DELETE
- acronym "CRUD" is considered fundamental set of features every data storage system must have

SQL - language for doing CRUD operations

- SQL Data Definition Language - DDL
 - to produce new tables or alter existing ones
 - SQL only knows tables, and every operation produces tables
 - by modifying an existing one
 - or it returns new temporary table as your data set
- SQL Data Manipulation Language – DML
- SQL Security

SELECT

sqlzoo.net

- 3 laboratory works
- A Gentle Introduction to
- Structured Query Language

Why is it called SQLzoo?

- The animals of this zoo are SQL engines - one of each species
- They have been caged and tamed
- The public can poke, prod and gawp - the exhibits and the public are protected from each other

Who runs this site?

- Andrew Cumming is a lecturer at Napier University in Edinburgh, Scotland, UK
- is the zoo keeper, he feeds the animals and shovels away the waste

Can I shake his hand?
buy him a drink?

- Next time you are visiting Edinburgh, UK you can shake his hand at the Tollcross State Circus which meets on Monday, 7pm to 9pm in Tollcross primary school

Can I buy him a drink?

- You can buy him a drink afterwards, in the Blue Blazer, Spittal Street; a pint of 80 Shilling please
- I (not Technical University of Cluj Napoca) will reimburse the expenses
 - bring a proof

Notes for teachers

- *This material was designed to be used in supervised tutorials at Napier University*
- Teachers from other institutions are welcome to use it in any way they see fit, however I have a few requests / suggestions:
- Reliability of the SQL Engines
 - Sometimes long running processes accumulate
- Feedback
 - If you are making use of this material please let me know what worked well and what didn't
 - Let me know if any of the questions are poorly phrased or confusing. I have tried to keep the language as simple as possible. Many students do not have English as their first language - I would like to hear more from them

Acknowledgements

- The CIA
 - Not just spying and wet work, they help you with your geography homework
- BBC News - Country Profiles
- Internet Movie Database
 - wonderful way to waste hours
- Scotland's Parliament
 - They've got my vote.
- TRAVELINE, Edinburgh City Council
- Amazon, New Scientist

SQL Basic Concepts and Principles

BASIC CONCEPTS AND PRINCIPLES

- in 1970, Dr. E. F. Codd of IBM's research laboratory published paper titled "A Relational Model of Data for Large Shared Data Banks" that proposed that data be represented as relations, sets of *tables*
- redundant data used to link records in different tables

- each table in relational database includes information that uniquely identifies row in that table (known as ***primary key***), along with additional information needed to describe entity completely
- some tables also include information used to navigate to another table; this is where “redundant data” mentioned earlier comes in
- these columns are known as ***foreign keys***, and they serve purpose as lines that connect entities

Terminology

- Entity - something of interest to database user community; examples include customers, parts, geographic locations, etc.
- Column - individual piece of data stored in table
- Row - set of columns that together completely describe entity; also called record
- Table - set of rows, held on permanent storage (persistent)
- Result set - another name for nonpersistent table, generally result of SQL query
- Primary Key - one or more columns that can be used as unique identifier for each row in table
- Foreign Key - one or more columns that can be used together to identify single row in another table

Terminology

- Result set - another name for nonpersistent table, generally result of SQL query
- SQL goes with relational model - result of SQL query is table (also called, in this context, *result set*)
- table can be created in relational database simply by storing result set of query
- query can use both permanent tables and result sets from other queries as inputs

Example

EXAMPLE

Northwind sample database

- Northwind Traders, Access database, sample database
- contains sales data for fictitious company called Northwind Traders, which imports and exports specialty foods from around the world
- like real world application but names of companies, products, employees, and any other data used or mentioned do not in any way represent any real world individual, company, product, etc.
- can use Northwind to explore nearly every important aspect of database

- Northwind Sample Databases
- <https://archive.codeplex.com/?p=northwinddatabase>
- Northwind and pubs Sample Databases for SQL Server
- <https://www.microsoft.com/en-us/download/details.aspx?id=23654>
- <https://github.com/fsprojects/SQLProvider/blob/master/docs/files/msaccess/Northwind.accdb>

categories

- CategoryID, integer, Primary Key
- CategoryName, varchar(15)
- Description, text
- Picture, varbinary(MAX)

- CategoryID CategoryName Description
- 1 Beverages Soft drinks, coffees, teas, beers, and ales
- 2 Condiments Sweet and savory sauces, relishes, spreads, and seasonings
- 3 Confections Desserts, candies, and sweet breads
- 4 Dairy Products Milk and Cheeses
- 5 Grains/Cereals Breads, crackers, pasta, and cereal
- 6 Meat/Poultry Prepared meats
- 7 Produce Dried fruit and bean curd
- 8 Seafood Seaweed and fish

customers

- CustomerID, varchar(5), Primary Key
- CompanyName, varchar(40)
- ContactName, varchar(30)
- ContactTitle, varchar(30)
- Address, varchar(60)
- City, varchar(15)
- Region, varchar(15)
- PostalCode, varchar(10)
- Country, varchar(15)
- Phone, varchar(24)
- Fax, varchar(24)

•	CustomerID	Region	CompanyName	PostalCode	Country	ContactName	Phone	Fax	ContactTitle	Address	City
•	ALFKI	Alfreds Futterkiste	12209	Germany	Maria Anders	Sales Representative	030-0074321	030-0076545		Obere Str. 57 Berlin	NULL
•	ANATR	Ana Trujillo Emparedados y helados	México D.F.	NULL	05021	Ana Trujillo	Mexico	(5) 555-4729	(5) 555-3745	Avda. de la Constitución 2222	
•	ANTON	Antonio Moreno Taquería	D.F.	NULL	05023	Antonio Moreno	Mexico	(5) 555-3932	NULL	Mataderos 2312	México
•	AROUT	Around the Horn	London	NULL		Thomas Hardy	WA1 1DP	UK	Sales Representative	120 Hanover Sq.	
•	BERGS	Berglunds snabbköp	Luleå	NULL		Christina Berglund	S-958 22	Sweden	Order Administrator	Berguvsvägen 8	
•	BLAUS	Blauer See Delikatessen	68306	Germany		Hanna Moos	0621-08460	0621-08924	Sales Representative	Forsterstr. 57 Mannheim	NULL
•	BLONP	Blondesddsl père et fils	Strasbourg	NULL		Frédérique Citeaux	67000	France	Marketing Manager	24, place Kléber	
•	BOLID	Bólido Comidas preparadas	Madrid	NULL	28023	Martín Sommer		Spain	Owner	C/ Araquil, 67	
•	BONAP	Bon app'	13008	France		Laurence Lebihan	91.24.45.40	91.24.45.41	Owner	12, rue des Bouchers	Marseille NULL
•	BOTTM	Bottom-Dollar Markets	Tsawassen	BC		Elizabeth Lincoln	T2F 8M4	Canada	Accounting Manager	23 Tsawassen Blvd.	

employees

- EmployeeID, integer, Primary Key
- LastName, varchar(20); FirstName, varchar(10)
- Title, varchar(30); TitleOfCourtesy, varchar(25)
- BirthDate, date; HireDate, date
- Address, varchar(60)
- City, varchar(15)
- Region, varchar(15)
- PostalCode, varchar(10)
- Country, varchar(15)
- HomePhone, varchar(24); Extension, varchar(4)
- Photo, varchar(50)
- Notes, text
- ReportsTo, integer Foreign Key refer PK

EmployeeID	LastName Region	FirstName PostalCode	Title Country	TitleOfCourtesy HomePhone	BirthDate Extension	HireDate Notes	Address ReportsTo	City PhotoPath
1	Davolio 507 - 20th Ave. E.	Nancy	Sales Representative	Ms.	1948-12-08 00:00:00.000		1992-05-01 00:00:00.000	
Apt. 2A	Seattle	WA	98122	USA	(206) 555-9857	5467	Education includes a BA in psychology	
	from Colorado State University in 1970. She also completed "The Art of the Cold Call." Nancy is a member of Toastmasters International. http://accweb/emmployees/davolio.bmp							2
2	Fuller 908 W. Capital Way	Andrew Tacoma	Vice President, Sales	Dr. WA	1952-02-19 00:00:00.000 98401	USA (206) 555-9482	1992-08-14 00:00:00.000 3457	Andrew received
	his BTS commercial in 1974 and a Ph.D. in international marketing from the University of Dallas in 1981. He is fluent in French and Italian and reads German. He joined the company as a sales representative, was promoted to sales manager in January 1992 and to vice president of sales in March 1993. Andrew is a member of the Sales Management Roundtable, the Seattle Chamber of Commerce, and the Pacific Rim Importers Association. http://accweb/emmployees/fuller.bmp							
3	Leverling 722 Moss Bay Blvd.	Janet Kirkland	Sales Representative	Ms. WA	1963-08-30 00:00:00.000 98033	USA (206) 555-3412	1992-04-01 00:00:00.000 3355	Janet has a BS
	degree in chemistry from Boston College (1984). She has also completed a certificate program in food retailing management. Janet was hired as a sales associate in 1991 and promoted to sales representative in February 1992. http://accweb/emmployees/leverling.bmp							2
4	Peacock 4110 Old Redmond Rd.	Margaret	Sales Representative	Mrs. Redmond	1937-09-19 00:00:00.000 WA	98052 USA	1993-05-03 00:00:00.000 (206) 555-8122	5176
	Margaret holds a BA in English literature from Concordia College (1958) and an MA from the American Institute of Culinary Arts (1966). She was assigned to the London office temporarily from July through November 1992. http://accweb/emmployees/peacock.bmp							2
5	Buchanan 14 Garrett Hill	Steven London	Sales Manager	Mr. NULL	1955-03-04 00:00:00.000 SW1 8JR	UK (71) 555-4848	1993-10-17 00:00:00.000 3453	Steven Buchanan
	graduated from St. Andrews University, Scotland, with a BSC degree in 1976. Upon joining the company as a sales representative in 1992, he spent 6 months in an orientation program at the Seattle office and then returned to his permanent post in London. He was promoted to sales manager in March 1993. Mr. Buchanan has completed the courses "Successful Telemarketing" and "International Sales Management." He is fluent in French. http://accweb/emmployees/buchanan.bmp							2
6	Suyama Coventry House	Michael	Sales Representative	Mr.	1963-07-02 00:00:00.000		1993-10-17 00:00:00.000	
Miner Rd.	London	NULL	EC2 7JR	UK	(71) 555-7773	428	Michael is a graduate of Sussex	
	University (MA, economics, 1983) and the University of California at Los Angeles (MBA, marketing, 1986). He has also taken the courses "Multi-Cultural Selling" and "Time Management for the Sales Professional." He is fluent in Japanese and can read and write French, Portuguese, and Spanish. http://accweb/emmployees/davolio.bmp							5
7	King Edgeham Hollow	Robert	Sales Representative	Mr.	1960-05-29 00:00:00.000		1994-01-02 00:00:00.000	
Winchester Way	London	NULL	RG1 9SP	UK	(71) 555-5598	465	Robert King	
	served in the Peace Corps and traveled extensively before completing his degree in English at the University of Michigan in 1992, the year he joined the company. After completing a course entitled "Selling in Europe," he was transferred to the London office in March 1993. http://accweb/emmployees/davolio.bmp							5
8	Callahan 00:00:00.000 4726 - 11th Ave. N.E.	Laura Seattle	Inside Sales Coordinator	Ms. WA	1958-01-09 00:00:00.000 98105	USA (206) 555-1189	1994-03-05 2344	Laura received a
	BA in psychology from the University of Washington. She has also completed a course in business French. She reads and writes French. http://accweb/emmployees/davolio.bmp							2
9	Dodsworth 7 Houndstooth Rd.	Anne London	Sales Representative	Ms. NULL	1966-01-27 00:00:00.000 WG2 7LT	UK (71) 555-4444	1994-11-15 00:00:00.000 452	Anne has a BA
	degree in English from St. Lawrence College. She is fluent in French and German. http://accweb/emmployees/davolio.bmp							5

- EmployeeID LastName FirstName ReportTo
 - 1 Davolio Nancy 2
 - 2 Fuller Andrew NULL
-
- values and reference between Foreign Key ReportTo and Primary Key EmployeeID means that Nancy Davolio report to Andrew Fuller and that last one report to none

products

- ProductID, integer, Primary Key
- ProductName, varchar(40)
- SupplierID, integer; CategoryID integer - Foreign Keys
- QuantityPerUnit, varchar(20)
- UnitPrice, double
- UnitsInStock, integer; UnitsOnOrder integer
- ReorderLevel, integer
- Discontinued, enum('y','n')

shippers

- ShipperID, integer, Primary Key
- CompanyName, varchar(40)
- Phone, varchar(24)

suppliers

- SupplierID, integer, Primary key
- CompanyName, varchar(40)
- ContactName, varchar(30)
- ContactTitle, varchar(30)
- Address, varchar(60)
- City, varchar(15)
- Region, varchar(15)
- PostalCode, varchar(10)
- Country, varchar(15)
- Phone, varchar(24); Fax, varchar(24)
- HomePage, varchar(255)

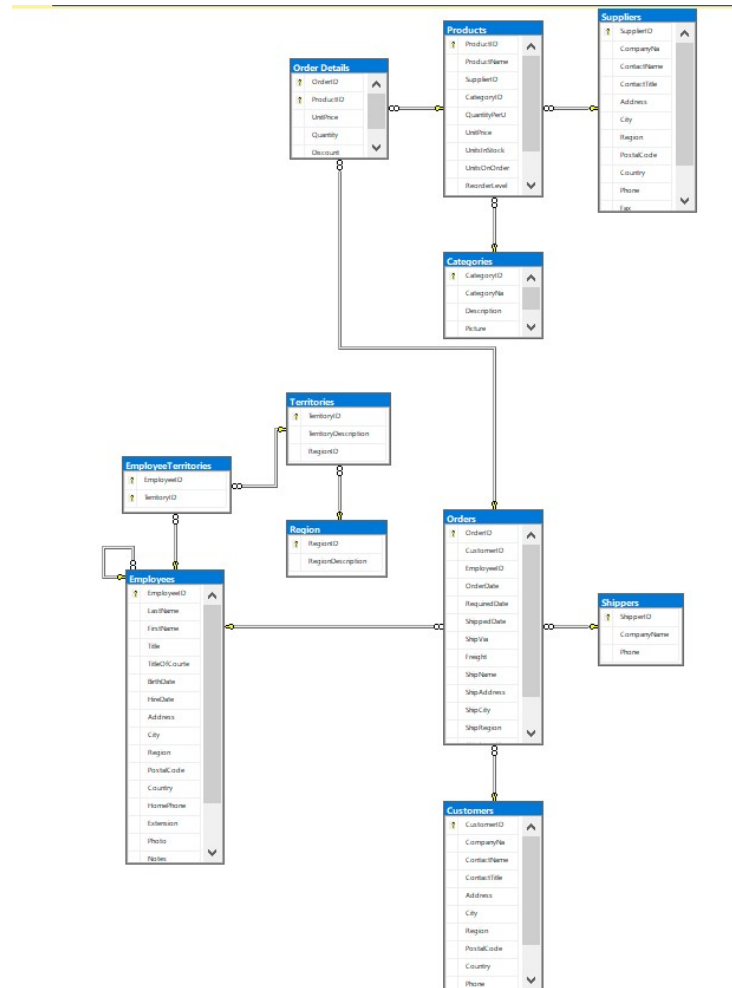
orders

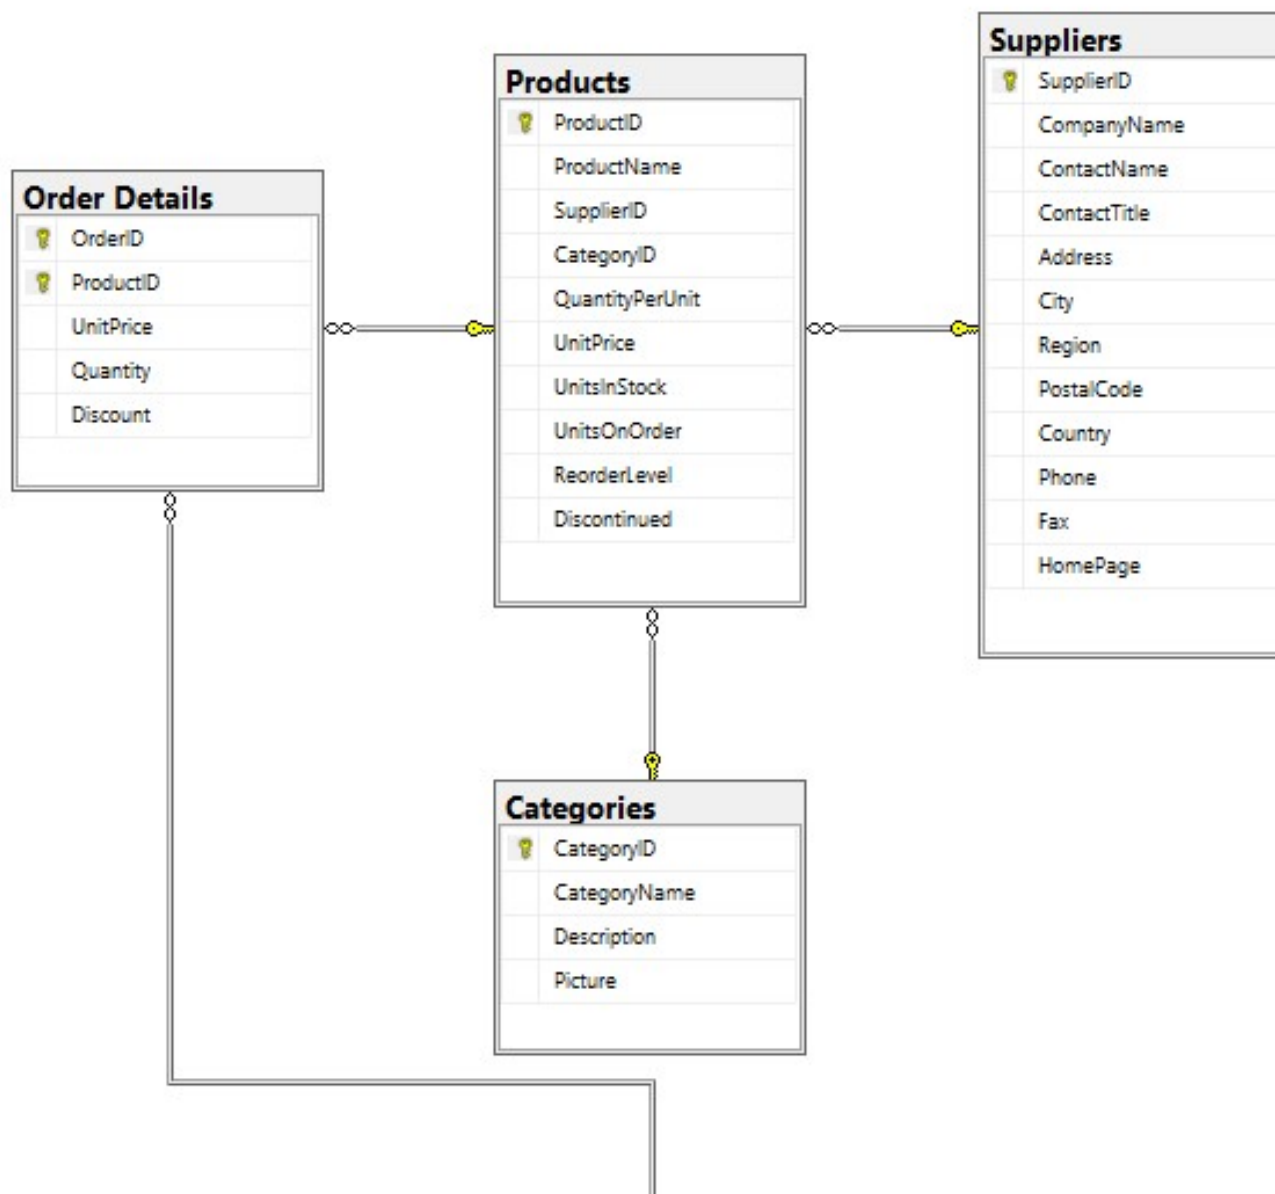
- OrderID, integer, Primary Key
- CustomerID, varchar(5); EmployeeID, integer - Foreign Keys
- OrderDate, date
- RequiredDate, date; ShippedDate, date
- ShipVia, integer Foreign Key
- Freight, double
- ShipName, varchar(40)
- ShipAddress, varchar(60)
- ShipCity, varchar(15)
- ShipRegion, varchar(15)
- ShipPostalCode, varchar(10)
- ShipCountry, varchar(15)

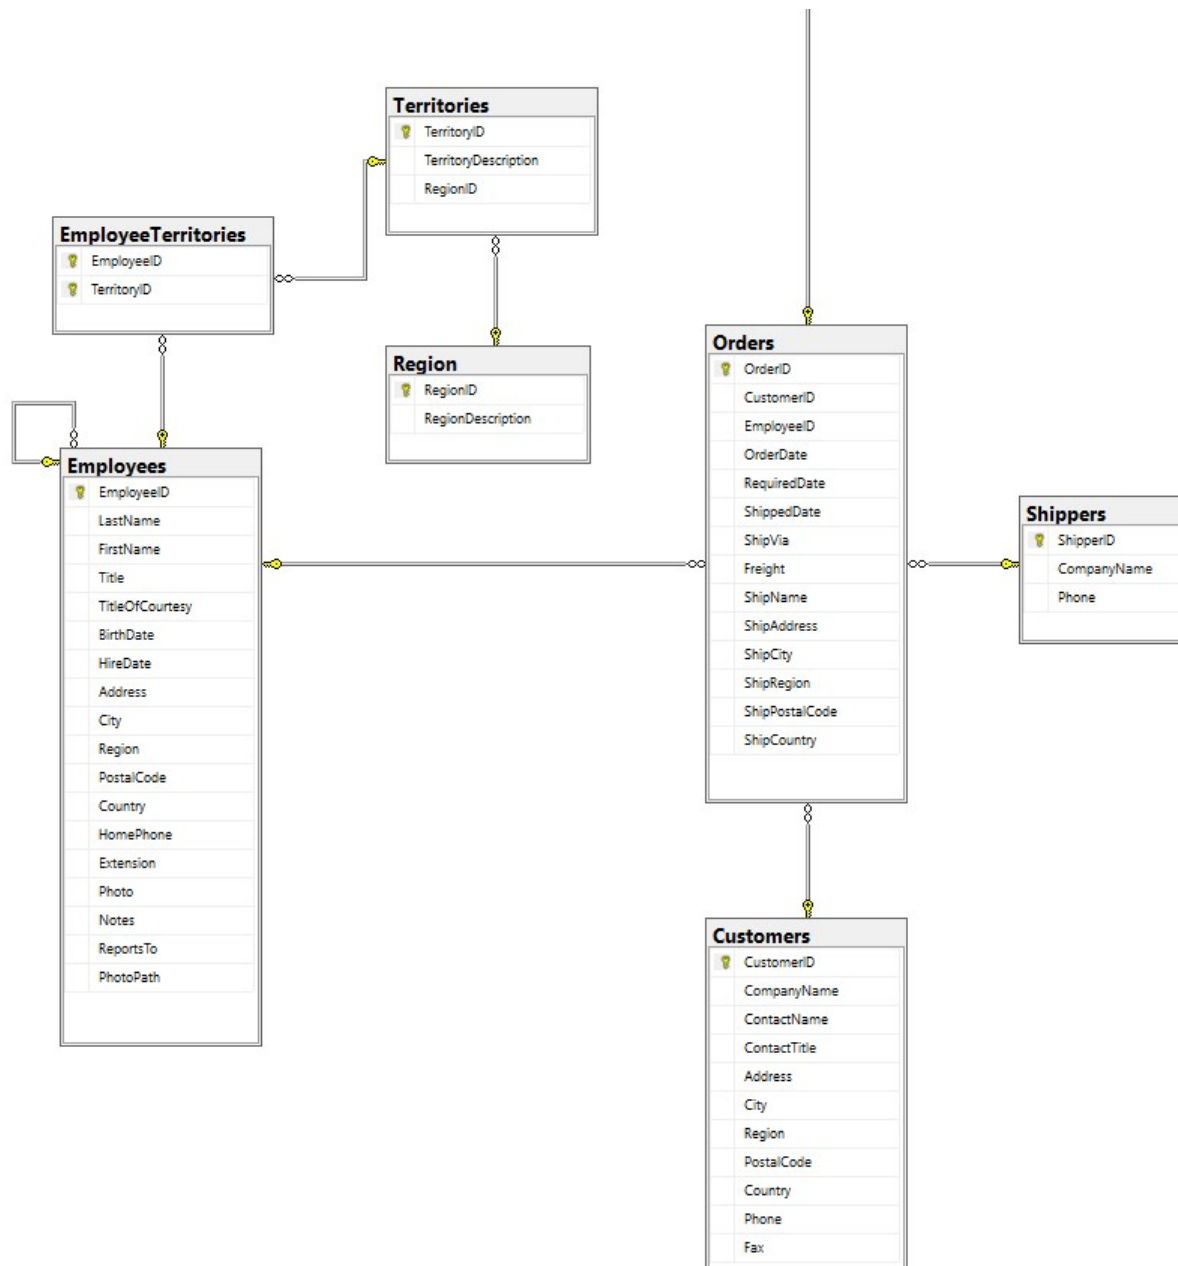
order_details

- ID, integer, OrderID, integer (Foreign Key refer order table) – Primary Key
- ProductID, integer – Foreign Key
- UnitPrice, double
- Quantity, integer
- Discount, float

DataBase Diagram







SQL Basic Concepts and Principles

Tables

- basic building unit of relational database
- fairly intuitive way of organizing data
 - has been around for centuries
- consists of rows and columns
 - called records and fields in database jargon
- each table has unique name in database
 - unique fully qualified name includes schema or database name as prefix

- the dot (.) notation in fully qualified name is commonly used in programming world to describe hierarchy of objects and their properties
- for example, table field in MS SQL Server database could be referred
`ACME.DBO.CUSTOMER.CUST_ID_N`
 - where ACME is database name
 - DBO is table owner (Microsoft standard)
 - CUSTOMER is name of table
 - CUST_ID_N is column name

column, domain

- each column has unique name within table, and any table must have at least one column
- records in table are not stored or retrieved in any particular order
- record is composed of number of cells, where each cell has unique name and might contain some data
- data within column must be of same type
 - for example, field AMOUNT contains only numbers
 - field DESCRIPTION, only words
- set of data within one field is said to be column's *domain*

Primary key

- primary role is to uniquely identify each record in table
 - based on idea of field (or fields) that contains set unique values
- *in the days of legacy databases, the records were always stored in some predefined order; if such an order had to be broken (because somebody had inserted records in a wrong order or business rule was changed), then the whole table (and, most likely, the whole database) had to be rebuilt*
- *RDBMS abolishes fixed order for records, but it still needs some mechanism of identifying the records uniquely, and primary key serves exactly this purpose*
- by its very nature, PK cannot be empty
 - means that in table with defined primary key, PK fields must contain data for each record

Primary key

- is a requirement to have primary key on each and every table
- many RDBMS implementations would warn you if you create table without defining PK
- purists go even further, specifying that PK should be *meaningless* in sense that they would use some generated unique value (like EMPLOYEE_ID) instead of, say, Social Security numbers (despite that these are unique as well)
- primary key could consist of one or more columns, i.e., though some fields may contain duplicate values, their combination (set) is unique through the entire table
- key that consists of several columns is called *composite key*

Relationships, Foreign key

- RDBMS is built upon parent/child relationship based solely on values in table columns
 - relationships meaningful in logical terms, not in low-level computer specific pointers
- take the example of our fictitious order entry database
- ORDER_HEADER table is related to CUSTOMER table since both of these tables have a *common set of values*
 - ORDHDR_CUSTID_FN (customer ID) in ORDER_HEADER (and its values) corresponds to
 - CUST_ID_N in CUSTOMER

Relationships, Foreign key

- CUST_ID_N is said to be *primary key* for CUSTOMER table
- and *foreign key* for ORDER_HEADER table
- ORDER_HEADER has its own primary key — ORDHDR_ID_N which uniquely identifies orders
- in addition it will have foreign key ORDHDR_CUSTID_FN field
- values in that field correspond to values in CUST_ID_N primary key field for CUSTOMER table
- unlike primary key, foreign key is not required to be unique
 - one customer could place several orders

Relationships, Foreign key

- by looking into ORDER_HEADER table you can find which customers placed particular orders
- table ORDER_HEADER became related to table CUSTOMER
- became easy to find customer based on orders, or find orders for customer
- no longer need to know database layout, order of records in table, or master some low-level pointers or proprietary programming language to query data
- possible to run ad-hoc queries formulated in standard English-like language — Structured Query Language

SELECT

SELECT

- SELECT Employees.FirstName,
Employees.LastName, Employees.BirthDate
- FROM Employees
- WHERE TitleOfCourtesy = 'Mr.'

- | • FirstName | Last Name | BirthDate |
|-------------|-----------|------------|
| • Steven | Buchanan | 1955-03-04 |
| • Michael | Suyama | 1963-07-02 |
| • Robert | King | 1960-05-29 |

Query Clauses

- Several components or *clauses* make up SELECT
- 1. SELECT - determines which columns to include in query's result set
- 2. FROM - identifies tables from which to draw data and how tables should be joined
- 3. WHERE - filters out unwanted data
- GROUP BY, HAVING, ORDER BY

Query Clauses

- SELECT, FROM, WHERE
- 4. GROUP BY - used to group rows together by common column values
- 5. HAVING –filters out unwanted groups
- 6. ORDER BY - sorts rows of final result set by one or more columns

SELECT * FROM Categories

- *Show me all the columns (*) and all the rows (No WHERE) in the Categories table*

- choose to include only subset of columns in the Categories table as well:
- SELECT CategoryName, Description
- FROM Categories
- *SELECT clause determines which of all possible columns should be included in query's result set*

Include in SELECT clause

- columns from table or tables named in FROM clause
- Expressions
 - such as `transaction.amount * -1`
- Function calls
 - built-in function calls
 - such as `ROUND(transaction.amount, 2)`
 - User-Defined Function calls

- SELECT ProductName, UnitPrice,
- UnitsInStock + UnitsOnOrder,
- ROUND(UnitPrice * (UnitsInStock + UnitsOnOrder), 2)
- FROM Products

Column Aliases

- SELECT ProductName,
- UnitPrice,
- UnitsInStock + UnitsOnOrder AS
TotalUnitsInStocks,
- ROUND(UnitPrice * (UnitsInStock +
UnitsOnOrder),2) TotalPriceInStocks
- FROM Products

Removing Duplicates

```
SELECT CustomerID  
FROM Orders
```

- 830 rows
- since some customers have more than one order, you will see same customer ID once for each order owned by that customer
- ALL keyword is default and never needs to be explicitly named

```
SELECT DISTINCT CustomerID  
FROM Orders
```

- 89 rows

from Clause

- list of one or more tables
- *from clause defines tables used by query, along with means of linking tables together*

Tables

- 1. Permanent tables (created using CREATE TABLE statement)
- 2. Temporary tables (rows returned by subquery)
- 3 Virtual tables (created using CREATE VIEW statement)

Subquery-generated tables

- is query contained within another query
- surrounded by parentheses and can be found in various parts of select statement
- within from clause subquery serves role of generating temporary table that is visible from all other query clauses and can interact with other tables named in from clause

Ms. Employees from WA

- SELECT e.FirstName, e.LastName
- FROM (SELECT Employees.FirstName,
Employees.LastName,
Employees.TitleOfCourtesy, Employees.Region
FROM Employees WHERE region = 'WA') e
- WHERE e.TitleOfCourtesy = 'Ms.'

Employees from WA

- SELECT
- Employees.FirstName, Employees.LastName, Employees.TitleOfCourtesy, Employees.Region
- FROM Employees
- WHERE region = 'WA')

Views

- query that is stored in data dictionary - looks and acts like table, but there is no data associated (stored) with view (*virtual* table)

- CREATE VIEW WAEmployees AS
- SELECT Employees.FirstName,
Employees.LastName,
Employees.TitleOfCourtesy, Employees.Region
- FROM Employees
- WHERE region = 'WA'
- SELECT FirstName, LastName
- FROM WAEmployees

Table Links

- second deviation from simple from clause definition is that if more than one table appears in from clause, conditions used to *link* tables must be included as well
- method of joining multiple tables
- SELECT Products.ProductName,
Categories.CategoryName
- FROM Products INNER JOIN Categories ON
Categories.CategoryID = Products.CategoryID

Table Aliases

- when multiple tables are joined in single query, you need way to identify which table you are referring to when you reference columns
- 1. use entire table name, such as Products.ProductName
- 2. assign each table *alias* and use alias throughout query

- SELECT Products.ProductName,
Categories.CategoryName FROM Products
INNER JOIN Categories ON
Categories.CategoryID = Products.CategoryID
- SELECT ProductName, CategoryName
- FROM Products INNER JOIN Categories ON
Categories.CategoryID = Products.CategoryID

- SELECT
- s.FirstName, s.LastName, s.Title,
- m.BossFirstName, m.BossLastName, m.BossTitle
- FROM Employees s INNER JOIN (SELECT
EmployeeID, FirstName AS BossFirstName ,
LastName AS BossLastName, Title AS BossTitle
FROM Employees) m
- ON s.ReportsTo = m.EmployeeID

Alias

- actually JOIN two copies of Employees, alias s and m, ON s.ReportsTo = m.EmployeeID
- SELECT s.FirstName, s.LastName, s.Title and
- m.FirstName, m.LastName, m.Title
- cause there will be conflicts add aliases
- m.BossFirstName, m.BossLastName, m.BossTitle

where Clause

- most of the time you will not wish to retrieve *every* row from table but will want way to filter out those rows that are not of interest
- where *clause mechanism for filtering out unwanted rows from your result set*
- contains single *filter condition*, but you can include many conditions as required, separated using operators such *and, or, not*

Ms. Employees from WA

- SELECT FirstName, LastName
- FROM Employees
- WHERE
- Region = 'WA' AND TitleOfCourtesy = 'Ms.'

Suppliers Products Categories

- SELECT
- Suppliers.CompanyName,
Products.ProductName,
Categories.CategoryName
- FROM Products
- JOIN Suppliers ON
Products.SupplierID=Suppliers.SupplierID
- JOIN Categories ON
Products.CategoryID=Categories.CategoryID

Suppliers of Beverages

- SELECT
- Suppliers.CompanyName, Products.ProductName, Categories.CategoryName
- FROM Products
- JOIN Suppliers ON Products.SupplierID=Suppliers.SupplierID
- JOIN Categories ON Products.CategoryID=Categories.CategoryID
- WHERE Categories.CategoryName = 'Beverages'

Suppliers of Condiments

- SELECT
- Suppliers.CompanyName, Products.ProductName, Categories.CategoryName
- FROM Products
- JOIN Suppliers ON Products.SupplierID=Suppliers.SupplierID
- JOIN Categories ON Products.CategoryID=Categories.CategoryID
- WHERE Categories.CategoryName = 'Condiments'

Suppliers of Beverages OR Condiments

- SELECT Suppliers.CompanyName
- FROM Products
- JOIN Suppliers ON
Products.SupplierID=Suppliers.SupplierID
- JOIN Categories ON
Products.CategoryID=Categories.CategoryID
- WHERE Categories.CategoryName = 'Beverages'
OR Categories.CategoryName = 'Condiments'

Gotcha ?!

- In programming, a gotcha is a feature of a system, a program or a programming language that works in the way it is documented but is counter-intuitive and almost invites mistakes because it is both enticingly easy to invoke and completely unexpected and/or unreasonable in its outcome.

Suppliers of Beverages OR Condiments

- SELECT DISTINCT Suppliers.CompanyName
- FROM Products
- JOIN Suppliers ON
Products.SupplierID=Suppliers.SupplierID
- JOIN Categories ON
Products.CategoryID=Categories.CategoryID
- WHERE Categories.CategoryName = 'Beverages'
OR Categories.CategoryName = 'Condiments'

- CompanyName
- Exotic Liquids
- Exotic Liquids
- Refrescos Americanas LTDA
- Bigfoot Breweries
- Bigfoot Breweries
- Aux joyeux ecclésiastiques
- Aux joyeux ecclésiastiques
- Leka Trading
- Bigfoot Breweries
- Pavlova, Ltd.
- Plutzer Lebensmittelgroßmärkte AG
- Karkki Oy
- Exotic Liquids
- New Orleans Cajun Delights
- New Orleans Cajun Delights
- Grandma Kelly's Homestead
- Grandma Kelly's Homestead
- Mayumi's
- Leka Trading
- Forêts d'érables
- Pavlova, Ltd.
- New Orleans Cajun Delights
- New Orleans Cajun Delights
- Plutzer Lebensmittelgroßmärkte AG

- CompanyName
- Aux joyeux ecclésiastiques
- Bigfoot Breweries
- Exotic Liquids
- Forêts d'érables
- Grandma Kelly's Homestead
- Karkki Oy
- Leka Trading
- Mayumi's
- New Orleans Cajun Delights
- Pavlova, Ltd.
- Plutzer Lebensmittelgroßmärkte AG
- Refrescos Americanas LTDA

Suppliers of Beverages AND Condiments

- SELECT DISTINCT Suppliers.CompanyName
- FROM Products
- JOIN Suppliers ON
Products.SupplierID=Suppliers.SupplierID
- JOIN Categories ON
Products.CategoryID=Categories.CategoryID
- WHERE Categories.CategoryName = 'Beverages'
AND Categories.CategoryName = 'Condiments'

- Exotic Liquids
- Leka Trading
- Pavlova, Ltd.
- Plutzer Lebensmittelgroßmärkte AG

Gotcha ?!

- In programming, a gotcha is a feature of a system, a program or a programming language that works in the way it is documented but is counter-intuitive and almost invites mistakes because it is both enticingly easy to invoke and completely unexpected and/or unreasonable in its outcome.

Suppliers of Beverages AND Condiments

- SELECT DISTINCT Suppliers.CompanyName
- FROM Suppliers
- JOIN Products p1 ON p1.SupplierID=Suppliers.SupplierID
- JOIN Categories c1 ON p1.CategoryID=c1.CategoryID
- JOIN Products p2 ON p2.SupplierID=Suppliers.SupplierID
- JOIN Categories c2 ON p2.CategoryID=c2.CategoryID
- WHERE
- c1.CategoryName = 'Beverages' AND
c2.CategoryName = 'Condiments'

Suppliers of Beverages AND Condiments

- SELECT DISTINCT Suppliers.CompanyName
- FROM Products
- JOIN Suppliers ON Products.SupplierID=Suppliers.SupplierID
- JOIN Categories ON Products.CategoryID=Categories.CategoryID
- WHERE Categories.CategoryName = 'Beverages'
- INTERSECT
- SELECT DISTINCT Suppliers.CompanyName
- FROM Products
- JOIN Suppliers ON Products.SupplierID=Suppliers.SupplierID
- JOIN Categories ON Products.CategoryID=Categories.CategoryID
- WHERE Categories.CategoryName = 'Condiments'

Suppliers of Beverages OR Condiments

- SELECT DISTINCT Suppliers.CompanyName
- FROM Products
- JOIN Suppliers ON Products.SupplierID=Suppliers.SupplierID
- JOIN Categories ON Products.CategoryID=Categories.CategoryID
- WHERE Categories.CategoryName = 'Beverages'
- UNION
- SELECT DISTINCT Suppliers.CompanyName
- FROM Products
- JOIN Suppliers ON Products.SupplierID=Suppliers.SupplierID
- JOIN Categories ON Products.CategoryID=Categories.CategoryID
- WHERE Categories.CategoryName = 'Condiments'

use parentheses to group conditions
together in WHERE clause

- use parentheses to group conditions together in WHERE clause
- Sales Representative hired after 1994
- and
- Sales Manager hired before 1994

- SELECT
- FirstName, LastName, Title, HireDate FROM Employees
- WHERE
- (Title = 'Sales Representative' AND HireDate > '1994-01-01')
- OR
- (Title = 'Sales Manager' AND HireDate < '1994-01-01')

group by and having Clauses

- all queries thus far have retrieved raw data without any manipulation
- sometimes you will want to find trends in your data that will require database server to prepare data before you retrieve your result set
- such mechanism is group by clause, which is used to group data by column values
- for example, rather than looking at list of employees and titles you might want to look at list of titles along with number of employees assigned to each
- when using group by clause, you may also use having clause, which allows you to filter group data in same way the where clause lets you filter raw data

- SELECT Title, COUNT(EmployeeID)
- FROM Employees GROUP BY Title

- Title (No column name)
- Inside Sales Coordinator 1
- Sales Manager 1
- Sales Representative 6
- Vice President, Sales 1

- SELECT Title, COUNT(EmployeeID)
 - FROM Employees
 - GROUP BY Title
 - HAVING COUNT(EmployeeID) > 1
-
- Title (No column name)
 - Sales Representative 6

order by Clause

- rows in result set returned from query are not in any particular order
- if you want your result set in particular order, you will need to instruct database server to sort results using order by clause:
- *order by clause is mechanism for sorting your result set using either raw column data or expressions based on column data*

- SELECT Products.ProductName,
Categories.CategoryName,
Suppliers.CompanyName
- FROM Products
- JOIN Suppliers ON
Products.SupplierID=Suppliers.SupplierID
- JOIN Categories ON
Products.CategoryID=Categories.CategoryID
- ORDER BY Products.ProductName

- SELECT Products.ProductName,
Categories.CategoryName,
Suppliers.CompanyName
- FROM Products
- JOIN Suppliers ON
Products.SupplierID=Suppliers.SupplierID
- JOIN Categories ON
Products.CategoryID=Categories.CategoryID
- ORDER BY Categories.CategoryName,
Products.ProductName

Sorting, Group By via Expressions

- `SELECT RIGHT (Products.ProductName, 1),
COUNT(*)`
- `FROM Products`
- `GROUP BY RIGHT (Products.ProductName, 1)`
- `ORDER BY COUNT(*) DESC`

•	(No column name)	(No column name)
•	e	17
•	s	9
•	t	9
•	a	7
•	d	6
•	u	6
•	r	5
•	i	4
•	n	3
•	g	3
•	x	3
•	o	2
•	l	2
•	p	1

Sorting via Numeric Placeholders

- can reference columns, in sorting, by their *position* in select clause rather than by name
- SELECT RIGHT (Products.ProductName, 1),
COUNT(*)
- FROM Products
- GROUP BY RIGHT (Products.ProductName, 1)
- ORDER BY 2 DESC

FILTER

- where clause may contain one or more *conditions*, separated by
- binary operators *and* and *or* or *or*
- unary operator *not*
- if where clause includes three or more conditions using both *and* and *or* or *not* operators, should use parentheses to make your intent clear

- typically difficult for person to evaluate where clause that includes not operator, which is why you won't encounter it very often
- WHERE NOT (TitleOfCourtesy = 'Mr.')
- WHERE TitleOfCourtesy != 'Mr.'

Condition

- made up of one or more *expressions* coupled with one or more *operators*
- expression can be any of the following:
 - number
 - string literal
 - column in table or view
 - built-in function
 - list of expressions, such as ('Mrs.', 'Ms.', 'Mr.')
 - subquery
- operators used within conditions include:
 - comparison operators
 - arithmetic operators

Condition

- large percentage of filter conditions will be ***equality condition*** of the form
 - '*column = expression*'
 - '*expression = expression*'
- another fairly common type of condition is ***inequality condition***, which asserts that two expressions are *not* equal
 - '*expression != expression*'
 - '*expression <> expression*'

Condition

- you can build conditions that check whether expression falls within certain ***range*** (common when working with numeric or temporal data)
- ranges of dates and numbers are easy to understand, you can also build conditions that search for ranges of strings
- `SELECT FirstName, LastName FROM Employees`
- `WHERE HireDate > '1993-01-01' AND
HireDate < '1994-01-01'`

Condition

- when you have *both* an upper and lower limit for your range, you may choose to use the ***between*** operator rather than using two separate conditions
- `SELECT FirstName, LastName FROM Employees`
- `WHERE HireDate BETWEEN ('1993-01-01', '1994-01-01')`

Condition

- in some cases, you will not be restricting an expression to single value or range of values, but rather to finite set of values
- *in* operator checks *membership*
- SELECT FirstName, LastName FROM Employees
- WHERE TitleOfCourtesy = 'Mrs.' OR
TitleOfCourtesy = 'Ms.'
- WHERE TitleOfCourtesy IN ('Mrs.', 'Ms.')

Condition

- with *in* operator, you can write single condition no matter how many expressions are in set
- along with writing your own set of expressions you can use subquery to generate set for you
- sometimes you want to see whether particular expression exists within set of expressions, and sometimes you want to see whether expression does *not* exist
- can use *not in* operator

Condition

- `SELECT ProductName, QuantityPerUnit, UnitsInStock, ReorderLevel`
- `FROM Products`
- `WHERE SupplierID IN (7,24)`
- `SELECT ProductName, QuantityPerUnit, UnitsInStock, ReorderLevel`
- `FROM Products`
- `WHERE SupplierID IN (`
- `SELECT SupplierID FROM Suppliers WHERE Country = 'Australia')`

Condition

- when searching for partial string matches, you might be interested in:
 - strings beginning/ending with certain character
 - strings beginning/ending with substring
 - strings containing certain character anywhere within string
 - strings containing substring anywhere within string
 - strings with specific format, regardless of individual characters
- can build search expressions to identify these and many others partial string ***matches*** by using ***wildcard*** characters

Condition

- wildcard character `_` matches exactly one character
- underscore character takes place of single character
- wildcard character `%` matches any number of characters (including 0)
- percent sign can take place of variable number of characters
- when building conditions that utilize search expressions, you use *like* operator

Condition

- `SELECT FirstName, LastName, Title FROM Employees WHERE EmployeeID IN (`
- `SELECT DISTINCT ReportsTo FROM Employees)`
- `SELECT FirstName, LastName, Title FROM Employees WHERE`
- `Title LIKE '%president%' OR`
- `Title LIKE '%manager%'`

Null

- Null is absence of value, various flavors of null:
- *Not applicable*
 - such as employee ID column for transaction that took place at ATM machine
- *Value not yet known*
 - such as federal ID is not known at the time customer row is created
- *Value undefined*
 - such as account is created for product that has not yet been added to the database

- expression can *be* null, but it can never *equal* null
- two nulls are never equal to each other
- test whether expression is null, you need to use the ***is null*** operator
- SELECT FirstName, LastName, Title
- FROM Employees
- WHERE ReportsTo IS NULL
- WHERE ReportsTo IS NOT NULL

- SELECT FirstName, LastName, Title FROM Employees
- WHERE ReportsTo !=2
- account for possibility that some rows might contain null in ReportsTo column
- SELECT FirstName, LastName, Title FROM Employees
- WHERE ReportsTo !=2 OR ReportsTo IS NULL

Null Values and Three Valued Logic

- Any comparison with *null* returns *unknown*
 - E.g. $5 < null$ or $null <> null$ or $null = null$
- Three-valued logic using the truth value *unknown*:
 - OR: $(unknown \text{ or } true) = true$, $(unknown \text{ or } false) = unknown$
 $(unknown \text{ or } unknown) = unknown$
 - AND: $(true \text{ and } unknown) = unknown$, $(false \text{ and } unknown) = false$,
 $(unknown \text{ and } unknown) = unknown$
 - NOT: $(\text{not } unknown) = unknown$
 - “*P* is unknown” evaluates to true if predicate *P* evaluates to *unknown*
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*

Null Values and Aggregates

- Total all amounts

```
select sum (amount)  
from table
```

- Above statement ignores null amounts
 - result is null if there is no non-null amount, that is the
- All aggregate operations except **count(*)** ignore tuples with null values on the aggregated attributes.

SQL Functions

Standard SQL Functions

- BIT_LENGTH (expression)
 - returns the length of the expression, usually string, in bits.
- CAST (value AS data type)
 - converts supplied value from one data type into another *compatible* data type
- CHAR_LENGTH (expression)
 - returns length of expression, usually string, in characters
- CONVERT (expression USING conversion)
 - returns string converted according to rules specified in conversion parameter
- CURRENT_DATE
 - returns current date of system

Standard SQL Functions

- CURRENT_TIME (precision)
 - returns current time of system, of specified precision
- CURRENT_TIMESTAMP (precision)
 - returns current time and current date of system, of specified precision
- EXTRACT (part FROM expression)
 - extracts specified named part of expression
- LOWER (expression)
 - converts character string from uppercase (or mixed case) into lowercase letters
- OCTET_LENGTH (*expression*)
 - returns length of expression in *bytes* (byte containing 8 bits)

Standard SQL Functions

- POSITION (*char expression IN source*)
 - returns position of the char expression in source.
- SUBSTRING (*string expression, start, length*)
 - returns string part of *string expression*, from *start* position up to specified *length*
- TRANSLATE (*string expression USING translation rule*)
 - returns string translated into another string according to specified rules
- TRIM(LEADING | TRAILING | BOTH *char expression* FROM *string expression*)
 - returns string from string expression where *leading, trailing, or both* char expression characters are removed
- UPPER (*expression*)
 - converts character string from lowercase (or mixed case) into uppercase letters

Numeric functions

- ABS (n)
 - returns absolute value of a number n
- CEILING (n)
 - returns smallest integer that is greater than or equal to n .
- EXP (n)
 - returns exponential value of n
- FLOOR (n)
 - returns largest integer less than or equal to n
- MOD (n,m) or %
 - returns remainder of n divided by m

Numeric functions

- POWER.(*m*,*n*)
 - returns value of *m* raised into *n*th power
- RAND (*n*)
 - returns a random number between 0 and 1
- ROUND (*n*,*m*,[0])
 - returns number *n* rounded to *m* decimal places; last argument - 0 - is a default
 - similar to TRUNCate
- SIGN(*n*)
 - returns -1, if *n* is negative number, 1 if it is positive number, and 0 if number 0

String functions

- ASCII (string)
 - returns ASCII code of the first character of string
- CHAR (number) NCHAR (number)
 - returns character for ASCII code
- CONCAT (string1, string2) or '+'
 - returns result of concatenation of two strings
- CHARINDEX (string1,string2, n)
 - returns position of occurrence of substring within string
- LEFT (string, n)
 - returns n number of characters starting from left
- LENGTH (string) or LEN (string)
 - returns number of characters in string

String functions

- DATALENGTH (expression)
 - returns number of bytes in expression, which could be any data type
- LTRIM (string)
 - returns string with leading blank characters removed
- REPLACE (string1, string2, string3)
 - replaces all occurrences of *string1* within *string2* with *string3*

String functions

- RTRIM (string)
 - returns string with trailing blank characters removed
- STR (expression)
 - converts argument expression into a character string
- SUBSTRING (string, n , m)
 - returns a part of a string starting from n^{th} character for the length of m characters

Date and time functions

MS SQL Server 2000

- DATEADD (month, number, date)
 - returns date plus date part (year, month, day)
- GETDATE
 - returns current date in session's time zone
- CONVERT or CAST
 - returns date from value according to specific format
- DAY (MONTH, YEAR)
 - returns DAY part (integer) of specified datetime expression

Date and time functions

MS SQL Server 2000

- DATEPART (date part, datetime)
 - returns requested date part (day, month, year)
- DATEDIFF
 - calculates difference between two dates
- DATEADD (day, n, m)
 - calculates what day would be next relative to some other supplied date
- DATEADD (datepart, n, m)
 - calculates what date would be next relative to some other supplied date

Time zone functions

- deal with Earth's different time zones
- functions always return time zone in which machine is located

Miscellaneous functions

- COALESCE (expression1, expression2, expression3 ...)
 - returns first argument on list that is not NULL
- CASE (expression)
WHEN <compare value>
THEN <substitute value>
ELSE END
 - compares input expression to some predefined values, and outputs substitute value, either hard coded or calculated
- ISNULL (expression, value)
 - checks whether expression is NULL, and if it is returns specified value

Querying Multiple Tables

JOIN

JOIN

- queries against single table are certainly not rare, but you will find that most of your queries will require two, three, or even more tables

SELECT * FROM Categories

- CategoryID CategoryName Description
- 1 Beverages Soft drinks, coffees, teas, beers, and ales
- 2 Condiments Sweet and savory sauces, relishes, spreads, and seasonings
- 3 Confections Desserts, candies, and sweet breads
- 4 Dairy Products Milk and Cheeses
- 5 Grains/Cereals Breads, crackers, pasta, and cereal
- 6 Meat/Poultry Prepared meats
- 7 Produce Dried fruit and bean curd
- 8 Seafood Seaweed and fish

SELECT * FROM Products

- | | ProductID | ProductName | CategoryID | QuantityPerUnit |
|-------|---------------------------------|--------------|---------------------|-----------------|
| | UnitPrice | UnitsInStock | UnitsOnOrder | ReorderLevel |
| • 1 | Chai | 1 | 10 boxes x 20 bags | 18.00 |
| • 2 | Chang | 1 | 24 - 12 oz bottles | 19.00 |
| • 3 | Aniseed Syrup | 2 | 12 - 550 ml bottles | 10.00 |
| • 4 | Chef Anton's Cajun Seasoning | 2 | 48 - 6 oz jars | 22.00 |
| • 5 | Chef Anton's Gumbo Mix | 2 | 36 boxes | 21.35 |
| • 6 | Grandma's Boysenberry Spread | 2 | 12 - 8 oz jars | 25.00 |
| • 7 | Uncle Bob's Organic Dried Pears | 7 | 12 - 1 lb pkgs. | 30.00 |
| • 8 | Northwoods Cranberry Sauce | 2 | 12 - 12 oz jars | 40.00 |
| • ... | | | | |

- retrieve data from both tables - answer lies in Products.CategoryID which holds ID of category to which each products is assigned (in more formal terms, Products.CategoryID column is *Foreign Key* to Category table – values match Primary Key)

Cartesian Product

- put Product and Category tables into from clause of query and
- SELECT Products.ProductID, Products.ProductName, Products.CategoryID, Categories.CategoryID, Categories.CategoryName
- FROM Products, Categories
- because query didn't specify *how* tables should be joined, server generated *Cartesian product*, which is *every* permutation of two tables
- (8 Categories × 77 Products = 616 permutations)

- | ProductID | ProductName | CategoryID | CategoryID | CategoryName |
|-----------|---------------------------------|------------|------------|--------------|
| 1 | Chai | 1 | 1 | Beverages |
| 2 | Chang | 1 | 1 | Beverages |
| 3 | Aniseed Syrup | 2 | 1 | Beverages |
| 4 | Chef Anton's Cajun Seasoning | 2 | 1 | Beverages |
| 5 | Chef Anton's Gumbo Mix | 2 | 1 | Beverages |
| 6 | Grandma's Boysenberry Spread | 2 | 1 | Beverages |
| 7 | Uncle Bob's Organic Dried Pears | 7 | 1 | Beverages |
| 8 | Northwoods Cranberry Sauce | 2 | 1 | Beverages |
| ... | | | | |

Cartesian Product

- this type of join is known as ***cross join***, and it is rarely used

Inner Joins

- to modify previous query so that 77 rows are included in result set (one for each product), you need to describe how two tables are related

Inner Joins

- `SELECT Products.ProductID,
Products.ProductName, Products.CategoryID,
Categories.CategoryID, Categories.CategoryName`
- `FROM Products, Categories WHERE
Products.CategoryID = Categories.CategoryID`
- or
- `SELECT Products.ProductID,
Products.ProductName, Products.CategoryID,
Categories.CategoryID, Categories.CategoryName`
- `FROM Products INNER JOIN Categories ON
Products.CategoryID = Categories.CategoryID`

- | ProductID | ProductName | CategoryID | CategoryID |
|-----------|---------------------------------|------------|--------------|
| | CategoryName | | |
| 1 | Chai | 1 | 1 Beverages |
| 2 | Chang | 1 | 1 Beverages |
| 3 | Aniseed Syrup | 2 | 2 Condiments |
| 4 | Chef Anton's Cajun Seasoning | 2 | 2 Condiments |
| 5 | Chef Anton's Gumbo Mix | 2 | 2 Condiments |
| 6 | Grandma's Boysenberry Spread | 2 | 2 Condiments |
| 7 | Uncle Bob's Organic Dried Pears | 7 | 7 Produce |
| 8 | Northwoods Cranberry Sauce | 2 | 2 Condiments |
| ... | | | |

Inner Joins

- if value exists for CategoryID column in one table but *not* the other, then join fails for rows containing that value and those rows are excluded from result set
- this type of join is known ***inner join***, and it is most commonly used type of join
- if you want to include all rows from one table or the other regardless of whether match exists, you need to specify *outer join*
- if you do not specify type of join, server will do inner join by default

- SELECT Products.ProductID, Products.ProductName, Products.CategoryID, Categories.CategoryID, Categories.CategoryName
- FROM Products INNER JOIN Categories ON Products.CategoryID = Categories.CategoryID
- WHERE Products.ProductName LIKE 'q%'
- Join conditions and filter conditions are separated into two different clauses (on subclause and where clause, respectively), making query easier to understand

Joining Three or More Tables

- is similar to joining two tables
- with two-table join, there are two tables and one join type in from clause, and a single on subclause to define how tables are joined
- with three-table join, there are three tables and two join types in from clause, and two on subclauses

- SELECT Products.ProductID, Products.ProductName, Products.CategoryID, Categories.CategoryID, Categories.CategoryName, Products.SupplierID, Suppliers.SupplierID, Suppliers.CompanyName
- FROM Products INNER JOIN Categories ON Products.CategoryID = Categories.CategoryID
- INNER JOIN Suppliers ON Products.SupplierID = Suppliers.SupplierID
- WHERE Products.ProductName LIKE 'q%'

- order in which tables are named it is irrelevant
- Products Categories Suppliers
- Products Suppliers Categories
- Categories Products Suppliers
- Suppliers Products Categories
- will get exact same results

- order in which tables are named it is irrelevant
- Categories Suppliers Products
- Suppliers Categories Products
- will not work because we cannot JOIN
Categories and Suppliers – there is no link

- order in which tables are named it is irrelevant
- using statistics gathered from your database objects, server must pick one of three tables as a starting point (chosen table is known as *driving table*), and then decide in which order to join remaining tables
- therefore, order in which tables appear in your from clause is not significant
- if, however, you believe that tables in your query should always be joined in particular order, you can place tables in the desired order and then
 - specify keyword STRAIGHT_JOIN in MySQL
 - request FORCE ORDER option in SQL Server

Using Subqueries As Tables in JOIN

- SELECT Products.ProductID, Products.ProductName, Products.CategoryID, Categories.CategoryID, Categories.CategoryName, Products.SupplierID, Suppliers.SupplierID, Suppliers.CompanyName
- FROM Products INNER JOIN Categories ON Products.CategoryID = Categories.CategoryID
- INNER JOIN Suppliers ON Products.SupplierID=Suppliers.SupplierID
- WHERE Categories.CategoryName = 'Beverages' AND
- (Suppliers.CompanyName LIKE '%Liquid%' OR Suppliers.CompanyName LIKE '%Brew%')

- SELECT Products.ProductID, Products.ProductName, Products.CategoryID, c.CategoryID, c.CategoryName, Products.SupplierID, s.SupplierID, s.CompanyName
- FROM Products INNER JOIN
- (SELECT * FROM Categories WHERE Categories.CategoryName = 'Beverages') c
- ON Products.CategoryID = c.CategoryID
- INNER JOIN
- (SELECT * FROM Suppliers WHERE Suppliers.CompanyName LIKE '%Liquid%' OR Suppliers.CompanyName LIKE '%Brew%') s
- ON Products.SupplierID = s.SupplierID

- first subquery is given alias c, finds all Beverages
- second subquery is given alias s, finds suppliers of liquids and brews
- results are the same
- lack of where clause in main query; since all filter conditions are all inside subqueries

Performance

- JOIN 77 rows (Products) with 8 rows (Categories) with 29 rows (Suppliers)
- filter 17.864 rows, result 5 rows
- JOIN 77 rows (Products) with 1 rows (Categories) with 2 rows (Suppliers)
- result 5 rows

Using Same Table Twice: Self-Joins

- for this to work, you will need to give each instance of table different alias so that server knows which one you are referring to in various clauses
- `SELECT s.EmployeeID, s.FirstName, s.LastName, s.Title, s.ReportsTo, m.EmployeeID, m.FirstName, m.LastName, m.Title`
- `FROM Employees s INNER JOIN Employees m ON s.ReportsTo = m.EmployeeID`

Using Same Table Twice: Self-Joins

- not only can you include same table more than once in same query, but you can actually join table to itself
- Employees table, for example, includes *selfreferencing Foreign Key*, which means that it includes column ReportsTo that points to *Primary Key* within same table
 - column points to employee's manager

•	EmployeeID	FirstName	LastName	Title	ReportsTo	Title
	EmployeeID	FirstName	LastName			
• 1	Nancy Fuller	Davolio	Sales Representative Vice President, Sales	2	2	Andrew
• 3	Janet Fuller	Leverling	Sales Representative Vice President, Sales	2	2	Andrew
• 4	Margaret Fuller	Peacock	Sales Representative Vice President, Sales	2	2	Andrew
• 5	Steven Fuller	Buchanan	Sales Manager Vice President, Sales	2	2	Andrew
• 6	Michael Buchanan	Suyama	Sales Representative Sales Manager	5	5	Steven
• 7	Robert Buchanan	King	Sales Representative Sales Manager	5	5	Steven
• 8	Laura Fuller	Callahan	Inside Sales Coordinator Vice President, Sales	2	2	Andrew
• 9	Anne Steven	Dodsworth	Sales Representative Sales Manager	5	5	

Equi-Joins vs Non-Equi-Joins

- all of queries shown thus far have employed ***equi-joins***, meaning that values from two tables must match for join to succeed
- ***equi-join*** always employs an equals sign ON
Products.CategoryID =Categories.CategoryID
- for example, let's say that managers has decided to have tennis tournament for all their people and you have been asked to create list of all pairings

- SELECT m.FirstName, m.LastName,
- s1.FirstName + ' ' + s1.LastName,
- ' vs. ',
- s2.FirstName + ' ' + s2.LastName
- FROM Employees m
- INNER JOIN Employees s1 ON s1.ReportsTo = m.EmployeeID
- INNER JOIN Employees s2 ON s2.ReportsTo = m.EmployeeID
- WHERE s2.LastName <> s1.LastName

Non-Equi-Joins

- `SELECT s1.FirstName + ' ' + s1.LastName,`
- `' vs. ', s2.FirstName + ' ' + s2.LastName`
- `FROM Employees s1 INNER JOIN Employees s2`
`ON s1.EmployeeID != s2.EmployeeID`
- Join Employees table to itself and return all rows where EmployeeID don't match (since person can't play tennis against himself)

- problem is that for each pairing (e.g., Andrew Fuller vs. Nancy Davolio), there is also a reverse pairing (e.g., Nancy Davolio vs. Andrew Fuller)
- way to achieve desired results is to use join condition
- `SELECT s1.FirstName + ' ' + s1.LastName,`
- `' vs. ', s2.FirstName + ' ' + s2.LastName`
- `FROM Employees s1 INNER JOIN Employees s2`
`ON s1.EmployeeID < s2.EmployeeID`
- $36 \text{ rows} = 9 * 8 / 2$

SETS

Set Operators

- SQL language includes three set operators
- each set operator has two flavors, one that includes duplicates and another that removes duplicates (but not necessarily *all* of duplicates)
- data sets should be union compatible
- union
- intersection
- except

union Operator

- allow you to combine multiple data sets
- union sorts combined set and removes duplicates,
- union all does not (number of rows in final data set will always equal sum of number of rows in sets)

- SELECT DISTINCT Suppliers.CompanyName
- FROM Categories JOIN Products ON
Categories.CategoryID = Products.CategoryID
- JOIN Suppliers ON Products.SupplierID =
Suppliers.SupplierID
- WHERE CategoryName = 'Beverages'

- CompanyName
 1. Aux joyeux ecclésiastiques
 2. Bigfoot Breweries
 3. Exotic Liquids
 4. Karkki Oy
 5. Leka Trading
 6. Pavlova, Ltd.
 7. Plutzer Lebensmittelgroßmärkte AG
 8. Refrescos Americanas LTDA

- SELECT DISTINCT Suppliers.CompanyName
- FROM Categories JOIN Products ON
Categories.CategoryID = Products.CategoryID
- JOIN Suppliers ON Products.SupplierID =
Suppliers.SupplierID
- WHERE CategoryName = 'Condiments'

- CompanyName
 1. Exotic Liquids
 2. Forêts d'érables
 3. Grandma Kelly's Homestead
 4. Leka Trading
 5. Mayumi's
 6. New Orleans Cajun Delights
 7. Pavlova, Ltd.
 8. Plutzer Lebensmittelgroßmärkte AG

- SELECT DISTINCT Suppliers.CompanyName
- FROM Categories JOIN Products ON Categories.CategoryID = Products.CategoryID
- JOIN Suppliers ON Products.SupplierID = Suppliers.SupplierID
- WHERE CategoryName = 'Beverages'
- UNION
- SELECT DISTINCT Suppliers.CompanyName
- FROM Categories JOIN Products ON Categories.CategoryID = Products.CategoryID
- JOIN Suppliers ON Products.SupplierID = Suppliers.SupplierID
- WHERE CategoryName = 'Condiments'

- CompanyName
 1. Aux joyeux ecclésiastiques
 2. Bigfoot Breweries
 3. Exotic Liquids
 4. Forêts d'érables
 5. Grandma Kelly's Homestead
 6. Karkki Oy
 7. Leka Trading
 8. Mayumi's
 9. New Orleans Cajun Delights
 10. Pavlova, Ltd.
 11. Plutzer Lebensmittelgroßmärkte AG
 12. Refrescos Americanas LTDA

intersect Operator

- SELECT DISTINCT Suppliers.CompanyName
- FROM Categories JOIN Products ON Categories.CategoryID = Products.CategoryID
- JOIN Suppliers ON Products.SupplierID = Suppliers.SupplierID
- WHERE CategoryName = 'Beverages'
- INTERSECT
- SELECT DISTINCT Suppliers.CompanyName
- FROM Categories JOIN Products ON Categories.CategoryID = Products.CategoryID
- JOIN Suppliers ON Products.SupplierID = Suppliers.SupplierID
- WHERE CategoryName = 'Condiments'

- CompanyName
 1. Exotic Liquids
 2. Leka Trading
 3. Pavlova, Ltd.
 4. Plutzer Lebensmittelgroßmärkte AG

except Operator

- Find Customers who Ordered Products from all Categories

CREATE VIEW CustomCateg AS

- SELECT Customers.CompanyName,
Categories.CategoryName
- FROM Customers JOIN Orders ON
Orders.CustomerID = Customers.CustomerID
- JOIN OrderDetails ON OrderDetails.OrderID =
Orders.OrderID
- JOIN Products ON Products.ProductID =
OrderDetails.ProductID
- JOIN Categories ON Categories.CategoryID =
Products.CategoryID

SELECT * FROM CustomCateg

• CompanyName	CategoryName
1. Alfreds Futterkiste	Produce
2. Alfreds Futterkiste	Beverages
3. Alfreds Futterkiste	Seafood
4. Alfreds Futterkiste	Condiments
5. Alfreds Futterkiste	Dairy Products
6. Ana Trujillo Emparedados y helados	Dairy Products
7. Ana Trujillo Emparedados y helados	Beverages
8. Ana Trujillo Emparedados y helados	Produce
9. Ana Trujillo Emparedados y helados	Grains/Cereals
10. Ana Trujillo Emparedados y helados	Seafood
11. ...	

- `SELECT CategoryName FROM CustomCateg
WHERE CompanyName = 'Alfreds Futterkiste'`
- CategoryName
 1. Beverages
 2. Condiments
 3. Dairy Products
 4. Produce
 5. Seafood

- `SELECT CategoryName FROM CustomCateg WHERE
CompanyName = 'QUICK-Stop'`
- CategoryName
 1. Beverages
 2. Condiments
 3. Confections
 4. Dairy Products
 5. Grains/Cereals
 6. Meat/Poultry
 7. Produce
 8. Seafood

- SELECT Categories.CategoryName FROM Categories EXCEPT
- SELECT CategoryName FROM CustomCateg WHERE CompanyName = 'Alfreds Futterkiste'
- CategoryName
 1. Confections
 2. Grains/Cereals
 3. Meat/Poultry

- `SELECT Categories.CategoryName FROM Categories EXCEPT`
- `SELECT CategoryName FROM CustomCateg WHERE CompanyName = 'QUICK-Stop'`
- `CategoryName`

- result it is the \emptyset empty set

- as always it is not possible to close whiteout
thank you for your kindly attention!
- *If you get half as much pleasure - the guilty variety, to be sure - from reading this slides as I get from writing it, we're all doing pretty well.*