

The background features a large, light blue watermark logo of the Technical University of Cluj-Napoca. The logo consists of a shield-like shape with stylized vertical bars inside, and the text 'TECHNICAL UNIVERSITY' at the top, 'OF CLUJ-NAPOCA' in the middle, and 'Computer Science' at the bottom.

Fundamental Algorithms

Lecture #7

Cluj-Napoca

November 13, 2019

Agenda

- **Red-Black Trees**
- **Disjoint Sets**

Red-Black trees

- Balanced trees
- Both insert/delete operations take $O(\lg n)$, with constant time for rebalancing

Def: is a BST with the following properties:

P_0 : the root is **black**

P_1 : each node is colored either **black** or **red**

P_2 : each leaf (nil!) is **black**

P_3 : both children of a red node are **black**

P_4 : every path from any node to a leaf has the **same number of black nodes** (black height)

Red-Black trees

Theorem: A RB tree with n internal nodes has height at most $2\lg(n+1)$

Proof: Denote **bh(x)** – the black height (without x) of node x

Step 1: We will prove the following statement, denoted $P(bh)$:

$P(bh)$: $\forall x \in RBT$, the tree rooted by x has at least $2^{bh(x)} - 1$ nodes

Induction:

$P(0) \quad 2^0 - 1 = 1$

Assume $P(bh)$ true $\Rightarrow P(bh+1)$ true?

x has 2 children; each has the black height $bh(x)$ (if x is red) or
 $bh(x)-1$ (if x is black)

nb of internal nodes of x = nb of internal nodes of children(x) + 1 (itself) \Rightarrow

\Rightarrow at least $(2^{bh(x)-1} - 1) + (2^{bh(x)-1} - 1) + 1 = 2^{bh(x)} - 1$ nodes — *q.e.d* (end of step 1)

Red-Black trees

Step 2:

We know $P(bh)$ is true, ie

$P(bh)$: $\forall x \in \text{RBT}$, the tree rooted by x has at least $2^{bh(x)} - 1$ nodes (1)

By P_3 of RBT def (use contradiction to prove) $bh(x) \geq h/2$ (2)

//since after each red node comes a black one

$$\Rightarrow n \geq 2^{bh(x)} - 1 \quad (\text{from (1)})$$


$$\geq 2^{h/2} - 1 \quad (\text{from (2)})$$

$$n \geq 2^{h/2} - 1 \Leftrightarrow n+1 \geq 2^{h/2} \Leftrightarrow h/2 \leq \lg(n+1) \Leftrightarrow h \leq 2\lg(n+1)$$


(q.e.d., end of Theorem proof)

Red-Black trees - rotations

Similar to single rotations for AVL
They are symmetric



Red-Black trees – rotations



Left-Rotate(T, x)

```
//x root of rotation (points on k1)
y<-right[x]           //y saves k2
right[x]<-left[y]      //right of k1 goes on beta
if left[y]<>nil        //if beta exists = is not nil
    then p[left[y]]<-x //beta's parent becomes k1
p[y]<-p[x]             //k2's parent what was k1's parent
if p[y]=nil           //k1 used to be the root of the tree
    then root[T]<-y
    else if x=left[p[x]] // the parent of k1 becomes the parent of k2
        then left[p[x]]<-y
        else right[p[x]]<-y

left[y]<-x             //k1 goes the left child of k2
p[x]<-y               //k2 becomes the parent of k1
```

RB-INSERT


- **Insert** like in ANY other BST
 - As a LEAF, as for any other BST
- **Assign** it a color
 - **RED** (why?)
- Check the properties
- Re-balance if needed (**RB-INSERT-FIXUP** – check the textbook for the complete code)
- P_3 both children of a red node are black
 - True for the children (*nils*) of the inserted node
 - Not true for the inserted node in case its parent is **RED**
- 3 cases to analyze and remove inconsistencies

RB-INSERT-FIXUP - Case#1

- **B** – inserted node (pointed by **x**)
- Parent (**A**)=RED, uncle (**D**)=RED, grandparent (**C**)=BLACK
- $\alpha, \beta, \gamma, \delta, \varepsilon$ are RB trees (β, γ empty at first)
- Swap colors between grandparent (**C**) and parent/uncle


```
parent<-black
uncle<-black
grandparent<-red
x<-grandparent
```

```
//no more P3 conflict A-B
//to preserve P4
//to preserve P4
```



RB-INSERT-FIXUP - Case#1 - eval

- P_3 may still be invalid, for the new x (i.e. C)
- Problem transferred 2 levels up in the tree (now β , γ not empty any longer)
- It takes (in the worst case) $O(h)$ to rebalance ($2\lg(n+1)/2 = \lg n$)




RB-INSERT-FIXUP - Case#2

- **B – inserted node (pointed by x)**
- Parent(**A**)=**RED**, uncle (δ 's root)=**BLACK** (here is the difference compared to case #1), grandparent (C)=**BLACK**
- $\alpha, \beta, \gamma, \delta$ are RB trees; δ 's root is **BLACK**

left_rotate(p(x))


//no more P3 conflict B-parent conflict

x ← left(x)



RB-INSERT-FIXUP - Case#2-eval

- Case #2 takes just $O(1)$ to apply, but
- P_3 is still invalid, for the new x denoted node (i.e. A-B conflict)
- \Rightarrow case #3 to the rescue




RB-INSERT-FIXUP - Case#3

- **Inserted A** /coming from #2 (node pointed by x)
- Parent (**B**)=**RED**, uncle (γ 's root)= **BLACK**, grandparent (**C**)=**BLACK**
- $\alpha, \beta, \gamma, \delta$ are RB trees

parent <- black

grandparent <- red

right_rotate(p[p[x]])



RB-INSERT-FIXUP - Case#3 - eval

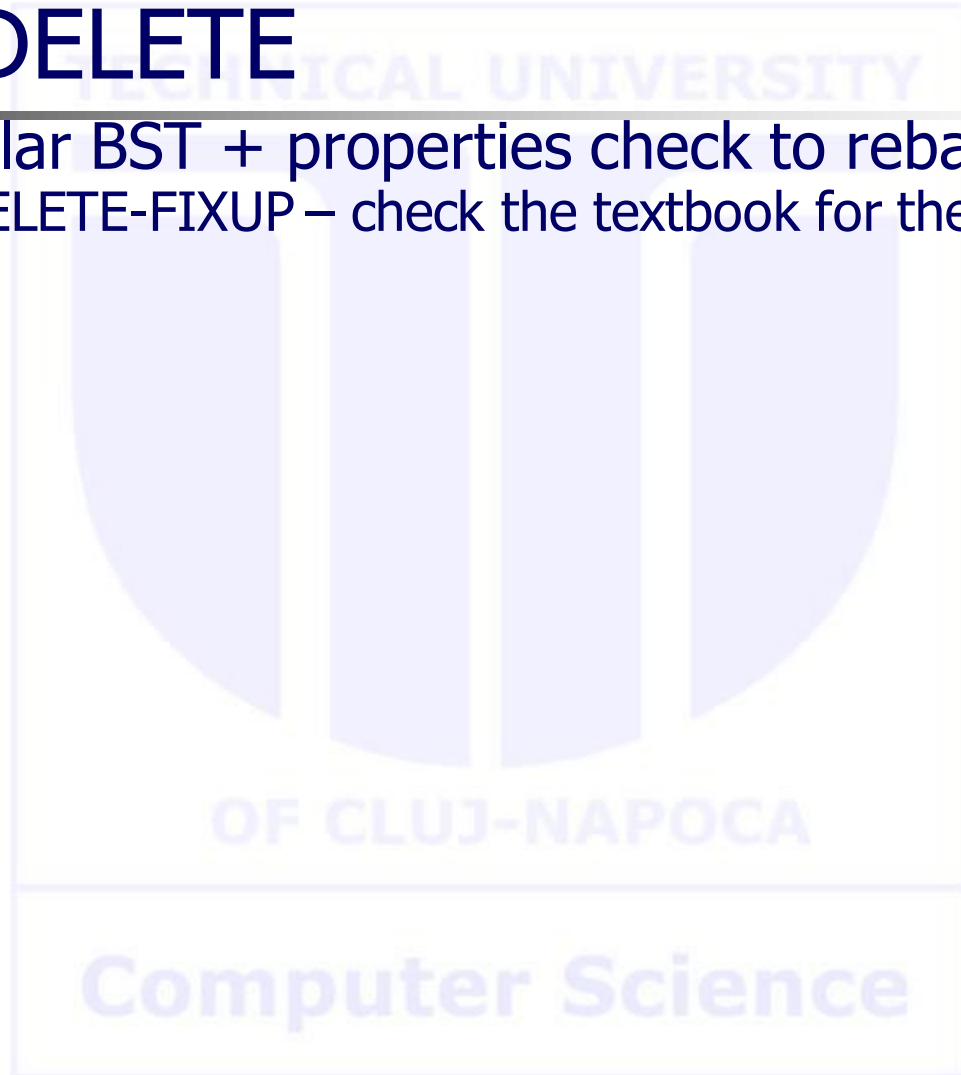
- Problem solved
- Each individual case takes $O(1)$
- Case #1 may repeat (up in the tree)
- Case #2 is followed by #3
- Case #3 solves the problem

RB-insert – Rebalancing eval

- Case #1 repeats up to the root $O(h)$
- Case #2+#3 \Rightarrow problem fixed $O(1)$
- Case #3 \Rightarrow problem fixed $O(1)$
- Insert $O(\lg n)$ + rebalancing
 - Worst case: #1 repeats $O(\lg n)$
 - Best case: #3 \Rightarrow 1 rotation $O(1)$
 - Other case: #2+3 \Rightarrow 2 rotations $O(1)$
- **$O(\lg n)$ overall worst time** (case 1 repeats), **at most 2 rotations** (case 2)

RB-DELETE

- Del as in regular BST + properties check to rebalance, if needed (RB-DELETE-FIXUP – check the textbook for the code)



RB-DELETE

- Del as in regular BST + properties check to rebalance, if needed (RB-DELETE-FIXUP – check the textbook for the code)

P4 (black height) is an issue!!

RB-DELETE

- Del as in regular BST + properties check to rebalance, if needed (RB-DELETE-FIXUP – check the textbook for the code)

P4 (black height) is an issue!!

```
rb_delete(T, z)  
  tree_delete(T, z)  
  if color[y]=black  
    then RB-DELETE-FIXUP(T, x)
```

RB-DELETE

- Del as in regular BST + properties check to rebalance, if needed (RB-DELETE-FIXUP – check the textbook for the code)

P4 (black height) is an issue!!

rb_delete(*T*, *z*)

tree_delete(*T*, *z*)

if color[*y*]=black

then RB-DELETE-FIXUP(*T*, *x*)

z=node to be removed (see picture on the blackboard)

y=node actually removed ($y \equiv z$ in case *z* has at most 1 child);
info in *y* is placed in *z*'s node

x=*y*'s only child before the delete process takes place (could be nil, in case *y* has no children). After *y* is deleted, *x* becomes the child of *y*'s parent (thus, *x*'s parent could have now both children, one being *x*)

w=*x*'s brother (after delete operation takes place; it's *y*'s brother before the deletion)


RB-DELETE-FIXUP

On x's branch check P4 property
x is y's (the removed node) only child

```
if color[x]=red  
    then color[x]<-black  
        //problem fixed; DONE!  
        //x brings its former father color  
    else color[x]<-double_black  
        //P1 property issue!
```

Computer Science

RB-DELETE-FIXUP - Case#1





- Issue at node **A** (pointed by x) which is **double black**!
- A = was the only child of the deleted node
- Parent (**B**) = **Black**, brother-w(**D**) = **Red**
- $\alpha, \beta, \gamma, \delta, \epsilon, \xi$ are RB trees
- $B \leftrightarrow D$ color interchange + left rotate \Rightarrow case 2 or 3 or 4 (w is black, differ w's children)

```
parent[x] <- red
brother[x] <- black
left_rotate(p[x])
1/18/2021
```

//colors interchanged

RB-DELETE-FIXUP - Case#2



Parent (**B**)=**Red** (if after case #1) or **Black**, brother-w (**D**)=**Black**,
both children of w (**C** and **E**) are **Black**

brother[x] ← red

if p[x]=red



then p[x] ← black

//when case#2 comes after case#1; problem solved

else p[x] ← double_black

x ← p[x] //the same problem as at the beginning of case #2, just 1 level
above; case 2 **repeats**; in lgn problem solved

RB-DELETE-FIXUP - Case#3






- Parent (**B**)=**Red** (if after case #1) or **Black**, brother-w (**D**)=**Black**, w's left child (**C**) = **Red**, w's right child (**E**) = **Black**
- $\alpha, \beta, \gamma, \delta, \epsilon, \xi$ are RB trees
- A=child of the deleted node (double Black A (pointed by x))
- C \leftrightarrow D color interchange +right rotate => **case 4**

`brother[x] <- red`

`left[brother[x]] <- black`

`right_rotate(brother[x])`

RB-DELETE-FIXUP - Case#4



- Parent(**B**)=**Red** (if after case #1) or **Black**, brother(**D**)=**Black**
- Node C is either **Red** or **Black**; node **E** is **Red**
- $\alpha, \beta, \gamma, \delta, \epsilon, \xi$ are RB trees
- A=child of the deleted node (double Black A (pointed by x))
- B \leftrightarrow D color interchange +left rotate, color E **Black** => probl. solved

```
brother[x] <- color[parent[x]]
```

```
parent[x] <- black
```

```
left_rotate(p[x])    //1 more black node on x's branch
```

```
color[right[w]] <- black
```

```
x <- root[T]
```


RB-del – Rebalancing eval

- Case #1 rotation followed by any other case
 - $1+2 \Rightarrow$ problem solved $O(1)$
 - $1+3+4 \Rightarrow$ problem solved $O(1)$
 - $1+4 \Rightarrow$ problem solved $O(1)$
- Case #2 (**no rotation**, only recoloring) repeats 1 level up in the tree
 - Worst case $O(\lg n)$
 - Best case $O(1)$
- Case #3 rotation followed by case #4 $O(1)$
- Case #4 rotation; solves the problem $O(1)$
- Delete $O(\lg n)$ + rebalancing
 - Worst case: #2 repeats (recoloring only) $O(\lg n)$
 - Best case: #4 \Rightarrow 1 rotation $O(1)$
 - Other cases: #1+2 or $1+3+4 \Rightarrow$ 2 or 3 rotations $O(1)$
- $O(\lg n)$ overall, at most 3 rotations

RB-DELETE-FIXUP - procedure

RB-DELETE-FIXUP (T, x)

while $x \neq \text{root}[T]$ and $\text{color}[x] = \text{black}$
do

if $x = \text{left}[p[x]]$

//cases on the left

then

//else case symmetric on the right; not discussed

$w \leftarrow \text{right}[p[x]]$

//w=x's brother

if $\text{color}[w] = \text{red}$

then $\text{color}[w] \leftarrow \text{black}$



//case #1 APPLY ; coloring+rotation

$\text{color}[p[x]] \leftarrow \text{red}$

$\text{left_rotate}(T, p[x])$

$w \leftarrow \text{right}[p[x]]$ //end case #1;

//another case comes





RB-DELETE-FIXUP – contd.

if color[left[w]] = black and color[right[w]] = black
then //case #2

color[w] ← red
x ← p[x]


else



RB-DELETE-FIXUP – contd.

```

else                                     //color[left[w]]≠ black or color[right[w]] ≠black
    if color[right[w]] = black
    then                                //case #3
        color[left[w]] ← black
        color[w] ← red
        right_rotate(T, w)
        w ← right[p[x]] //end case #3
    
```



RB-DELETE-FIXUP – contd.

```
color[w] <- color[p[x]]           //case #4
```



```
color[p[x]] <- black
```

```
color[right[w]] <- black
```

```
left_rotate(T, p[x])
```

```
x <- root[T]
```

else ... //x=right[p[x], all 4 cases symmetric to the right
color[x] <- black



Conclusions on balanced search trees

Tree	Height	Ins	Del
BST	$[\lg n, n]$	$O(h)$	$O(h)$
RBT	$[\lg n, 2\lg n]$	2 rot	3 rot
AVL	$[\lg n, 1.45\lg n]$	1 rot	$\lg n$ rot
PBT	$\lg n$	n rot	n rot

For RBT, at most $\lg n/2$ color updates needed

Disjoint Sets

- Collection of dynamic DS $S = \{S_1, \dots, S_k\}$
- $\exists n$ elements (objects) in all k sets ($n \geq k$)
- each set S_i is identified by its *representative* element, $x \in S_i$;
- Basic operations:
 - **MAKE-SET** (x)
Generates a new set, with a single element $\Rightarrow n$ sets initially, each object has its own set, and it is its own representative el.
 - **UNION** (x, y)
joins 2 disjoint sets, represented by x and y ; builds $S_x \cup S_y$ (and destroys S_x and S_y); the representative becomes any of the 2 representatives;
 - **FIND-SET** (x)
Returns a pointer to the representative element of the set containing element x .

Disjoint Sets – contd.

n = nb. of objects in the whole S

m = total nb. of operations (MAKE-SET, UNION, FIND-SET)

$m \geq n$ (as we have n MAKE-SET operations)

Utility/Applications:

- speeds up execution when we need to find/group items with similar features
- graphs (connected components; MST)
- many other

Disjoint Sets - implementation

- Linked List
- A set = a **linked list**
- representative = the first element (head) of the list
- An object in such a list contains
 - The element from the set;
 - The pointer to the next element in the list (LL)
 - Pointer to the representative (ex: blackboard)
- **MAKE-SET(x)** — builds a list with a single element $O(1)$
- **FIND-SET(x)** — returns the representative $O(1)$
- **UNION(x, y)** — adds x's list at the end of y's list;
 - representative = former y's representative
 - all x's elements have to update representative pointer (ex: blackboard)

Disjoint Sets – implementation – contd.

- Worst case: $O(m^2)$ for all operations
 - n MAKE-SET (1 for each element)
 - UNION
 - n times (to get to a single set)
 - $1 + 2 + 3 + \dots + n-1 = O(n^2)$ (show on the blackboard)
 - $n \sim m$ (actually $m > n$, yet n is linear in m)
 - On average, $O(m)$ for a call of UNION, m calls
 $\Rightarrow O(m^2)$

Computer Science

Disjoint Sets – implementation increase efficiency

- Update pointers for the shorter lists
- Keep as knowledge their length (similar to Order Statistic Trees)
- **Theorem:** For n objects in LL with weighted union, for m MAKE-SET, FIND-SET and UNION takes $O(m + n \lg n)$
- **Proof:** (check the textbook – identify an **informal** justification)

Forest of Disjoint Sets

- Set = **tree** with root; keep parent pointer
- 1 node = 1 element (=1 obj) from the set
- 1 tree = a set
- The root = *representative* el.
- Basic Implem. ~ to lists (no improvement)
 - MAKE-SET (x) build the tree with root only
 - FIND-SET (x) goes up and return the representative
 - UNION (x, y) Ex: (blackboard)

Forest of Disjoint Sets – Heuristics (to increase performance)

- UNION by **rank**

- Similar to weighted unify on lists
- The tree with less nodes will point to the tree with many nodes
- Info kept at root level = rank = max height of the tree
- $\text{rank} \cong \lg(\text{dim})$ (is an approximation, not an exact value; a guarantee that value is never exceeded)

- PATH **compression**

- Within the **Find-Set**, each node on the *search path* will update the parent node to the *representative* (instead of parent), and leave the **rank unchanged!**
- **Shrink does NOT change rank!** Why? Ex: blackboard

$\text{rank} \cong \lg(\text{dim})$ It is an approximation, ONLY

Forest of Disjoint Sets – Heuristics

- $\text{Rank}[x]$
 - = max height of the subtree rooted by x
 - = nb. of edges on the longest path from x to a leaf
 - $\text{rank}[\text{leaf}] = 0$
- Find-Set leave ranks unchanged
- $O(m * \text{Ack}(n)) \rightarrow O(m)$

Computer Science

Forest of Disjoint Sets - Implementation

MAKE-SET (x)

```
p[x] <- x  
rank[x] <- 0
```

UNION (x, y)

```
LINK (FIND-SET (x), FIND-SET (y))
```

LINK (x, y)

```
if rank [x] > rank [y]  
    then p[y] <- x  
    else p[x] <- y  
if rank [x] = rank [y]  
    then rank [y] = rank [y] + 1
```

Forest of Disjoint Sets - Implementation

FIND-SET (x)

```
if x != p[x]  
  then p[x] <- FIND-SET (p[x])  
return p[x]
```


Required Bibliography

- From the Bible – Chapter 13 (Red Black Trees), Chapter 21 (Data Structures for Disjoint Sets)