



Fundamental Algorithms

Lecture #5

Cluj-Napoca 30.10.2019
Rodica Potolea, CS, UTCN

Agenda

- **Hash Tables**
- **Trees**
 - **Binary and multiway**
 - **Representation**
 - **Basic operations**

Hash Tables

- Stores a dynamic set of data for fast access = data whose content varies (ex symbol table in a compiler)
- Frequent operation = search
- DS that maintains a set of items (identified by a key) subject to following operations:
 - insert (item): add item to set
 - delete (item): remove item from set
 - search (key): return item with key if it exists
- goal: $O(1)$ time per operation.

Hash Tables - direct access table

- Items stored in an array (hash table) indexed by key (identifier of item)

1	/
2	/
...	...
key1	item1
...	...
key2	item2
...	...

Limitations:

keys must be nonnegative integers

large key range = large storage space

Solution

Reduce universe U of all keys down to reasonable range for table

\Leftrightarrow project U onto a table of size m

Hash Tables


- m = table's size
- $n = |U|$ (universe's size, number of possible keys)
- $h: U \rightarrow \{0, 1, \dots, m-1\}$ mapping. Properties?
- h calculates key's location in the table
- $h(\text{key1}) = h(\text{key2})$
 - Is it possible?
 - Collision
 - What happens?
 - Deal with collision

Hash Tables – collision handling

- Chaining (linked lists) ~ find alternative solution
 - Hash functions
 - Universal hash
- Open addressing (in table) ~ find alternative place
 - Linear probing
 - Double hashing

Chaining (linked lists)

- Linked list of colliding elements in each slot of table



Picture taken from MIT OpenCourseWare, Introduction to Algorithms
<http://ocw.mit.edu6.006Introduction>

Chaining (linked lists)

- $h(k)$ =Dispersion value
- Search must go through the whole list
- Load factor: $\alpha = n/m$ =average number of keys per slot ($n=|K|$, K = key set, $m=|T|$, T =table)
- Expected performance of chaining assuming *simple uniform hashing* $O(1 + \alpha) \Rightarrow O(1)$ iff $\alpha=1$, i.e. $n=m$!
- Worst case: all keys in K to the same slot $\Rightarrow O(n)$

iff =
if and only

Hash functions

- **Division Method:** $h(k) = k \bmod m$
- k_1 and k_2 collide when $k_1 = k_2 \pmod{m}$ i.e. when m divides $|k_1 - k_2|$
- OK if keys are randomly distributed
- Not OK if they are on a *pattern* distribution
- **Good Practice:** to *avoid* common regularities in keys make m a *prime number* that is not close to a power of 2 or 10.
- Drawbacks:
 - find prime numbers (finding small/reasonable prime values is not a problem – just take them from tables),
 - division is slow

Hash functions - cont

- **Multiplication Method:**

$$h(k) = m\{kA\}$$
$$= m(kA - [kA]),$$

fractional part from kA

KA – integer part of KA

where $0 < A < 1$, $A = ct$.

- **Good practice:**

- considering w = the length of the word of the machine,
 - $m = 2^p$ for some int p so that m fits a single word.
 - Thus, $h(k)$ easy to calculate (check the textbook for justification).
- Knuth shown **$A = (\sqrt{5} - 1)/2$** it's good value
 - malicious keys \Rightarrow all keys in the SAME location! $\Rightarrow O(n)$
 \Rightarrow any possible hash function is vulnerable

Universal hash

- **Random** select the hash function at the execution time, from a set of functions

Note: again, randomness helps efficiency

- **Theorem:** if $n \leq m$, the average number of collisions per key < 1 if a class of Universal hash functions is considered
- Hw: Check the textbook for **a class of Universal hash functions**

Open addressing

- All keys kept in the table (no linked lists),
1 key/slot $\Rightarrow m \geq n$
- The hash function specifies the order of slots to try for a key, not just one slot
- Sequence to try for a key k :
 $\langle h(k,0), h(k,1), \dots, h(k,m-1) \rangle$
- The sequence should be a permutation of $\langle 0,1,\dots,m-1 \rangle$

/
item ₃
item ₁
...
item ₂
...

Probing Strategies

- **Linear probing**

$h(k, i) = (h'(k) + i) \bmod m$, where $h'(k)$ is ordinary hash function, $i = 0, 1, \dots, m-1$

- Drawback: *clustering* = consecutive group of filled slots grows \Rightarrow average search time grows

- **Double hashing**

$h(k, i) = (h'(k) + ih''(k)) \bmod m$, where $h'(k)$ and $h''(k)$ are ordinary hash functions

$h''(k)$ should be relatively prime to m

Open addressing - eval

- $\alpha < n/m < 1$, α load factor
- **Theorem**
average **un**successful search time is $1/(1 - \alpha)$
- **Theorem**
average successful search time is $1/\alpha \ln(1/(1 - \alpha))$

Trees

- Dynamic structures
 - Target
 - Faster (than on linked lists) retrieval of elements
 - Maintain good running time for other operations
 - Basic operations ($n = \# \text{nodes in } T$; $h = \text{height of } T$)
 - Traversal $O(n)$
 - (pre, in, post) order
 - Search $O(n)$ regular; $O(h)$ BST
 - Insert $O(h)$
 - Remove $O(n)$ regular; $O(h)$ BST
- $h \in [\lg n, n]$ Why? Best? Worst?
- Find: min, max, pred, succ in BST only $O(h)$

Trees (binary) - representation


- Dynamic linked structures
- Minimal data representation:
 - key, left, right
 - parent, info
 - other (like balance, size, ... depending on type)
- Empty tree = nil
- Types of nodes:
 - root (just one in a tree),
 - internal (non root, non leaf)
 - leaves (all nodes with both children nil)

Trees – representation - cont

- Multiway trees
 - A node has more than just 2 children (unspecified; unknown; variable)
 - Represented as:
 - a tree with just one child (linked list)
 - a binary tree:
 - left link= first child (proper tree link)
 - right link = brother (next child of the parent's node; right links form a singly linked list of brothers)
- Transformation?
- Ex? See blackboard

Binary Search Trees (BST)

- Binary Trees – if no order imposed on keys, NO improvement over lists! Why to have them?
- BST=BT with a **total order relation** defined on the key's set.
- $\forall x \in \text{Left}, \forall y \in \text{Right}, x \leq \text{Key} \leq y$
- Any subtree of a BST is a BST
- In general, the properties of a structure with recurrent definition are shared by the component structures (subtrees in our case)



BST traversal

- Preorder: **Key**, Left, Right
- Inorder: Left, **Key**, Right
- Postorder: Left, Right, **Key**

(pre)=> **7** 3 2 5 4 6 11


(in) => 2 3 4 5 6 **7** 11

(post)=> 2 4 6 5 3 11 **7**

Root boldface

Left underlined

Inorder: keys are in nondecreasing order



BST traversal - code

```
tree_walk(x, order) //x=root; order=in, pre,post
if x<>nil
    then
        if order= pre
            then write key[x]
            tree_walk(left[x],order)
        if order= in
            then write key[x]
            tree_walk(right[x],order)
        if order= post
            then write key[x]
```

Note: Just ONE write statement is executed (one color)
Look for the nonrecursive implementation!!!

BST traversal - eval

- $\text{order} \in \{\text{in}, \text{pre}, \text{post}\}$
- Only one of the 3 statements `write key[x]` is executed
- $O(n)$ (assuming constant time for the operation(s) performed at the level of each individual node – write in our case)

BST – search -recursive

```
r_tree_search(x, k)    //x=root; k=searched  
if x=nil or k=key[x]  
    then return(x)  
    else if k<key[x]  
        then r_tree_search(left[x], k)  
        else r_tree_search(right[x], k)
```

Running time: $O(h)$

In a BST, $h \in [\lg n, n]$

Discussion on recursive vs iterative implementation

Recursive implementation: where to place the conditional statement (if) & why

BST – search -iterative

```
i_tree_search(x, k)    //x=root; k=searched  
while x<>nil and k<>key[x]  
do  
    if k<key[x]  
        then x<-left[x]  
        else x<-right[x]  
return x
```

How does the time differ between iterative vs recursive implementation?

Same efficiency (big Oh), smaller machine time for iterative version (reason: overhead with stack)

BST – insert

Always as a leaf, regardless the particularity of the BT!!!!

NEVER EVER internal node. There is NO exception!

Running time: $O(h)$.

Range of h LARGE for regular trees

Rooted tree = tree as a DS, $\text{root}[T]$ its root

BST – insert - code


```
tree_insert(T,z)    //x=root; z=new node, already allocated
y<-nil              //y=x's parent; stays behind x;
x<-root[T]

while x<>nil          //search loop to find the position to insert
  do y<-x            // y=x at the prev step
  if key[z]<key[x]
    then x<-left[x]
    else x<-right[x]
p[z]<-y              //position found; x=nil; y=new node (z)'s parent
if y=nil            //in case the tree was empty before this insertion
  then root[T]<-z
  else if key[z]<key[y]
    then left[y]<-z
    else right[y]<-z
```


BST - delete

- Find the node
- Remove the node
- Cases:
 - Case 1: Leaf – remove it
 - Case 2: 1-successor node – skip it (its only child will become its parent's child)
 - Case 3: 2-successors nodes!
 - Chain the tree (fast, unbalances the tree)
 - Replace the operation with an easier one:
 - Keep the structure= keep the node,
 - place a different (appropriate) content,
 - remove of the node with the given content


BST – delete – ex: leaf




BST – delete – ex: a single successor node



BST – delete – ex: a single successor node



BST – delete – ex: root



Advantage: fast. How fast?

Drawback: increases the height

BST – delete - root

- Replace it with an easy-to-delete node
- Node to replace: *previous* or *next*. Why?
- $\text{Left} < \text{Key} < \text{Right}$
- $\text{Left}' < \text{prev}(\text{Key}) < \text{Key} < \text{next}(\text{Key}) < \text{Right}$
- Is it *prev/next* easier to del? Why?
- *prev*=max in Left= \Rightarrow no successor to the right
- *next*=min in Right= \Rightarrow no successor to the left
- \Rightarrow both are nodes with at most 1 child (easy to del)

BST – delete - code

```

tree_delete(T, z)           //z=node to delete; y physically deleted
if left[z]=nil or right[z]=nil
    then y<-z                //Case 1 OR 2; z has at most 1 child => del z
    else y<-tree_successor(z) //find replacement=min(right)
if left[y]<>nil                //we are in Case 2; y is a single child node
    then x<-left[y]           //y has no child to the right; x=y's child
    else x<-right[y]          //case 2 or 3. Why?
if x<>nil                     //y is not a leaf;
    then p[x]<-p[y]            // y's child redirected to y's parent = x's parent
    //becomes the former single (why?) grandparent
if p[y]=nil                  //means y were the root
    then root[T]<-x           //y's child becomes the new root
    else if y=left[p[y]]      //link y's parent to x which becomes its child
        then left[p[y]]<-x
        else right[p[y]]<-x
return[y]                    //outside the procedure: copy y's info into z; dealloc y

```


BST – delete - eval

- Find node to delete $O(h)$
- Find successor $O(h)$
- BUT:
 - if finding node to delete takes $O(h) \Rightarrow$ case 1 \Rightarrow leaf (no succ needed)
 - if node to delete not a leaf \Rightarrow case 2 or 3 \Rightarrow succ searched from that place down \Rightarrow find node + find succ $= O(h)$
- Delete takes $O(h)$ overall