

The background features a large, light blue watermark logo of the Technical University of Cluj-Napoca. The logo consists of a shield-like shape with the letters 'T' and 'U' integrated into its design. Above the shield, the words 'TECHNICAL UNIVERSITY' are written in a sans-serif font. Below the shield, the words 'OF CLUJ-NAPOCA' are written in a smaller sans-serif font. At the bottom of the logo, the words 'Computer Science' are written in a stylized, rounded font.

Fundamental Algorithms

Lecture #6

Cluj-Napoca

November, 06, 2019

Agenda

- **Trees**
 - **Basic operations**
 - walk, search, insert, delete– review
 - min, max, pred, succ
 - **Special types**
 - **Balanced trees**
 - PBT (seminar #3)
 - AVL
 - Red-Black (next lecture)
 - **Augmented Trees**
 - Order-statistic trees
 - Tree/lists

BST – walk, search, insert

- **Walk**

- pre/in/post-orders **$O(n)$** if $O(1)$ outside recursive calls
- else apply master theorem

- **Search**

- **$O(n)$** for BT
- **$O(h)$** for BST, $h \in [\lg n, n]$
- **$O(\lg n)$** for **balanced** BST

- **Insert**

- **Search** for it and reach a leaf/1-child node (parent for the new node)
- Insert as **leaf** always, as child of the given leaf/1-child node

BST - delete

- Remove the node
- Cases:
 - Leaf – remove it
 - 1-successor node – link parent with the only child
 - 2-successors nodes!
 - Chain the tree (fast, unbalances the tree)
 - Replace the node with an appropriate one (content of predecessor/successor), and remove (the location of) that one (same time, better balance)

BST – delete - eval

- Find node to delete $O(h)$
- Find successor/predecessor $O(h)$
- BUT:
 - if finding node to delete takes $O(h) \Rightarrow$ the node is a leaf \Rightarrow case 1 \Rightarrow no succ needed
 - if node to delete not a leaf, succ searched from that place down \Rightarrow find node+find succ= $O(h)$
- Delete takes only $O(h)$

Find-min/max $O(h)$

- Root's leftmost/rightmost leaf in the tree rooted at x;

find_tree_min(x) //x=root;

```
while left[x] <> nil  
do   x <- left[x]  
return x
```

Q: what if left[x]=nil?

find_tree_max(x) //x=root;


```
while right[x] <> nil  
do   x <- right[x]  
return x
```

Find-pred/succ

- $\text{pred} = \text{max in the left subtree} \Rightarrow$
find_tree_max(left[x])
- $\text{succ} = \text{min in the right subtree}$
find_tree_min(right[x])
- Any other situation possible?
 - What if the node has no left/right subtree?
Possible?
 - It has no pred/succ?
 - Not necessarily: counterexample!

Find-pred/succ- counterexample

- 6 has no right child.
- It means it has no successor?
 - False! 7 is its successor!
- 5 has no left/right child.
- It means it has no predecessor/succ?
 - False! 4 is its predecessor/6 its pred!
- How can we find pred/succ for such nodes?



(identify the property such nodes possess)

succ=lowest level ancestor whose left child is an ancestor as well

pred=lowest level ancestor whose right child is an ancestor as well

Determine (for succ) a triangle:

node-upwards while on a right child link

the first time the node is a left child= it is the succ node

Find-succ-code

find_tree_successor(x) //returns x's successor

if right[x] <> nil //regular case; the succ belongs to the same subtree
then return find_tree_min(right[x])

y <- p[x] //y keeps a pointer 1 level above x

while y <> nil and x = right[y]

// as long as we haven't reached the root and not changed the direction

// along the upwards path, go upwards 1 level

do x <- y


y <- p[y]

return y

Note: 2's successor is 4 (in find_tree_min)

6's successor is 7 (take **while twice** and change direction)

5's successor is 6 (0 while, exit while without going upwards at all)



Find-succ $O(h)$

- Cases:
 - *find_tree_min(right[x])*, worst case: $x = \text{root}$, succ lowest leaf $\Rightarrow O(h)$
 - x has no right child; worst case: $x = \text{leaf}$ on the lowest level, direction changes at the root level \Rightarrow succ root of the tree $\Rightarrow O(h)$
- *find_tree_successor* $O(h)$
- Find the predecessor is symmetric (change right with left and min with max) -

Homework

BST-eval

- Theorem: All operations in a BST (except traversal) take $O(h)$
- Adv: faster than on lists!
- Limitation: h ? Worst case $h=n$ (why?).
Therefore, no improvement at all!
- Enhancement?
 - **Balanced trees!**

Balanced trees

- Augmented BST to keep the height under control
- No matter the balance type, the height is proportional to $\lg n$ (**$c \cdot \lg n$** , with $c \geq 1$, but c a SMALL CONSTANT)
- The best possible balanced trees – PBT (perfect balanced trees) – seminar #3
- many other possibilities (for balance)

Balanced trees - PBT

- Perfect Balanced Trees
- Balance refers to nb of nodes, not to heights
- $b = n_R - n_L \in \{-1, 0, 1\}$
- $h = \lg n$
- Insert $O(n)$: as in regular BST $O(h) = O(\lg n)$
requires n rotations $O(n)$
- Delete $O(n)$: as from a regular BST $O(h) = O(\lg n)$
requires n rotations $O(n)$
- Best h property; difficult (costly) to maintain


Balanced trees - AVL

- On height (AVL=Adelson-Velskii, Landis)
- $b = h_R - h_L \in \{-1, 0, 1\}$
- PBTs are AVLs
- Most unbalanced out of AVL=Fibonacci trees (i.e. nb of left/right nodes specified by fib. numb.)
$$F_n = F_{n-1} + F_{n-2} + 1$$
 (b=-1 in every node)
- Insert $O(h)$: as in regular BST $O(h) = O(\lg n)$
Blackboard justification requires at most **1 rotation $O(1)$**
- Delete $O(h + \lg n)$: as from a regular BST $O(h) = O(\lg n)$
requires at most **$\lg n$ rotation $O(\lg n)$**
- $h \leq 1.45 \lg n \Rightarrow$ Good height property;
- easy to maintain for insertion; deletion might make many changes in the structure

AVL – rotations

- Preserve the search property
- Ensure the balance property
- Self-balancing:
 - Single rotation (blackboard)
 - Double rotation (blackboard)
 - Both take $O(1)$!
- After an insertion, at MOST 1 rotation may occur. Discussion.
- No other situation may occur. Why? Justification.
- After a rotation, the NEXT insertion along the same branch would NOT require a self-balancing (rotation)
- The same rotations are used for Red-Black trees!

BST-balanced trees relationship



Augmented DS


- Augmented = additional property and/or behavior to help (i.e. speed up) various tasks preserving ALL existing properties and behavior with (at least) the SAME performance
- Balanced BST are augmented trees (objective, keep the height under control)
- Current objective = better (=faster) select operations on BST
- **Order Statistic (OS) Tree**
- Augmentation= store at the node level as additional information the dimension of the tree (i.e. the number of nodes in the tree rooted by the given node)
- $\text{dim}[x] = \text{dim}[\text{left}[x]] + \text{dim}[\text{right}[x]] + 1$
- How is calculated? (if the information is not already stored?) – postorder.

Augmented DS – contd.

- How to maintain this information for the basic tasks (search, insert, delete, traversal, update)?
- What operations are improved?
- Other tasks: Selection and Ranking
 - Selection (i^{th} selection) = find the node which is the i^{th} one in inorder traversal
 - Selection
 - in arrays – ordered? Not ordered?
 - in lists – ordered.
 - in trees
 - Can we do better for BST?

Selection

- Returns the i^{th} smallest key in the tree
 - rank given (i)
 - key returned (pointer to the i^{th} smallest key in the tree)
- Input: rank (i.e. index in inorder),
- Output: node at the given rank
- Augmentation: dimension =
=nb of nodes rooted by the node.
- $\text{dim}[x] = \text{dim}[\text{left}[x]] + \text{dim}[\text{right}[x]] + 1$
- $\text{dim}[\text{nil}] = 0$



OS Select $O(h)$

Initial call with $\text{root}(T)$ and Returns pointer to the i^{th} key

Resembles Hoare's selection on unordered arrays (partition is missing, as we have a BST; so, just the recursive call)

OS_Select(x, i)

$r \leftarrow \text{dim}[\text{left}[x]] + 1$ //number of nodes on the left + root

if $i = r$ //found it

then return x

else if $i < r$ //ith smallest is on the left

then

return $\text{tree_select}(\text{left}[x], i)$

else //ith smallest is on the right

return $\text{tree_select}(\text{right}[x], i - r)$


Ranking

- Reverse problem:
 - key given
 - rank returned
- Input:
 - given an existing key from the tree (that is, a pointer to the node containing that key)
- Output:
 - Return its rank in the tree (i.e. its position in the inorder walk)
 - Rank = nb of keys smaller than the checked key in the tree. Approach: count them all (all before = all to left)

Ranking – contd.

Case #1 node is a right child of its parent (Ex: rank Key2)

$$\text{rank}(\text{Key2}) = \text{dim}(\text{RL}) + 1 + \text{dim}(\text{L}) + 1$$




While going upwards in the tree, evaluate what type of child the current node is:

- if a right child (case #1)

Count the nb of nodes in any subtree to the left of the branch starting from the current node (x) up to the root (T)

Ranking – contd.

Case #2 node is a left child of its parent (Ex: rank Key1)
 $\text{rank}(\text{Key1}) = \text{dim}(\text{LL}) + 1$



While going upwards in the tree, evaluate what type of the child the current node is:

- if a left child (case #2)

Count the nb of nodes in any subtree to the left of the branch starting from the current node (x) up to the root (T)

OS Rank $O(h)$

OS_Rank (*T*, *x*)

```
r<-dim[left[x]]+1
```

```
y<-x
```

```
while y<>root[T]
```

```
do
```

```
  if y=right[p[y]]
```

```
  then
```

//case #1

```
    r<-r+ dim[left[p[y]]]+1
```

```
    y<-p[y]
```

//case #2

```
return r
```


Augmented trees (by dimension)

- Evaluation (performance for select and rank)
- Worst case $O(h)$
- For balanced trees $h = \lg n \Rightarrow O(\lg n)$
- OS trees are Red-Black Trees (RBT – check lecture #7)
- What happens (what changes in the OStree, besides the regular info/tasks specific to RBT) when updates occur
 - Insert? Discussion/Analysis
 - Delete? Discussion/Analysis

Augmented trees (type 2)

- Requirements:
 - Regular operations are performed as (**same performance also**) in BST (walk ($O(n)$), search, ins, del ($O(h)$))
 - Several other operations are enhanced (i.e. performed faster)
 - Succ
 - Pred
 - Min
 - Max
 - **All required to be performed in $O(1)$!!!**
 - BUT NONE of the before-defined operations should degrade their performance

Augmented trees – contd.

- Info in a node:
 - Usual info:
 - key
 - left pointer
 - right pointer
 - parent pointer
 - Supplementary info (see picture on the blackboard):
 - succ pointer
 - pred pointer (together ensure walking through the list)
 - pp ensures min/max oper. (in regular BST, succ/pred calculated **either** based on min/max **or** pp which is determined at the execution time)

Augmented trees – contd.

- The structure acts BOTH as a BST and DLL!!
(check the blackboard for an example)
- Regular operations are:
 - done like in any other BST
 - in addition, need to make some updates
- They (the additional updates) refer to:
 - making the appropriate links within the DLL
(set/update the pred and succ pointers)
 - link the double pointer (set/update the pp
pointer)

Augmented trees - Insert

- Regular **insert** operation in a BST (x inserted) **+**

if $x = \text{right}[p[x]]$ //node inserted = right child

then //case #1

$pp[x] \leftarrow \text{succ}[p[x]]$

$\text{dl_list_ins_after}(p[x], x)$

else //case #2; node inserted = left child

$pp[x] \leftarrow \text{pred}[p[x]]$

$\text{dl_list_ins_after}(pp[x], x)$

Augmented trees – Delete

(z = node requested to be removed; it's content is replaced by y's content
y=node actually removed = at most 1 child node;
x = its only child/if any, might be nil;
z=y if z has at most one child)

- Apply regular **delete** operation in a BST + code below

if right[y]=nil

then

x<-left[y]

while x<>nil

do pp[x]<-pp[y] //update pp

x<-right[x]

dl_list_del(y)

else

//symmetric on the left

Augmented trees – Min

- **min** (based on *succ* and *pp* as opposed to regular BST where *succ* is calculated based on *min* **or** determined *pp*)

if $x = \text{left}[p[x]]$

then

return $\text{succ}[pp[x]]$

//on the leftmost branch, **HAS TO BE** $pp[x] = \text{nil}!!!$

else

return $\text{succ}[p[x]]$

Augmented trees – Max

- **max** (based on *pred* and *pp* as opposed to regular BST where *pred* is calculated based on *max* or determined *pp*)

if $x = \text{left}[p[x]]$

then

return $\text{pred}[p[x]]$

else

return $\text{pred}[pp[x]]$

//on the rightmost branch, **HAS TO BE** $pp[x] = \text{nil}$ and
 $\text{pred}[\text{nil}[T]] = \text{last node in inorder} = \text{last node in the list}$

Augmented trees – contd.

- Particular (initial) cases discussion on the blackboard!
- First **insert** (in the empty tree)

Tree_ins(*T*, *x*)

if *x* = root[*T*]

then //the node just inserted is the root = tree is empty

p[*x*] <- nil[*T*]

pp[*x*] <- nil[*T*]

dl_list_ins_after(*pp*[*x*], *x*)

else //the regular case described earlier (needs another if-then-else;
 //here only left case)

pp[*x*] <- *pred*[*p*[*x*]]

dl_list_ins_after(*pp*[*x*], *x*)

Homework: updates for delete!