

The background features a large, light blue watermark logo of the Technical University of Cluj-Napoca. The logo consists of a shield with a stylized 'T' and 'U' inside, with the text 'TECHNICAL UNIVERSITY' at the top, 'OF CLUJ-NAPOCA' in the middle, and 'Computer Science' at the bottom.

Fundamental Algorithms

Lecture #7&8

Cluj-Napoca

November 20, 2019

Agenda Lecture 7

- **MT**
- **Augmented Trees (type 2) – check Lecture 6 for notes**

Agenda Lecture 8

- **Red-Black Trees**
 - **Insert**
 - **Delete**
- **Balanced Trees**
 - **Comparative analysis and conclusions**
- **Disjoint Sets**
 - **Representation**
 - **Basic Operations**
 - **Implementation**

Red-Black trees

- Balanced trees
- Both insert/delete operations take $O(\lg n)$, with constant time for rebalancing

Def: is a BST with the following properties:

P_0 the root is **black**

P_1 each node is **colored** either black or red

P_2 each **leaf** (=nil!) is **black**

P_3 both **children of** a **red** node are **black**

P_4 every **path** from any node to a leaf has the **same number of black nodes**

Red-Black trees

Th: A RB tree with n **internal** nodes (without leaves = nil nodes) has its **height at most $2\lg(n+1)$**

Proof: notation: $bh(x)$ = the black height (without x) of node x

Step 1: Define the statement $P(bh)$ as follow:

$P(bh)$: $\forall x \in RBT$, the tree rooted by x has at least $2^{bh(x)} - 1$ nodes

Induction:

$$P(0) \quad 2^0 - 1 = 1$$

Assume $P(bh)$ true $\Rightarrow P(bh+1)$ true?

x has 2 children; each has the black height $bh(x)$ (if x is red)
OR $bh(x)-1$ (if x is black)

nb of internal nodes of x = nb of internal nodes of children(x) + 1

(itself) \Rightarrow at least $(2^{bh(x)-1} - 1) + (2^{bh(x)-1} - 1) + 1 = 2^{bh(x)} - 1$ qed (end of step1)

Red-Black trees

Step 2:

We know $P(bh)$ is true, ie

$P(bh)$: $\forall x \in RBT$, the tree rooted by x has at least $2^{bh(x)} - 1$ nodes (1)

By P_3 of RBT def (use contradiction to prove) $bh(x) \geq h/2$ (2)

//since after each red node comes a black one

$$\Rightarrow n \geq 2^{bh(x)} - 1 \quad (\text{from (1)})$$


$$\geq 2^{h/2} - 1 \quad (\text{from (2)})$$

$$n \geq 2^{h/2} - 1 \Leftrightarrow n+1 \geq 2^{h/2} \Leftrightarrow h/2 \leq \lg(n+1) \Leftrightarrow h \leq 2\lg(n+1)$$

(qed, end of Th proof)

Red-Black trees - rotations

Similar to single rotations right/left from AVL
They are symmetric



Picture from wiki

Red-Black trees - rotations

left_rotate(T, x)

//x root of rotation (points on P)

y←right[x] //y saves Q

right[x]←left[y] //right of P goes on B

if left[y]≠nil //if B exists = is not nil

then p[left[y]]←x //B's parent becomes P

p[y]←p[x] //Q's parent what was P's parent

if p[y]=nil //P used to be the root of the tree

then root[T]←y

else if x=left[p[x]] // the parent of P becomes the parent of Q


then left[p[x]]←y

else right[p[x]]←y

left[y]←x //P goes the left child of Q

p[x]←y //Q becomes the parent of P

12/3/2019





RB-insert

- **Insert** like in ANY other BST
 - As a LEAF, as for any other BST
- **Assign** it a color
 - RED (WHY?)
- Check the properties
- Re-balance if needed (**RB-INSERT-FIXUP** – check the textbook for the complete code)
- P_3 both children of a red node are black
 - True for the children (nils) of the inserted node
 - Not always true; in case the parent of the inserted node is RED colored \Rightarrow red (parent)-red (inserted)
 - Cases to analyze and remove inconsistencies

RB-insert- Case#1


- **B inserted node (pointed by x)**
- Parent (A)=RED, uncle (D)=Red, grandparent (C)=BLACK
- $\alpha, \beta, \gamma, \delta, \varepsilon$ are RB trees (β, γ empty at first)
- Swap colors between grandparent (C) and parent (A)/uncle (D)

parent<-black //no more P3 conflict A-B
uncle<-black //to preserve P4
grandparent<-red //to preserve P4
x<-grandparent



RB-insert- Case#1-eval

- **P_3 may still be invalid**, for the new x (i.e. **C**)
- Problem transferred **2 levels up** in the tree (now β, γ not empty any longer)
- It takes (in the worst case = all way up to the root) $O(h)$ to rebalance ($2\lg(n+1)/2 = \lg n$)





RB-insert- Case#2

- **B inserted node (pointed by x)**
- Parent(**A**)=RED, uncle (δ 's root)=BLACK (here is the difference compared to case #1), grandparent (C)=BLACK
- $\alpha, \beta, \gamma, \delta$ are RB trees; δ 's root is BLACK


left_rotate(p(x))
 $x \leftarrow \text{left}(x)$

//no more P3 conflict B-parent conflict



RB-insert- Case#2-eval

- Case #2 takes just $O(1)$ to apply, but
- **P_3 is still invalid**, for the new x denoted node (i.e. A-B conflict)
- \Rightarrow it is followed by case #3





RB-insert- Case#3

- **Inserted A** /coming from #2 (node pointed by x)
- Parent (**B**)=RED, uncle (γ 's root)= BLACK, grandparent (**C**)=BLACK
- $\alpha, \beta, \gamma, \delta$ are RB trees

parent<-black

grandparent<-red

right_rotate(p[p[x]])



RB-insert- Case#3-eval

- Problem solved
- Each individual case takes $O(1)$
- Case #1 may repeat (up in the tree)
- Case #2 is followed by #3
- Case #3 solves the problem

RB-insert – Rebalancing eval

- Case #1 repeats up to the root $O(h)$
- Case #2+#3 \Rightarrow problem fixed $O(1)$
- Case #3 \Rightarrow problem fixed $O(1)$
- Insert $O(\lg n)$ + rebalancing
 - Worst case: #1 repeats $O(\lg n)$
 - Best case: #3 \Rightarrow 1 rotation $O(1)$
 - Other case: #2+3 \Rightarrow 2 rotations $O(1)$
- **$O(\lg n)$ overall worst time** (case 1 repeats), **at most 2 rotations** (case 2)

RB-delete

- Del as in regular BST + properties check to rebalance, if needed (RB-DELETE-FIXUP – check the textbook for the code)

- P4 (black height) is an issue

rb_delete(T, z)

tree_delete(T, z)

if color[y]=black

then rb_del_fix(T, x)

//else NO issue at all

z=node containing the key to be removed (see picture on the blackboard)

***y=node actually removed ($y \equiv z$ in case z has at most 1 child);
y's info is placed in z 's node***

x=y's only child before the delete process takes place (could be nil, in case y has no children). After y is deleted, x becomes the child of y 's parent (thus, x 's parent could have now both children, one being x)



w=x's new brother (after delete operation takes place; it's y 's brother before the deletion)

RB-delete

On x's branch check P4 property
x is y's (= the removed node) only child

```
if color[x]=red  
    then color[x]<-black  
        //problem fixed; DONE!  
        //x brings its former father color  
else color[x]<-double_black  
        //does not exist something like this  
        //brings one additional black  
        //to keep the black height property (P4)  
        //yet, a P1 property issue!
```


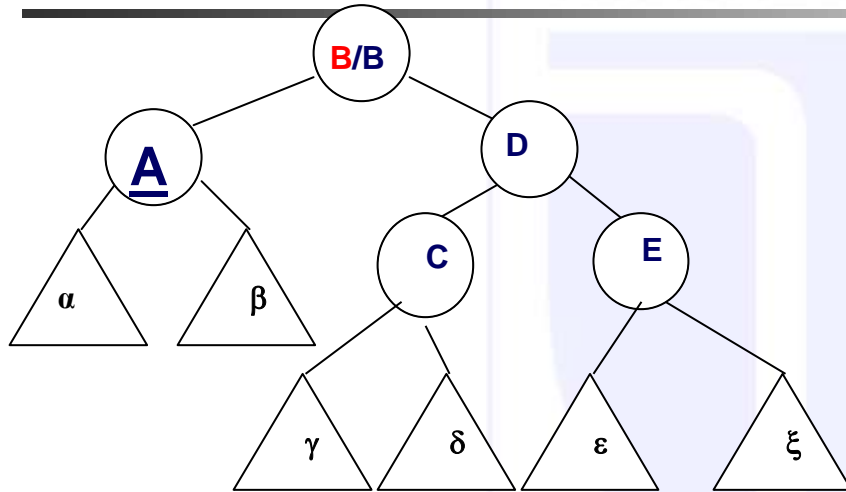
RB-delete- Case#1



- Issue at node **A** (pointed by x) which is **double black**!
- A= was the only child of the deleted node
- **Parent (B) = Black, brother (D) = Red**
- $\alpha, \beta, \gamma, \delta, \epsilon, \xi$ are RB trees
- B \leftrightarrow D color interchange +left rotate=>case 2 or 3 or 4 follows

```
parent[x] <- red
brother[x] <- black           //colors interchanged
left_rotate(p[x])
```

RB-delete- Case#2



**Parent (B)=Red (if after case #1) or Black (stand alone case #2),
brother (D)=Black, node C = Black**

brother[x] ← red

if p[x]=red



then p[x] ← black

//when case#2 comes after case#1; problem solved

else p[x] ← double_black



**x ← p[x] //the same problem as at the beginning of case #2, just
1level above; case 2 **repeats**; in lgn problem solved**

RB-delete- Case#3



- Parent (B)=Red (if after case #1) or Black (standalone case), brother (D)=Black, node C = Red
- $\alpha, \beta, \gamma, \delta, \epsilon, \xi$ are RB trees
- A=child of the deleted node (double Black A (pointed by x))
- $C \leftrightarrow D$ color interchange +right rotate \Rightarrow **case 4**
`brother[x] ← red`
`left[brother[x]] ← black`
`right_rotate(brother[x])`

RB-delete- Case#4



- Parent=Red (if after case #1) or Black, brother=Black, Node C is either Red or Black; node E = Red
- $\alpha, \beta, \gamma, \delta, \epsilon, \xi$ are RB trees
- A=child of the deleted node (double Black A (pointed by x))
- B \leftrightarrow D color interchange +left rotate => problem solved
`brother[x] <- color[parent[x]]`
`parent[x] <- black`
`left_rotate(p[x])` //1 more black node on x's branch

RB-del – Rebalancing eval

- Case #1 rotation followed by any other case
 - $1+2 \Rightarrow$ problem solved $O(1)$
 - $1+3+4 \Rightarrow$ problem solved $O(1)$
 - $1+4 \Rightarrow$ problem solved $O(1)$
- Case #2 **no rotation**, only recoloring - repeats 1 level up in the tree
 - Worst case $O(\lg n)$
 - Best case $O(1)$
- Case #3 rotation followed by case #4 $O(1)$
- Case #4 rotation; solves the problem $O(1)$
- Delete $O(\lg n)$ + rebalancing
 - Worst case: #2 repeats (recoloring only) $O(\lg n)$
 - Best case: #4 \Rightarrow 1 rotation $O(1)$
 - Other cases: #1+2 or 1+ 3+4 \Rightarrow 2 or 3 rotations $O(1)$
- **$O(\lg n)$ worst overall, at most 3 rotations**

RB-del - procedure

rb_del_fix(T, x)

while $x \neq \text{root}[T]$ ***and*** $\text{color}[x] = \text{black}$
do

if $x = \text{left}[p[x]]$

then

$w \leftarrow \text{right}[p[x]]$

if $\text{color}[w] = \text{red}$

then

$\text{color}[w] \leftarrow \text{black}$

$\text{color}[p[x]] \leftarrow \text{red}$

left_rotate($T, p[x]$)

$w \leftarrow \text{right}[p[x]]$ //end case #1;



//another case comes

//cases on the left

//else case symmetric on the right; not discussed

//w=x's brother

//case #1 APPLY ; coloring+rotation





RB-del - procedure

if color[left[w]]=black and color[right[w]]=black
then //case #2

color[w]←red
x←p[x]



else



RB-del – procedure - cont

```

else           //color[left[w]]≠ black or color[right[w]] ≠black
  if color[right[w]] = black
  then         //case #3
    color[left[w]] ← black
    color[w] ← red
    right_rotate(T, w)
    w ← right[p[x]] //end case #3
  
```



RB-del – procedure - cont

```
color[w] ← color[p[x]]
```

//case #4



```
color[p[x]] ← black
```

```
color[right[w]] ← black
```

```
left_rotate(T, p[x])
```

```
x ← root[T]
```

else ... //x=right[p[x], all 4 cases symmetric to the right
color[x] ← black



Conclusions on balanced search trees

Tree	Height	Ins	Del
BST	$[\lg n, n]$	$O(h)$	$O(h)$
RBT	$[\lg n, 2\lg n]$	2 rot	3 rot
AVL	$[\lg n, 1.45\lg n]$	1 rot	$\lg n$ rot
PBT	$\lg n$	n rot	n rot

For RBT, at most $\lg n/2$ color updates needed

Disjoint Sets

- Collection of dynamic DS $S = \{S_1, \dots, S_k\}$
- \exists n elements (objects) in all k sets ($n \geq k$)
- each set S_i is identified by its *representative* element, $x \in S_i$;
- Basic operations:
 - **Build-Set** (x)
Generates a new set, with a single element \Rightarrow n sets initially, each object has its own set, and it is its own representative el.
 - **Unify** (x, y)
joins 2 disjoint sets, represented by $x \in$ and y ; builds $S_x \cup S_y$ (and destroys S_x and S_y); the representative becomes any of the 2 representatives;
 - **Find-Set** (x)
Returns a pointer to the representative element of the set containing element x .

Disjoint Sets – contd.

n = nb. of objects in the whole S

m = total nb. of operations (Build-Set, Unify, Find-Set)

$m \geq n$ (as we have n Build-Set operations)

Utility/Applications:

- speeds up execution when we need to find/group items with similar features
- graphs (connected components; MST)
- many other

Disjoint Sets - implementation

- LL
- A set = a **linked list**
- representative= the first element (head) of the list
- An object in such a list contains
 - The element from the set;
 - The pointer to the next element in the list (LL)
 - Pointer to the representative (ex: blackboard)
- **Build-Set** (builds a list with a single element) $O(1)$
- **Find-Set** (returns the representative) $O(1)$
- **Unify** (x, y) — adds x 's list at the end of y 's list;
 - representative = former y 's representative
 - all x 's elements have to update representative pointer (ex: blackboard)

Disjoint Sets – implementation – contd.

- Worst case: $O(m^2)$ for all operations
- n Build-Set (1 for each element)
- Unify
 - n times (to get to a single set)
 - $1 + 2 + 3 + \dots + n = O(n^2)$ (show on the blackboard)
- $n \sim m$ (actually $m > n$, yet n is linear in m)
- On average, $O(m)$ for a call of Unify, m calls, $O(m^2)$

Disjoint Sets – implementation increase efficiency

- Update pointers for the shorter lists
- Keep as knowledge their length (similar to Order Statistic Trees)
- **Theorem:** For n objects in LL with weighted unify, for m Build-Set, Find-Set and Unify takes $O(m + n \lg n)$
- **Proof:** (check the textbook – identify an **informal** justification)

Forest of Disjoint Sets

- Set = **tree** with root; keep parent pointer
- 1 node = 1 element (=1 obj) from the set
- 1 tree = a set
- The root = *representative* el.
- Basic Implem. \sim to lists (no improvement)
 - Build-Set (x) build the tree with root only
 - Find-Set (x) goes up and return the representative
 - Unify (x, y) Ex: (blackboard)

Forest of Disjoint Sets – Heuristics

(to increase performance)

- Unify based on **rank**
 - Similar to weighted unify on lists
 - The tree with less nodes will point to the tree with many nodes
 - Info kept at root level = rank = max height of the tree
 - $\text{rank} \cong \lg(\text{dim})$ (is an approximation, not an exact value; a guarantee that value is never exceeded)
 - Tree shrink
 - Within the **Find-Set**, each node on the *search path* will update the parent node to the *representative* (instead of parent), and leave the **rank unchanged!**
 - **Shrink does NOT change rank!** Why? Ex: blackboard
- $\text{rank} \cong \lg(\text{dim})$ It is an approximation, ONLY

Forest of Disjoint Sets – Heuristics

- $\text{Rank}[x]$
 - = max height of the subtree rooted by x
 - = nb. of edges on the longest path from x to a leaf
 - $\text{rank}[\text{leaf}] = 0$
- Find-Set leave ranks unchanged

Forest of Disjoint Sets - Implementation

Build-Set (x)

```
p[x] <- x  
rank[x] <- 0
```

Reunion (x, y)

```
Unify (Find-Set(x), Find-Set(y))
```

Unify (x, y)

```
if rank [x] > rank [y]  
  then p[y] <- x  
  else p[x] <- y  
if rank [x] = rank [y]  
  then rank [y] = rank [y] + 1
```