

Chapter 3

User Application Interaction With an OS

Dealing With C Programs in Linux

Print Version of Lectures Notes of *Operating Systems*

Technical University of Cluj-Napoca (UTCN)
Computer Science Department

Adrian Colan and Ciprian Opris

March 11th, 2020

3.1

Purpose and Contents

The purpose of today's lecture

- Review few basic aspects regarding writing and running C programs
 - their relationship with the OS

3.2

Contents

1	Getting Executable from Source Code	1
2	Running the Executable	5
3	C Programs Debugging	12
4	Recommendations About Writing C Programs	14
5	Conclusions	16

3.3

1 Getting Executable from Source Code

Compiler Role

- CPU does not understand “C language”, not even “assembly”
 - **understands machine instructions** (its own language), i.e. bytes encoding certain actions
- **a program** (i.e. user application) that could be run should be
 - **a sequence of bytes** organized as **machine instructions**
 - machine instructions map 1:1 to assembly instructions

```
8b 10          mov     edx, DWORD PTR [eax]
89 d0          mov     eax, edx
c1 e0 02       shl     eax, 0x2
01 d0          add     eax, edx
c1 e0 02       shl     eax, 0x2
01 c8          add     eax, ecx
89 85 30 ff ff mov     DWORD PTR [ebp-0xd0], eax
```

- the **app developer** “knows” a **higher-level language** (e.g. C)
- **compiler** “**translates**” the **source code** into machine instructions
 - ⇒ binary executable

3.4

Calling the Compiler

- **gcc (Linux, Windows, Mac)** `gcc [opt] <source_name> -o <exec_name>`
- Visual C (Windows) `cl.exe [opt] <source_name> /link /OUT:<exec_name>`
- from an Integrated Development Environment (IDE)
 - interact with an interface, e.g. just click some button
 - ⇒ calls transparently commands like that above
- example
 - `gcc -Wall -Werror hellow.c -o hellow`
 - * `-Wall` option means *warnings=all*, i.e. displays all warnings
 - * `-Werror` option means *warnings=errors*, i.e. report warnings as errors

3.5

Never ignore compiler’s warnings!

3.6

Multiple Source File Applications

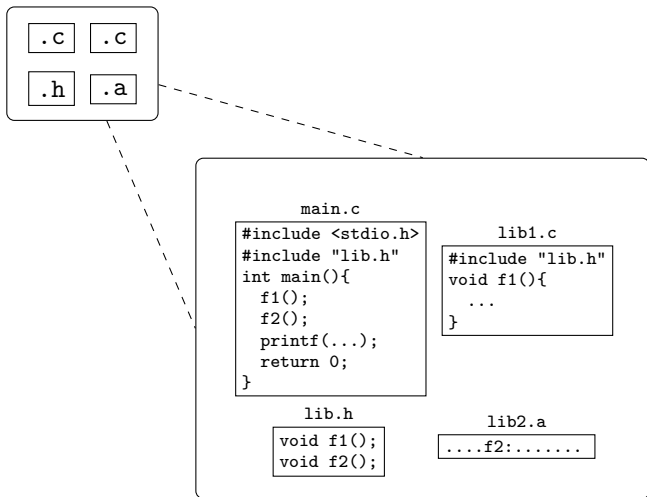
- source (text) files of a C-based application could be
 - **.c files**: code, **implementation (definition)**
 - **.h files**: type, function, constant **declaration**¹
 - * to be included (as text!) in other files
 - take care to not include the same .h file multiple times²!
 - take care to cyclic inclusions!
 - * usually **not for** variable and function **definition**!
- compilation process consists in more phases
 1. **pre-compilation**: expand / replace the preprocessor directives like `#include`, `#define`, ...
 2. **compilation**: compile each .c file ⇒ **object file .o**
 3. **linking**: object files linked together into a single executable
 - solve references to functions in other object files
 - solve references to functions in dynamic / **shared libraries** (.so files) or **static libraries** (.a files)

3.7

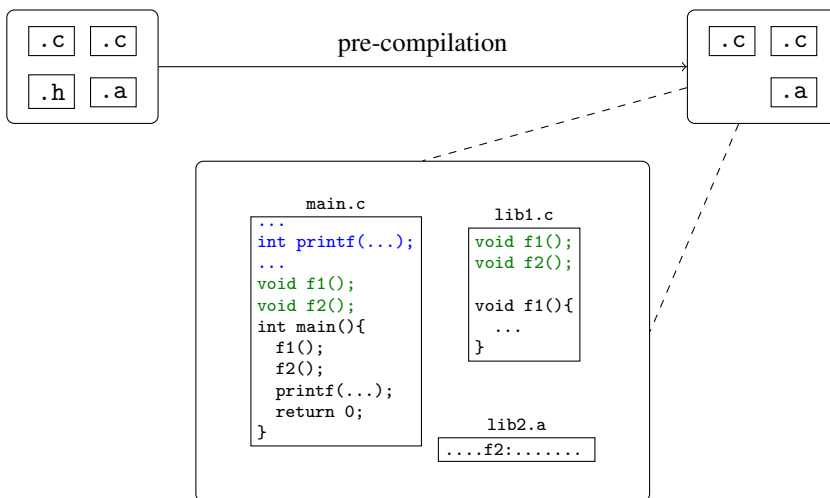
Compilation

¹Note the [difference](#) between definition and declaration!

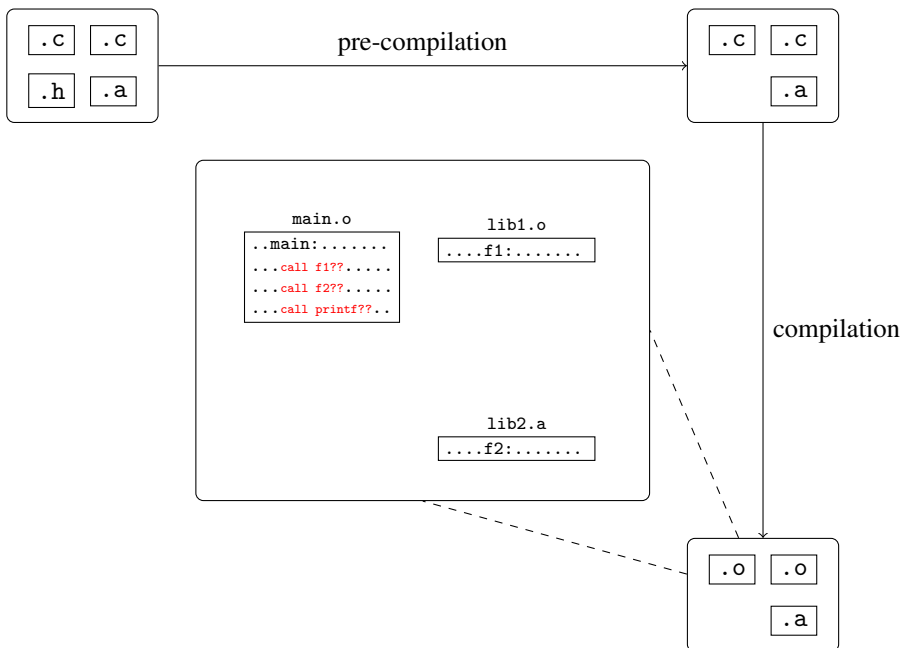
²See [gcc](#) and [Visual Studio](#) recommendations



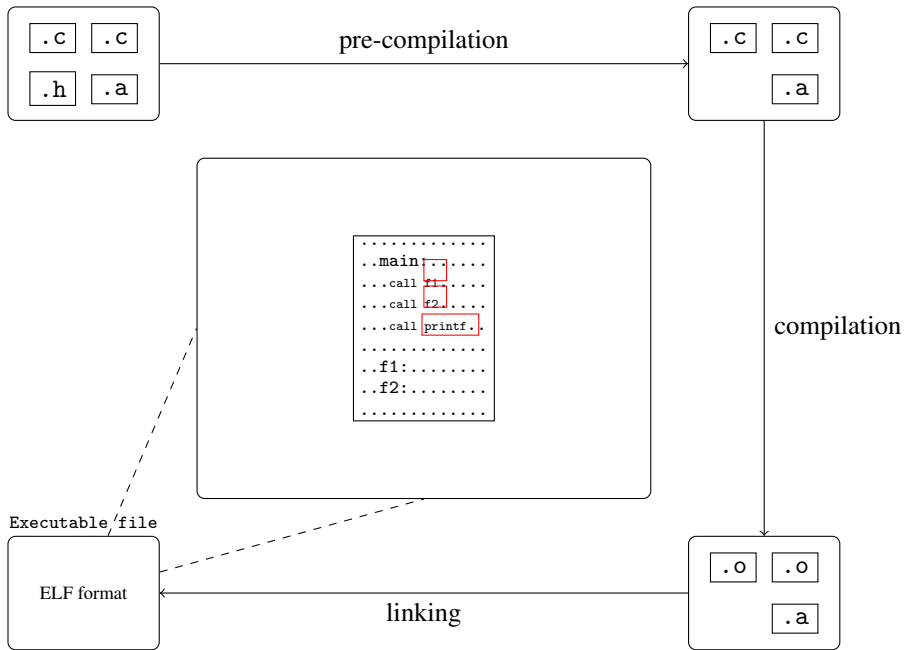
3.8



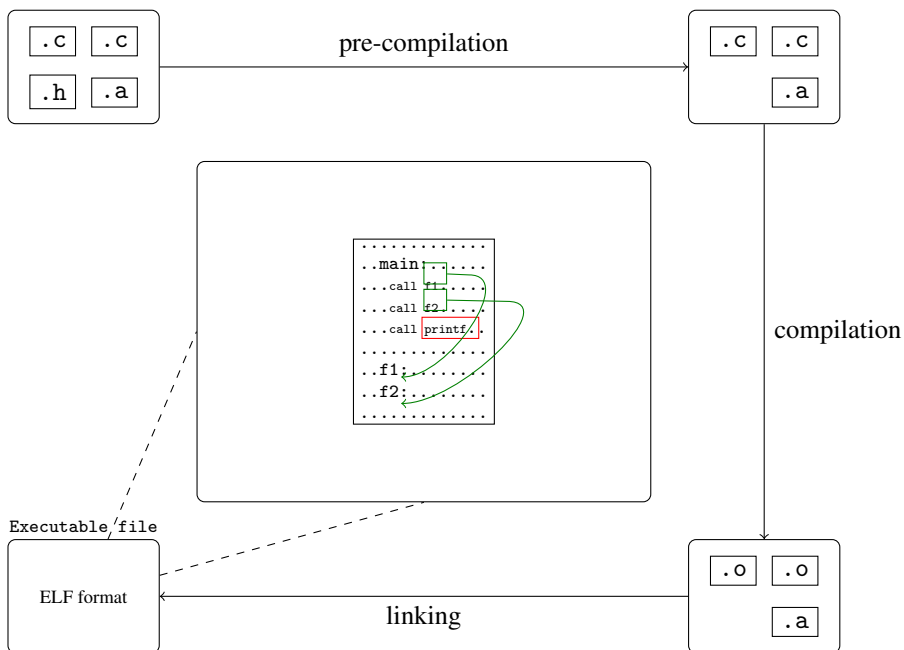
3.9



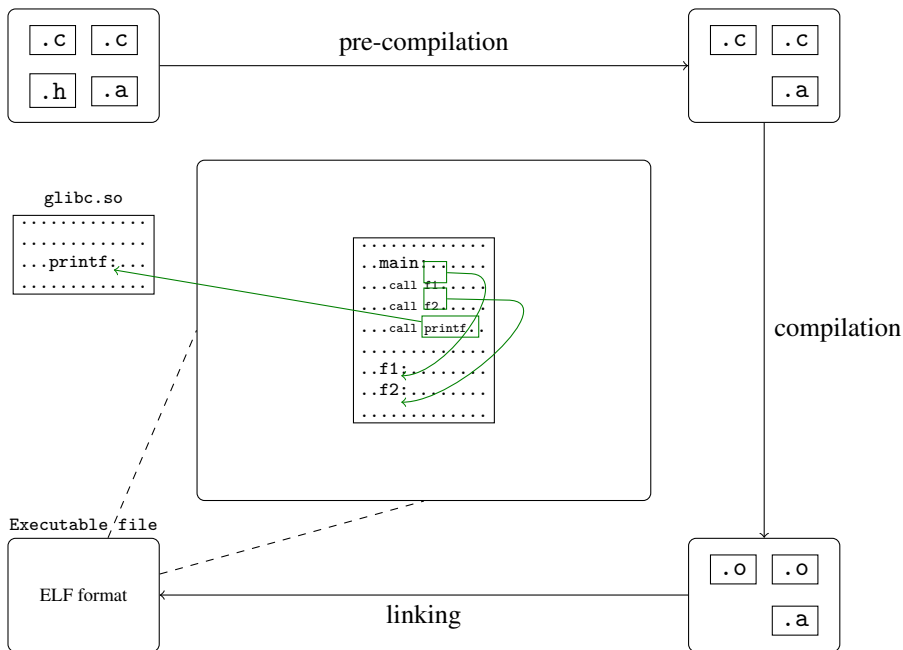
3.10



3.11



3.12



3.13

2 Running the Executable

From Executable File to the User Application

- the OS **allocates memory** for the new application
- the OS **loads** code and data from the **executable file** (ELF) into the allocated memory
 - ELF specification says what to load from the executable file
 - ELF specification says how many memory is needed
 - ELF specification says where to load into memory
- ⇒ **application's (virtual) address space**
 - complying a specific structure
 - different areas (segments): code, data, heap, stack etc.
 - there are also invalid areas (holes)
- **configure the CPU registers** with values corresponding to the new application's memory
- ⇒ CPU starts running the new application

3.14

Local variables and the stack (1)

test1.c:

```
#include <stdio.h>

int main(void){
    int x = 7;
    printf("x=%d\n", x);
    return 0;
}
```

Compiling: `gcc -Wall test1.c -o test1` Running: `./test1 x=7` Compiling with different

options: `gcc -g -m32 -Wall test1.c -o test1`

- `-g`: add debugging info in the executable
- `-m32`: generate 32-bit code

View executable as assembly code: `objdump -D test1 -M intel -S`

- `-D`: disassembly
- `-M intel`: Intel syntax
- `-S`: display source code also


3.15

Local variables and the stack (2)

```

;int main(void){
...
push ebp
mov  ebp,esp
...
push ebx
push ecx
sub  esp,0x10
...
;int x = 7;
mov  DWORD PTR [ebp-0xc],0x7
;printf("x=%d\n", x);
...
push DWORD PTR [ebp-0xc]
lea  edx,[eax-0x19e8]
push edx ; "x= %d\n"
mov  ebx,eax
call 3b0 <printf@plt>
...
;return 0;
mov  eax,0x0

```



3.16

Array Manipulation (1)

test2.c:

```

int main(void){
    int a[2];
    int b[] = {7, 8};
    a[0] = 3;
    a[1] = 4;
    printf("%d %d\n", a[0], a[1]);
    printf("%d %d\n", b[0], b[1]);
    return 0;
}

```

```
$ gcc -Wall test2.c -o test2$ ./test2 3 4 7 8
```

- could be initialized when declared
- element indexing starts at 0
- when declared as local variables, they are allocated on the stack

3.17

Array Manipulation (2)

```


...
push ebp
mov  ebp,esp
push ebx
push ecx
sub  esp,0x20
...
;int a[2];

```

```

;int b[2] = {7, 8};
mov DWORD PTR [ebp-0x14],0x7
mov DWORD PTR [ebp-0x10],0x8
;a[0] = 3;
mov DWORD PTR [ebp-0x1c],0x3
;a[1] = 4;
mov DWORD PTR [ebp-0x18],0x4
...

```



3.18

Memory Corruption (1)

test3.c:

```

int main(void){
    int a[2];
    int b[] = {7, 8};
    a[0] = 3;
    a[1] = 4;
    a[2] = 5;
    printf("%d %d\n", a[0], a[1]);
    printf("%d %d\n", b[0], b[1]);
    return 0;
}

```

\$ gcc -Wall test3.c -o test3\$./test3 3 4 5 8 What happens when access an array out of its bounds?

- Java: an array is an object and all the operations on it are controlled, it knows its bounds
⇒ throws an exception
- C: the array is just a memory address, where the array starts
- C standard does not define behavior when array accessed out of bounds
 - anything could happen

3.19

Memory Corruption (2)


```

...
;int a[2];
;int b[2] = {7, 8};
mov DWORD PTR [ebp-0x14],0x7
mov DWORD PTR [ebp-0x10],0x8
;a[0] = 3;
mov DWORD PTR [ebp-0x1c],0x3
;a[1] = 4;
mov DWORD PTR [ebp-0x18],0x4
;a[2] = 5;
mov DWORD PTR [ebp-0x14],0x5
...

```

Question

What happens if change a[8]?



3.20

Pointers (1)

test_ptr.c:

```
void f1(int x){
    x = x * 2;
}
void f2(int *x){
    *x = *x * 2;
}


int main(void){
    int x = 3;
    f1(x);
    printf("%d\n", x);
    f2(&x);
    printf("%d\n", x);
    return 0;
}
```

- a pointer is a variable that contains a memory address, e.g. the address of another variable
- memory addresses are just integers (on 32 or 64 bits)
- useful when need referring some data, not copying it (e.g. reference parameters of a function)
- & - reference (get the address of a variable)
- * - dereference (get memory contents from a memory address)

3.21

Pointers (2)

```
; int x = 3;
mov DWORD PTR [ebp-0x10], 0x3
; f1(x);
mov eax, DWORD PTR [ebp-0x10]
push eax
call 57d <f1>
add esp, 0x4
; printf("%d\n", x);
...
; f2(&x);
sub esp, 0xc
lea eax, [ebp-0x10]
push eax
call 590 <f2>
add esp, 0x10
; printf("%d\n", x);
...
```

Arrays Are Pointers (1)

test_array_ptr.c:


```
#include <stdio.h>
void f(int *v) {
    v[0] = 25; // *v = 25
    *(v + 2) = 17; // v[2] = 17
    *((char*)v + 5) = 1; // !!!
    3[v] = 44; // v[3] = 44
}
int main() {
    int v[] = {1, 2, 3, 4};
    int i, n = sizeof(v)/sizeof(v[0]);
    f(v);
    for(i=0; i<n; i++){
        printf("%d ", v[i]);
    }
    printf("\n");
    return 0;
}
```

- v variable (i.e. the array name)
 - is a pointer to the beginning of the array
 - points where the array starts
 - it is a pointer of the same type as the array elements
- adding N (an integer) to a pointer \Rightarrow add “N*sizeof(ptr. type)”

```
$ ./test_array_ptr
25 258 17 44
```

Arrays Are Pointers (2)

```
;void f(int *v) {
push ebp
mov ebp,esp
...
;v[0] = 25;
mov eax,DWORD PTR [ebp+0x8]
mov DWORD PTR [eax],0x19
;*(v + 2) = 17;
mov eax,DWORD PTR [ebp+0x8]
add eax,0x8
mov DWORD PTR [eax],0x11
;*((char*)v + 5) = 1;
mov eax,DWORD PTR [ebp+0x8]
add eax,0x5
mov BYTE PTR [eax],0x1
...
}
```



$$2 \cdot 256^0 + 1 \cdot 256^1 = 258$$

Note: due to the **little-endian representation**

- the value of $v[1]$, i.e. $0x00000002$, stored in memory as bytes $0x02\ 0x00\ 0x00\ 0x00$ (from smaller to bigger addresses)
- $\Rightarrow v[1]$ becomes $0x02\ 0x01\ 0x00\ 0x00$

Data Structures and Pointers

```
struct Point {
    int x;
    int y;
};

struct MyStruct {
    int a;
    short b;
    struct Point p;
    char c[5];
};

int main(void){
    ...
}
```

Assigning values to structure fields

```
struct MyStruct s;
s.a = 7;
s.b = 12;
s.p.x = 150;
s.p.y = -11;
s.c[1] = 10;
s.c[2] = 'a';
```

Initialization when declared

```
struct MyStruct s = {
    .a=7, .b=12,
    .p={.x=150, .y=11},
    .c={0, 0, 0, 0, 0}
};
```

Pointers to data structures:

```
struct MyStruct s = {...};
struct MyStruct *ps;
```

How do we access the a field?

- with s : $s.a$
- with ps (v1): $(*ps).a$
- with ps (v2): $ps->a$


Recommended style: $(->)$.

How do we access (sub)field x of field p ?

- $s.p.x$
- $ps->p.x$

Data structure's size

```
printf("Point size = %d\n",
    sizeof(struct Point));
printf("MyStruct size = %d\n",
    sizeof(struct MyStruct));
```



```
Point size = 8
MyStruct size = 24
```

Each data structure field is aligned at DWORD (4 bytes).

Data Structure Alignment

How to force BYTE-alignment, if needed?

gcc (Linux)

```
struct __attribute__((packed)) MyStruct {
    ...
};
```

Visual Studio (Windows)

```
#pragma pack(push, 1)
struct MyStruct {
    ...
};
#pragma pack(pop)
```

Cast to data structures

```
unsigned char v[] = {
    7, 1, 0, 0, 12, 0, 0, 0,
    150, 0, 0, 0, 11, 0, 0, 0,
    'a', 0x62, 'c', 100, 0, 0, 0, 0
};
struct MyStruct *ps = (struct MyStruct*)v;
printf("a=%d b=%d p=(%d, %d), c='%s'\n",
    ps->a, ps->b, ps->p.x, ps->p.y, ps->c);
```

What will be displayed?

a=263 b=12 p=(150, 11), c='abcd'

3.25

Dynamic Memory Allocation (1)

Functions malloc and free

```
int *v = (int*)malloc(10000 * sizeof(int));
struct MyStruct *ps = (struct MyStruct*)
    malloc(sizeof(struct MyStruct));
...
free(v);
free(ps);
```

- malloc allocates memory area on the *heap* and returns a pointer to the allocated memory
- free releases an allocated memory area having a pointer to it

3.26

Dynamic Memory Allocation (2)

When does it make sense to allocate memory dynamically?

- when working with dynamic structures like
 - linked lists
 - trees, graphs
- when do not know in advance the (maximum) data size
 - we cannot declare a certain (maximum) size
- when data size is too large to be stored on the stack
 - default stack size on Linux: 8MB
 - default stack size on Windows: 1MB
 - changing default size on Linux: use “ulimit -s ...” or setrlimit()
 - changing default size on Windows: use the dwStackSize parameter of the CreateThread(...)

NOTE: dynamic allocation is slower than allocation on the stack (as local variable)

3.27

Dynamic Memory Allocation (3)

Memory leaks

- any dynamically allocated memory should be explicitly released with `free()`
 - otherwise it remains allocated until the program ends
- C language has **no garbage collector**
- sometimes memory is not release in the same function where it is allocated
 - *example*: a function allocates memory (for storing some results) and returns a pointer to that memory
 - the function getting such a pointer should release the memory, when not needing it anymore
- example of a classical memory leak (**don't do like this!**)

```
int x[1000];
int *p = (int*)malloc(1000 * sizeof(int));
p = x; // the only pointer to the dynamically allocated memory is lost
```

3.28

String Operations

- in C a string is
 - a byte array ended with byte 0 (zero or null) \Rightarrow null-terminated strings
 - each byte value is the code of a printable character
 - last byte is the unprintable character with code `'\0'` \Rightarrow **NUL-terminated strings**
- can be handled like a normal array
- in `string.h` there is a collection of **string handling functions**
 - `strlen()`
 - `strncpy()` (**do not use `strcpy()`!**)
 - `strcmp()`
 - `strchr()`
 - `strstr()`
- **string handling functions assume NUL-terminated strings**
 - \Rightarrow do not use them when strings not terminated with NUL

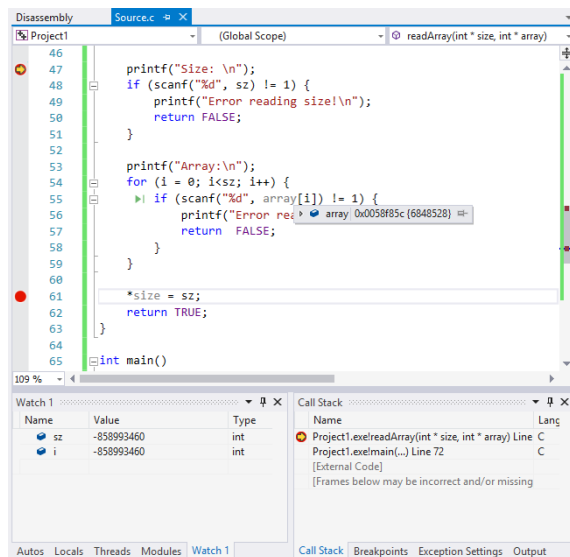
3.29

3 C Programs Debugging

Debugger

- a program that provides us the way to execute another program (app) step by step (instruction by instruction)
- usually in order to find and fix bugs
- common operations
 - execution step by step
 - setting breakpoints
 - monitor variable values
 - monitor memory contents
 - monitor the stack contents and evolution
- must compile the source code with explicit options to generate detailed (helpful) debugging information
 - like the `-g` option of `gcc`
 - debugging info: variable names, correlation to the source code

3.30



Debugging with a GUI (1)

- very convenient, simple and efficient
- e.g. set breakpoints by clicking in front of line of source code
- e.g. view variable value with *mouse over* or right-click → *Add Watch*
- trace / debug a program
 - *step into*: execute and go forward one instruction entering into functions
 - *step over* execute and go forward one line (even if a function, so not entering in such a function)
 - *continue*: execute and go forward to the next breakpoint or program's end

3.31

Debugging with a GUI (2)

3.32

Debugging with GDB

Running as: `gdb <executable_file>`

Displays an interactive shell where we can type commands like:

- `break test.c:12`
- `run`
- `continue`
- `step` (step into)
- `next` (step over)
- `bt` (*backtrace*: shows the stack frames)
- `print myvar`: displays the variable's contents

3.33

Postmortem Debugging

- provides us the way to investigate (debug) a crashed program's state
- useful when our programs run on other remote systems
- steps
 - activate core dumping on the remote system `ulimit -c unlimited`
 - run the program until its crash ⇒ a core file
 - loads the core dump into the gdb debugger `gdb <executable_file> core`

3.34

Using *valgrind* to Detect Memory Leaks

- allows analyzing a program during its runtime
- for finding out memory leaks it intercepts `malloc` and `free` (and other related functions) and keeps evidence of the allocated memory areas
- at program termination displays a report with memory leaks

```
$ valgrind ./<executable\_file> [<arguments>]
...
==9347== LEAK SUMMARY:
==9347==    definitely lost: 55 bytes in 5 blocks
==9347==    indirectly lost: 0 bytes in 0 blocks
==9347==    possibly lost: 0 bytes in 0 blocks
==9347==    still reachable: 43 bytes in 3 blocks
==9347==    suppressed: 0 bytes in 0 blocks
```

- when there are leaks run it again for detailed info `-leak-check=full -show-leak-kinds=all`

3.35

4 Recommendations About Writing C Programs

There Are No Perfect Programs

- it is almost impossible to write a complex bug-free program
 - at least from first try
- every day new bugs are found in real-life “professional” (commercial) software run by millions of users
- there are metrics that account bugs per thousands of lines of code

A good programmer could be allowed to not write perfect programs, but (s)he is required to be able to test / debug them, identify bugs and fix them!

- testing and debugging a program has same importance like writing code
 - understanding a written code is not at all the same like writing it by yourself from scratch

3.36

Using Code From Other Sources

- **Do not copy-paste code you do not understand!**
- actually, **never copy-paste any code**, at least for your school projects!
- there is (almost) nothing to learn from others’ code
- if really need to use someone else’ code (in real life)
 - make sure you understand and control that code
 - mention the source of your code (at least in a comment)

3.37

A Simple Editor or an IDE?

- IDE = Integrated Development Environment
- IDE provides an integrated environment for
 - code editing with auto-completion, suggestions
 - code compilation
 - program running and debugging
 - code refactoring
- IDE’s advantages
 - development efficiency
 - automatize and make efficient frequent and complex development processes
- IDE’s disadvantages
 - during the learning phase has “training wheels” effect
 - when learning a new language / technology is better to use basic editor and tools
 - could be slower than basic tools

3.38

Watch the Uninitialized Variables

- local variables are allocated on the stack
- just declaring a variable leads to assembly code (actually, machine code generated by compiler) that simply decrease the ESP register
 - i.e. reserve space on stack for the declared variable
- \Rightarrow initial value of the variable depends on the current contents of the reserved space
- **initialize variables before using them!**
 - where they are declared or as close as possible to their declaration
- dynamically allocated memory is also uninitialized, i.e. initialized with undefined values
- after calling `free(p)`, `p` will point to an undefined memory location
 - \Rightarrow **dangling pointer**
 - **after releasing the memory a pointer points to, assign it NULL value!**

3.39

Watch the Global Variables

- global variable not explicitly assigned values, are initialized with 0
- global variables are visible from any function of your program
 - \Rightarrow if not a constant, anyone can change them
- global variables can be changed concurrently from multiple threads
 - \Rightarrow unexpected, unpredicted values
 - such a code is called thread-unsafe
- **do not use global variable, when not really needed!**
 - better and safer to give a function the needed context information as parameters

3.40

Do not use uninitialized local variables!
Avoid using global variables!
Do not use dangling pointers!

3.41

Warning Handling

- compiler warns us when something in the compiled program is confusing regarding the use of some variables or functions
 - a logical expression (e.g. a condition in an `if`) using a single `'='` instead of two
 - code with no effect
 - usage of uninitialized variables
 - calling `printf()` with a strange or incorrect format string
- when we want to ignore some types of warnings, so not be reported about them, we can specify certain compiling options
 - be aware that too many irrelevant warnings, could hide from us the important ones
- for the release version of a program is recommended the `-Werror` option
 - just to report an error for any warning

3.42

5 Conclusions

What We Talked About!

- basic aspects related to C programs
- getting an executable from the source code
- local variables on the stack
- pointers, arrays, data structures
- debugging aspects
- coding recommendations

3.43

Lessons Learned

- local variables are allocated on stack
 - their initial value is unknown
 - stack could be corrupted due to overflowing local arrays
- dynamically allocated memory should be released (i.e. freed)
- pointers give us direct access to application's memory
 - take care at pointer arithmetic!
 - take care of memory overflowing or corruption!
 - do not use uninitialized pointers!
- array names are pointers to where the array is stored
- debugging and testing have same importance as writing code

3.44