

Chapter 4.1

The File System

The User Perspective

Print Version of Lectures Notes of *Operating Systems*

Technical University of Cluj-Napoca (UTCN)
Computer Science Department

Adrian Coleşa

March 18th, 2020

4.1.1

Purpose and Contents

The purpose of today's lecture

- General Overview of The File System Module
- File Concept
- Directory Concept

4.1.2

Bibliography

- Andrew Tanenbaum, *Modern Operating Systems*, 2nd Edition, 2001, Chapter 6, pg. 380 – 398.

4.1.3

Contents

1	File System (FS) Overview	1
2	Fundamental Concepts	2
2.1	File	2
2.2	Directory	6
3	Conclusions	8

4.1.4

1 File System (FS) Overview

Context

- users need to **store** and **retrieve**
 - **persistent data**
 - **large amount of data**
- users need to **share data**
- ⇒ OS should
 - **provide services** for such needs
 - **manage** storage devices (area) and stored data

4.1.5

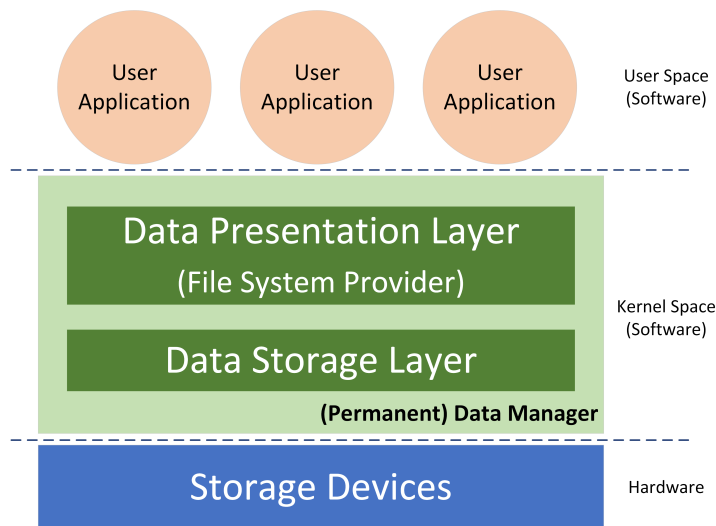


Figure 1: Data Manager's Layers. What we call "File System" is the way the "Data Manager" makes visible the data storage space to users.

Definition

- **an OS's component**
 - \Rightarrow part of OS
- the **interface** between the user and the physical storage devices
 - *provides* the users access to the storage area
 - *manages* the information on the storage area
- \Rightarrow the (permanent) **data manager**

4.1.6

Role

- **provides** the users a simplified, uniform and **abstract view of the storage area**
 - hides the complexity of physical storage devices
 - abstract concepts: *file* and *directory*
 - \Rightarrow *its name: File System*
- **manages** data on the storage area
 - interacts directly with storage devices
 - \Rightarrow contains a hardware dependent layer (code modules)

4.1.7

File System Architecture. General View

4.1.8

File System Architecture. Detailed View

4.1.9

2 Fundamental Concepts

2.1 File

Definition. User Perspective

- the basic **storage unit** of information
 - anything the user wants to store must be placed in a file
- provides a (convenient) way to
 - **store** the information on storage devices
 - **retrieve** the information back later
- **a container, i.e. a box**
 - a collection of related information defined by its creator

4.1.10

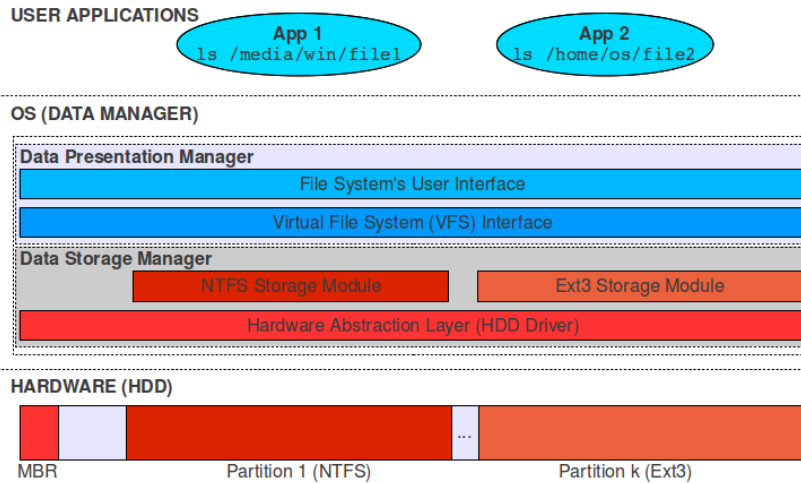



Figure 2: Data Manager's Components



The file is the basic abstract concept for storing user data!

4.1.11

Definition. OS (Internal) Perspective


- an abstraction mechanism: a *container*
 - models the way data is provided to users
 - does not necessarily correspond to the way data is stored (on physical devices)
- OS must
 - **store** the container's contents on storage area
 - **retrieve** back later the container's contents
- the file maps the user view of data to the physical data
 - identifies on physical devices and links together physical data (i.e. chunks of bytes) belonging to a file

4.1.12

File's Components

- **data**: user's useful information placed in the container
- **metadata**: description associated to stored data
 - needed, maintained and sometimes imposed by OS
 - some fields also used by user
- the two components **can be stored in different parts** on a storage area
 - there is a link maintained from metadata to data
 - \Rightarrow FS always **starts from a file's metadata** in order to access that file

4.1.13



File Metadata: Name

- the user's way to **identify** (refer to) a file
- **must be unique** (there are though some exceptions!)
- each file must have (at least) a name
- consists of a **string of characters**
 - upper vs. lower letters, char sets
 - its length is generally limited
- file name's **extension**
 - normally **not imposed by OS**
 - *rarely*: may be recognized and used by the OS
 - *usually*: a **user-space convention** to provide hints about file's contents
 - required by some application (e.g. *gcc* requires *.c*)
 - examples: *book.docx*, *letter.odt*, *setup.exe*, *arch.tar.gz*, *progr.c*, *image.jpg*

4.1.14

From the OS perspective every filename is just a string of (unrestricted) characters!

4.1.15

File Metadata: Type

- Regular files
 - Contain user data (text or binary)
- Directories
 - System files used to organize file space
- Links
 - System files used to redirect the access to other files
- Special files
 - Model I/O devices
 - *Character*: terminals, mouse, etc.
 - *Block*: disks
- Pipes
 - Inter-process communication (IPC) mechanisms

4.1.16

File Metadata: Structure

- **file's structure = the way the file's contents is organized and interpreted**
- possible strategies
 1. **no structure**: just a *sequence of bytes*
 - used by most OSes for normal files
 2. **specialized structures**: used normally for system files
 - sequence of (fixed-length) records
 - tree of records (e.g. B-Trees)

4.1.17

File Metadata: Type vs. Structure – Pros and Cons

- different **file types** \Leftrightarrow different **file structures**
- that is **used for system files** (directories, links, pipes etc.)
 - managed by OS transparently from the user
- **But ...** should the OS support **types and structures for regular files**?
 - **if, yes**
 - * +: **efficient** / particular file's contents manipulation
 - * +: convenient for novice users (not much / complex code)
 - * -: additional OS code
 - * -: restrictions, rigidity
 - * -: file not portable
 - **if, no**
 - * +: **flexible** \Rightarrow convenient for advanced users
 - * +: no additional OS code
 - * -: no OS support, additional application code
 - * \Rightarrow **each application must manage / interpret itself the contents of files it works with**

4.1.18

File Metadata: Type vs. Structure – Real Life OSes

- **no structure for regular files**
 - *text* vs. *binary* files: just a user convention
 - basically, **all files are binary**, i.e. each file's contents must be handled specifically by applications using it
 - \Rightarrow *text file* is just a particular file type
- \Rightarrow **flexibility** (*separate mechanism by policy*)
- **yet, every OS recognizes its own executable files**

4.1.19

From the OS perspective every file is just a sequence of bytes!

4.1.20

Question

Yet, from the user perspective ... how would you classify the following files in terms of **text** vs **binary**, based on their filename extension?

1. program.c
2. archive.zip
3. paper.pdf
4. index.html
5. persons.xml
6. app.java
7. cv.docx

4.1.21

File Metadata: Attributes

- system attributes
 - type
 - length
 - owner
 - access permissions (rights): basically *read*, *write* and *execute*
 - time stamps (e.g. creation time, last access time, last modification)

- addresses of allocated blocks for that file
 - * named during that course *Block Addresses Table* (BAT)
 - * the link between metadata and data
- ...
- user attributes (if supported)
 - anything the user wants
 - examples: *source Web address* (for an html file), *place* (for an image file), *author* (for a document) etc.

4.1.22

File Access Method

- **sequential** access
 - **impose an order** on the way the bytes are accessed
 - could be imposed by the storage device (e.g. tapes)
 - could be imposed by the nature of the file's contents
 - * e.g. **specific for directories**
- **random** access
 - accesses the bytes or records **out of order**
 - specific to storage devices like hard disks
 - normally, **specific to regular user files**

4.1.23

Operations On Files

- Access file data
 - *open*
 - **write**, i.e. *store* data
 - **read**, *get back* stored data
 - position (seek), i.e. *randomly* accessing stored data
 - *close*
- Manipulate files/Access file metadata
 - create
 - get/set attributes
 - rename
 - truncate
 - delete

4.1.24

2.2 Directory

Definition. User Perspective

- the way to **organize / classify files**
 - when too many, difficult to find them based just on their name
- a **collection / class** of related elements (files)
 - a file could be placed in a directory (or more?)
 - ⇒ directory's name is a sort of additional (user) metadata associated to the file
- reduces the size of the filename space
 - files in different directories could have the same name
- imposes a structure / **hierarchy** on the file space
 - helps the user locating files easier in huge file spaces
 - **finding files visually** in the hierarchy is based on **navigation**
 - * refine filters gradually, based on subdirectory names
 - does not help (too much) the **user applications locating a particular file**
 - * if not knowing in advanced a file's complete path
 - * ⇒ **file must be searched** in the entire FS tree

4.1.25

The directory is the abstract concept for organizing user files!

4.1.26

Definition. OS (Internal) Perspective

- an abstraction mechanism of grouping files
- a **container**: data (bytes) that must be stored
 - a **special file managed by OS** transparently from user
 - i.e. directory's bytes are interpreted by the OS and user-applications provided directory's elements
 - usually organized as specialized searching data structures, e.g. B-trees
- it also **consists of data and metadata**
 - very similar to a regular file
 - \Rightarrow just a FS-element from the OS perspective

4.1.27

Directory Hierarchy

- types
 - single-level directory systems
 - two-level directory systems
 - **hierarchical** directory systems
 - * the most general form: **tree** or **graph**
- **file paths**
 - the way of **identify a file in a hierarchy**
 - a list of consecutive nodes ending with the file name

4.1.28

File Paths

- **absolute paths**
 - starting node is the **root directory**
 - examples
 - * c:\Program Files\Application\run.exe (Windows)
 - * /home/students/adam/program.s (Linux)
 - usage: access files at known fixed places
- **relative paths**
 - not starting from the root directory
 - * starting node is the **current directory**
 - each application has its own current directory associated to it
 - examples
 - * Application\run.exe (Windows)
 - * students/adam/program.s (Linux)
 - special notations (directories)
 - * current directory: '.'
 - * parent directory: '..'
 - usage: access files in a subtree with a known structure, but located to an unknown location in the overall FS tree

4.1.29

The directory concept imposes a hierarchy on the file space, requiring a path in order to identify a file!

4.1.30

Question

Supposing the current working directory of an application is “/home/os”, which will be the corresponding absolute paths for the following relative paths?

1. file_1
2. ./file_2
3. project/file_3
4. ../file_4
5. ../../file_5
6. ../../../../file_6
7. ../../etc/file_7
8. ../../../../etc/apache2/../../../../../home/os/./file_8

4.1.31

Operations On Directories

- create
- delete
- rename
- opendir
- readdir
 - **sequential access**
 - positioning not allowed anywhere
- rewind
- “write” (**not directly allowed**)
 - add elements (create directories and files)
 - delete elements (delete directories and files)
- close

4.1.32

3 Conclusions

What We Talked About

- **file system (FS)** as an OS component
 - role
 - general structure
- **file** concept
 - used for **storing user data**
 - components: **data** (container component) and **meta-data** (description component)
 - meta-data fields: name, types, attributes etc.
 - **structure**: sequence of bytes, specialized collection of elements
- **directory** concept
 - used for **organizing the file space**
 - a system file: its bytes are interpreted by the OS and provided as a collection of elements
 - imposes a **hierarchy**
 - FS elements (e.g. files) specified by a **path**
 - absolute vs relative path

4.1.33

Lessons Learned

- the basic storage unit for the user (applications) is the file
 - \Rightarrow storing just a single byte of data needs placing it in a file
- OS usually not involved in the interpretation of a file's contents
 - \Rightarrow file structure (format) is the user-application business
- text vs binary files is just a user convention
 - from the OS perspective all files consist just in a sequence of bytes
- FS organization, i.e. classifying files, is a real need
 - helps the user to navigate visually in the FS space in order to find her files
 - does not help (too much) the user applications finding files