

Chapter 4.2

Linux and Windows File Systems

Permission Rights and System Calls

Print Version of Lectures Notes of *Operating Systems*

Technical University of Cluj-Napoca (UTCN)
Computer Science Department

Adrian Coleşa

March 25, 2020

4.2.1

Purpose and Contents

The purpose of today's lecture

- Presents and compare *permission rights* in Linux and Windows
- Presents and compare *system calls for files* in Linux and Windows

4.2.2

Bibliography

- A. Tanenbaum, *Modern Operating Systems*, 2nd Edition, 2001, Chapter 10. Case Study 1: Unix and Linux, pg. 732 - 744, p. 753 - 757
- A. Tanenbaum, *Modern Operating Systems*, 2nd Edition, 2001, Chapter 11. Case Study 2: Windows 2000, pg. 830 - 833, p. 844 - 847
- Lab texts related to Linux's and Windows' file system and their system calls.
- From [http://msdn.microsoft.com/en-us/library/aa364407\(v=85\).aspx](http://msdn.microsoft.com/en-us/library/aa364407(v=85).aspx) about File management and Directory Management (following the lecture slides)

4.2.3

Contents

1	Permission Rights	1
1.1	Linux Permission Rights	1
1.2	Windows' Permission Rights	5
2	Basic System Calls on Linux and Windows	9
2.1	Linux	9
2.2	Windows	19
3	Conclusions	21

4.2.4

1 Permission Rights

1.1 Linux Permission Rights

Basic Permission Rights

- defined for three classes of users
 - owner or user (u)

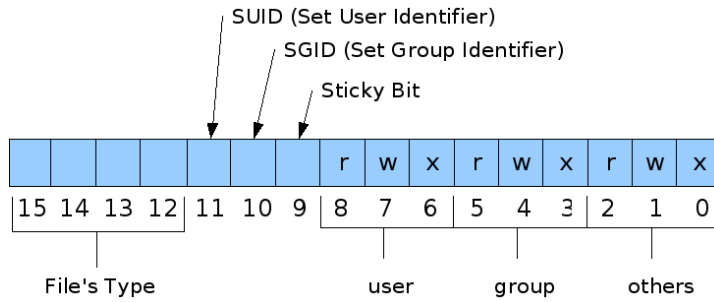


Figure 1: The way permission rights are stored

- groups (g) the owner belongs to
- others (o)
- operations (types)
 - *read* (r)
 - *write* (w)
 - *execute* (x)
- See examples running `ls -l` command on different files and directories

4.2.5

I-node Field Structure

- bit value: 0 / 1 → denied / allowed
- example
 - string: `rw-r--r--`
 - binary: `110100100`
 - octal: `0644`

4.2.6

Permissions on Regular Files

- **read**: read file's contents
- **write**: write (*modify, append to, truncate*) file's contents
- **execute**: execute file

4.2.7

Permissions on Directories

- **read**: read (list) directory's contents
- **write**: write (*modify, add and remove elements*) file's contents
 - confusing and too limited
- **execute**: **traverse directory**, i.e. search for an element in the directory
- ⇒ **Read and/or Write without Execute not so useful**
 - but ... **Execute without Read and/or Write makes sense**
 - when we want a directory to be traversed, but its contents not be visible
 - commonly used in practice for the `/home` directory

4.2.8

Basic permission rights for FS elements in Linux are r, w, x for u, g, o `rw-rw-rw-`

4.2.9

Questions (1)

Give the equivalent permission right representation for the following cases?

1. `rwxr-xr--`
2. `r--r--r--`
3. `0765`

4.2.10

Questions (2)

Which of the following operations

- (O1) `ls /home/os`
(O2) `cat /home/os/file.txt`
(O3) `rm /home/os/file.txt`

could be performed by the “os” user, supposing the directory “/home/os” is its home directory and has the following permission rights (all the file in “/home/os” have `r--r--r--` permissions)?

- (P1) `r-xr--r--`
(P2) `--x--x--x`
(P3) `rw-r--r--`

4.2.11

Special Techniques: SUID and SGID Bits

- SUID
 - has effect only for executable files
 - the process resulting from the corresponding executable file will have the effective UID that of the owner of the file
 - see the classical example of `/usr/bin/passwd` executable file, that can be run by any *non-privileged user*, modifying the `/etc/passwd` file belonging to *root*
- SGID
 - similar to SUID but applies to files’s GID

4.2.12

Special Techniques: Sticky Bit

- has effect only for directories
- allows elements to be removed only by their owner
- see `/tmp` directory

4.2.13

Extended Attributes and ACL-like Permission Rights

- extended attributes
 - extra attributes like “append-only”
 - see “man chattr”
- Access Control List (ACL)
 - list of complementary permission rights per user / group
 - see “man getfacl”

4.2.14

Security Considerations: Weak Permissions

- **permission rights are essential** for
 - file protection
 - **system security**
- pay attention to
 - **not allow** read / write permission on vital system files
 - **not allow** write permission on system directories
 - **not trust** files / directories writable by regular users (possible attackers)
- attack scenarios
 - readable “/etc/shadow” file allows for brute-force password guess
 - writable “/etc/passwd” allows for creation of new users
 - writable “/etc/sudoers” allows getting root (admin) permissions
 - writable “/sbin” allows replacement of system executables
 - ...

4.2.15

Security Considerations: Study Case (Bad Code)

Consider the following C program “fopen-permissions.c”

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>

int main( void)
{
    FILE *f;

    umask(0000);

    if ((f=fopen("TEST", "w+")) == NULL) {
        perror("File creation error");
        exit (1);
    }

    printf("File TEST created with the following permission rights\n");

    system("stat --format=%A TEST");

    unlink("TEST");

    return 0;
}
```

When run, the program displays

```
$ gcc -Wall fopen-permissions.c -o fopen-permissions
```

```
$ ./fopen-permissions
```

```
File TEST created with the following permission rights
-rw-rw-rw-
```

4.2.16

Security Considerations: Study Case (Good Code)

Consider the following C program “fopen-permissions.c”

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>

int main( void)
{
    FILE *f;

    umask(0022);

    if ((f=fopen("TEST", "w+")) == NULL) {
        perror("File creation error");
        exit (1);
    }

    printf("File TEST created with the following permission rights\n");

    system("stat --format=%A TEST");

    unlink("TEST");

    return 0;
}
```

When run, the program displays

```
$ gcc -Wall fopen-permissions.c -o fopen-permissions
$ ./fopen-permissions
File TEST created with the following permission rights
-rw-r--r--
```

4.2.17

Never trust the users! Protect files using the appropriate permission rights!

4.2.18

1.2 Windows' Permission Rights

Strategy

- Permission rights are defined for each user using ACLs (*Access Control Lists*)
- Basically, they are: *read* (r), *write* (w), *execute* (x)

4.2.19

Simplified (Synthesized) View

- Files
 - *read*: permits viewing or accessing of the file's contents
 - *write*: permits writing to a file
 - *read&execute*: permits viewing and accessing of the file's contents as well as executing of the file
 - *modify*: permits reading and writing of the file; allows deletion of the file
 - *full control*: permits reading, writing, changing and deleting of the file
- Directories
 - *read*: permits viewing and listing of files and subdirectories
 - *write*: permits adding of files and subdirectories
 - *read&execute*: permits viewing and listing of files and subdirectories, as well as executing of files
 - *modify*: permits reading and writing of files and subdirectories; allows deletion of the directory
 - *full control*: permits reading, writing, changing, and deleting of files and subdirectories

4.2.20

Advanced View

- Traverse Folder/Execute File
- List Folder/Read Data
- Read Attributes
- Read Extended Attributes
- Create Files/Write Data
- Create Folders/Append Data
- Write Attributes
- Write Extended Attributes
- Delete Subfolders and Files
- Delete
- Read Permissions
- Change Permissions
- Take Ownership

4.2.21

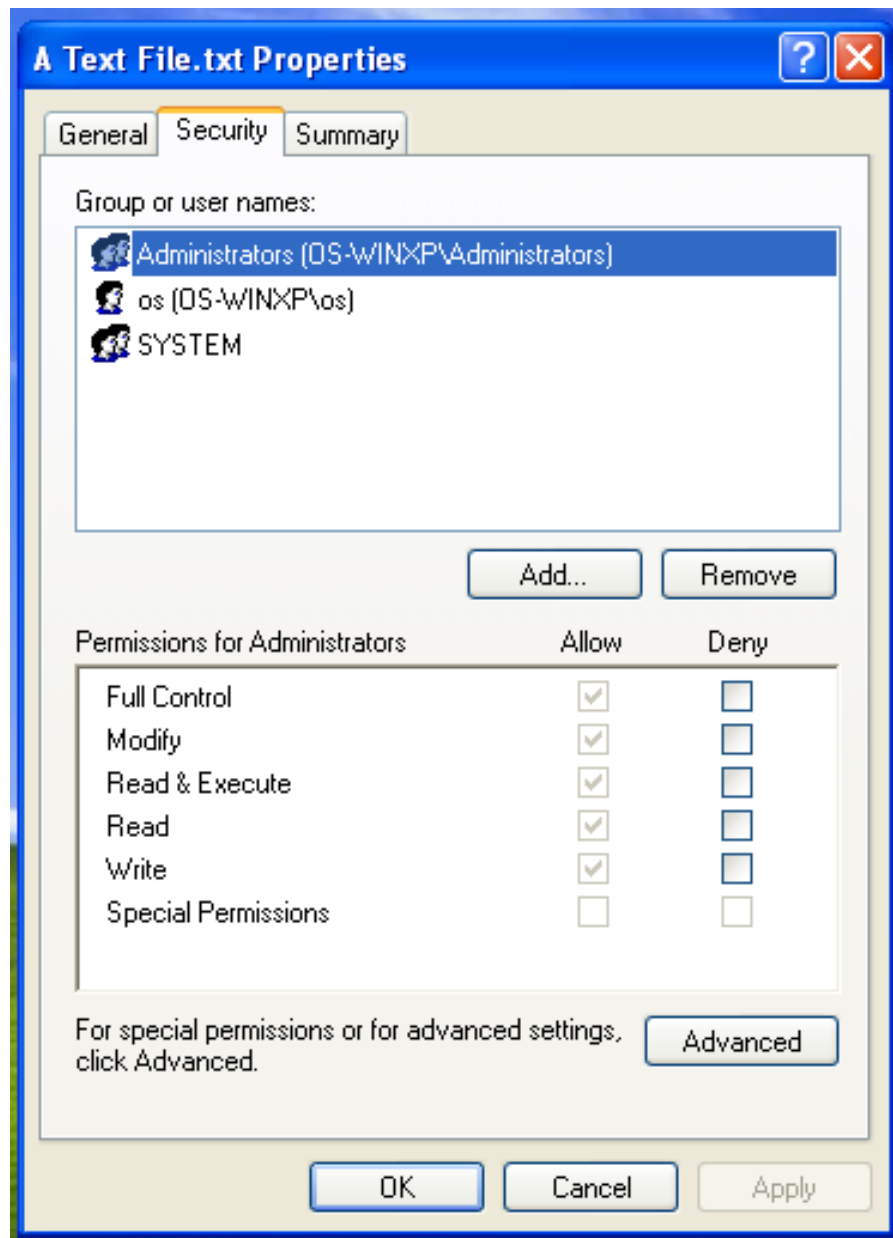


Figure 2: The Simplified View of Permission Rights

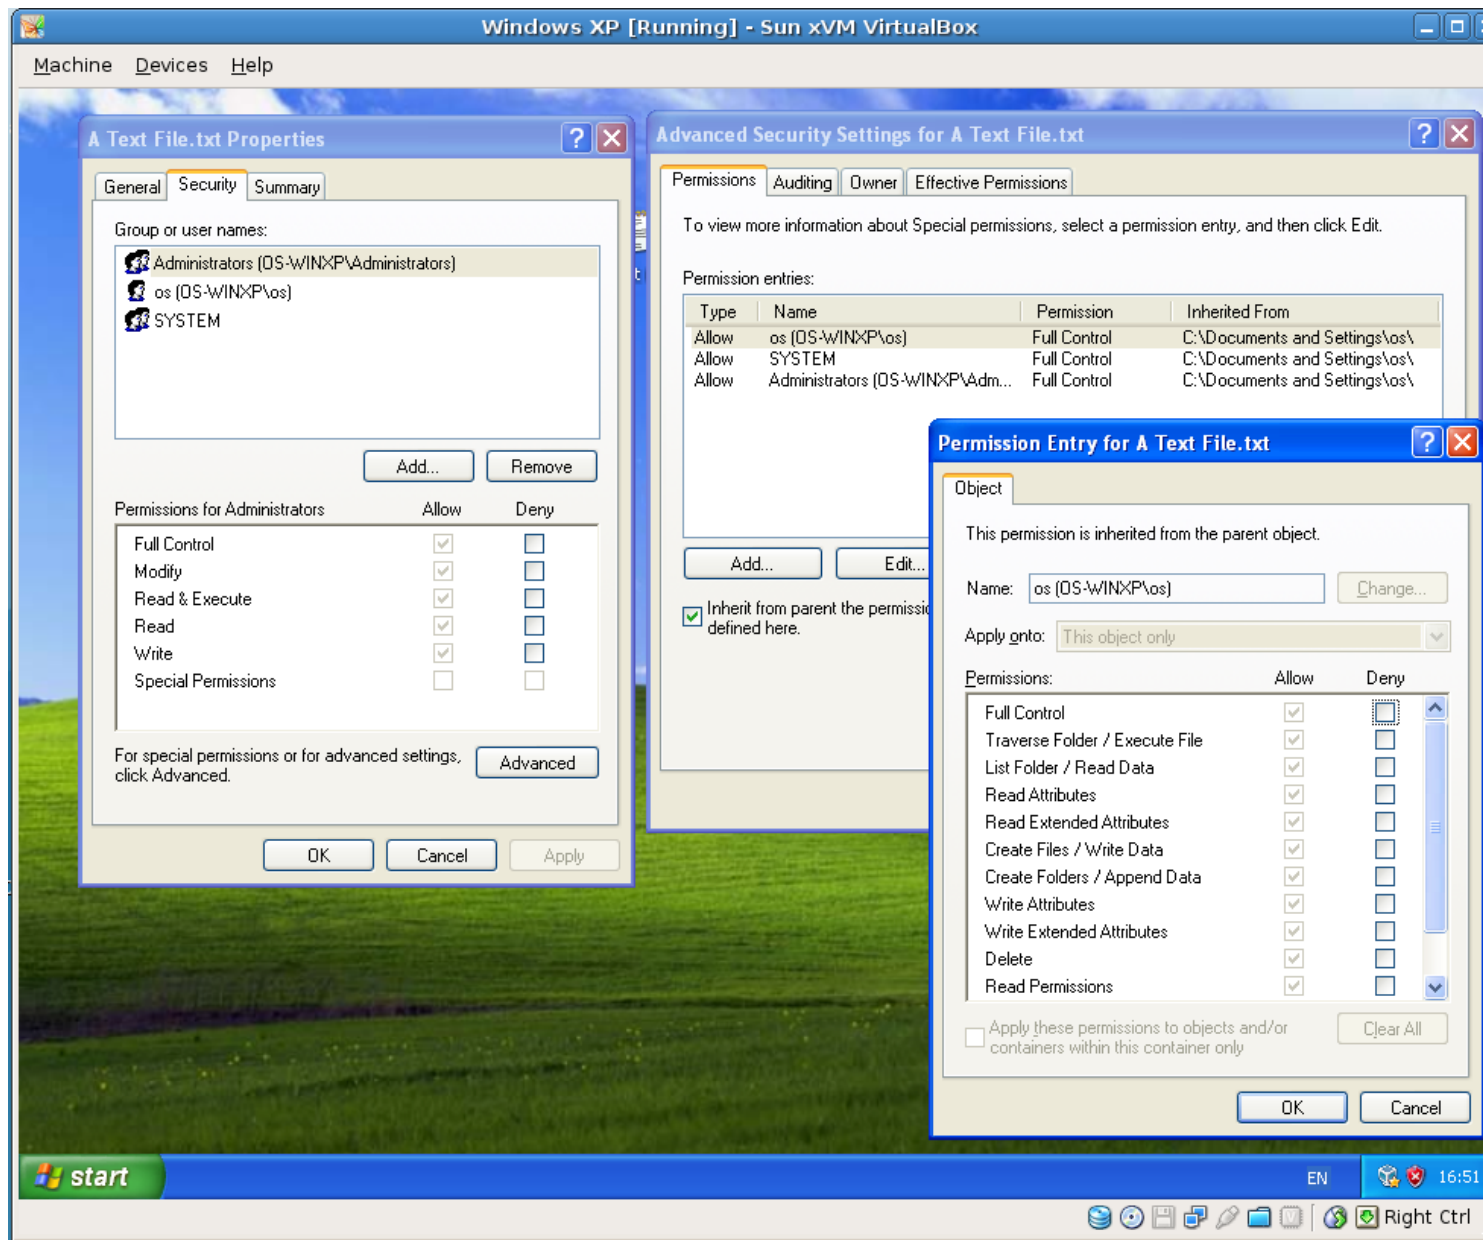


Figure 3: The Way of Setting Detailed Permission Rights (Windows XP)

Permission Entry for A Text File.txt

Principal: Administrators (ACOLESA-LPT\Administrators) [Select a principal](#)

Type: Allow

Advanced permissions:

<input checked="" type="checkbox"/> Full control	<input checked="" type="checkbox"/> Write attributes
<input checked="" type="checkbox"/> Traverse folder / execute file	<input checked="" type="checkbox"/> Write extended attributes
<input checked="" type="checkbox"/> List folder / read data	<input checked="" type="checkbox"/> Delete
<input checked="" type="checkbox"/> Read attributes	<input checked="" type="checkbox"/> Read permissions
<input checked="" type="checkbox"/> Read extended attributes	<input checked="" type="checkbox"/> Change permissions
<input checked="" type="checkbox"/> Create files / write data	<input checked="" type="checkbox"/> Take ownership
<input checked="" type="checkbox"/> Create folders / append data	

Add a condition to limit access. The principal will be granted the specified permissions only if conditions are met.

[Add a condition](#)

Figure 4: The Advanced View of Permission Rights (Windows 10)

	Full control	Modify	Read & Execute	List folder contents	Read	Write	Special permissions
Full control	X						
Traverse folder/Execute file	X	X	X	X			
List folder / Read data	X	X	X	X	X		
Read Attributes	X	X	X	X	X		
Read extended attributes	X	X	X	X	X		
Create files/Write data	X	X				X	
Create folders/Append data	X	X				X	
Write attributes	X	X				X	
Write extended attributes	X	X				X	
Delete subfolders and files	X	X					X
Delete	X	X					
Read permissions	X	X	X	X	X		
Change permissions	X						X
Take ownership	X						X

Relationship Between Synthesized and Advanced Permission Rights

4.2.22

2 Basic System Calls on Linux and Windows

2.1 Linux

Access File Data

```
int open(const char *pathname, int flags);
int open(const char *pathname, int flags, mode_t mode);
ssize_t read(int fd, void *buf, size_t count);
ssize_t write(int fd, const void *buf, size_t count);
off_t lseek(int fd, off_t offset, int whence);
int close(int fd);

int dup(int oldfd);
int dup2(int oldfd, int newfd);
```

4.2.23

Manipulate Files

```
int creat(const char *pathname, mode_t mode);
int rename(const char *oldpath, const char *newpath);
int truncate(const char *path, off_t length);
int ftruncate(int fd, off_t length);
int stat(const char *path, struct stat *buf);
int fstat(int fd, struct stat *buf);
int lstat(const char *path, struct stat *buf);
int chmod(const char *path, mode_t mode);
int fchmod(int fd, mode_t mode);
int chown(const char *path, uid_t owner, gid_t group);
int fchown(int fd, uid_t owner, gid_t group);
int lchown(const char *path, uid_t owner, gid_t group);
```

4.2.24

Manipulate Directories

```
int mkdir(const char *pathname, mode_t mode);
int rmdir(const char *pathname);
int link(const char *oldpath, const char *newpath);
int unlink(const char *pathname);
int symlink(const char *oldpath, const char *newpath);
DIR *opendir(const char *name);
struct dirent *readdir(DIR *dir);
void rewinddir(DIR *dir);
off_t telldir(DIR *dir);
void seekdir(DIR *dir, off_t offset);
int closedir(DIR *dir);
```

4.2.25

Create and Remove a File

```
int fd;

// create a new file
fd = creat("/home/os/file", 0600);
// file size = 0 (no space allocated)
// only the i-node (metadata) allocated
// if file exists, it is truncated
// the new file is opened for WRONLY
// note permissions: rw-----

// remove the file (remove a link to the file)
unlink("/home/os/file");
```

4.2.26

File Structure (Format): What Is The File Provided Like?

- no (special) structure
- just a sequence (stream) of bytes
 - each byte has its fixed position (offset)
- **no special byte(s)** inside the file to mark the end of file
 - every byte in the file could have any possible (user-provided) value
 - the file size kept as an i-node field

4.2.27

Relationship Between the OS and User Views of A File

- user (application) view
 - unstructured sequence of bytes
 - \Rightarrow a logically contiguous area
- OS logical view
 - sequence of blocks
 - \Rightarrow a logically contiguous area
- OS physical view
 - collection of blocks
 - \Rightarrow a collection of more physical contiguous areas

4.2.28


Illustration of Different File Views

4.2.29


Open Files Management

- accessing a file is done using open files
 - opening a file let the OS prepare for the next read/write operations
 - let the OS be very efficient
- OS maintains three types of tables (i.e. internal structures)
 - **i-node table (IT)**: one per system
 - **open file table (OFT)**: one per system
 - **file descriptor table (FDT)**: one for each process
- **every open \Rightarrow a new open file \Rightarrow different entry in OFT**
 - read operations are independent
 - write operations are independent, but ...
 - write effect is immediately visible (*one-copy semantic*)

4.2.30




User logical view =
a sequence of **BYTES**



OS logical view =
a sequence of **BLOCKS**



..... Block Address Table (BAT)



OS physical view =
a collection of **BLOCKS**

HDD (storage area)

Open a File

```
int fd;

// open an existing file
fd = open("/home/os/file_1", O_RDWR);
// file must exist
// opened for both RD and WR (if allowed by permission rights)

// ALWAYS CHECK return values of I/O operations!
if (fd < 0) {
    // -1 returned in case of error
    // e.g. file does not exists (wrong filepath)
    // e.g. permission denied
    perror("File cannot be opened"); // display the system err msg.
    exit(1);                         // terminate program
}

// create a file with "open"
fd = open("/home/os/file_2", O_CREAT | O_EXCL | O_RDWR, 0600);
// O_EXCL check if file does not already exist
```

4.2.31

Reading From/Writing To A File

- operations are performed relative to the *current position*
 - most of the cases at beginning of file after open
- current position is advanced (increased)
 - with the number of bytes successfully read or written
 - by each read and write, respectively
- both read and write syscalls
 - use *memory address* where bytes are
 - * written to (after being *read*)
 - * taken from (to be *written*)
 - return the number of bytes successfully read or written
- **end of file (EOF) detected when read returns 0 (zero)**

4.2.32

Example: Opening Files, Reading From Them, Creating Duplicates

```
// Process 1
int fd1, fd2;
char buf[100]="1234567890";

fd1=open("file1", O_RDWR); // fd1 = 3
fd2=open("file1", O_RDONLY); // fd2 = 4
write(fd1, buf, 10); // write "1234567890"
read(fd2, buf, 5); // read "12345"

// Process 2
int fd1, fd2, fd3;
char buf[100];

fd1=open("file1", O_RDWR); // fd1 = 3
fd2=open("file2", O_RDWR); // fd2 = 4
fd3=dup(fd2); // fd3 = 5
write(fd2, buf, 100); // write first 100 bytes
write(fd3, buf, 100); // write next 100 bytes
```

4.2.33

Example: Open File Tables Illustration

4.2.34

Each “open()” leads to a new, independent open file, i.e. a new entry in OFT and FDT.

4.2.35

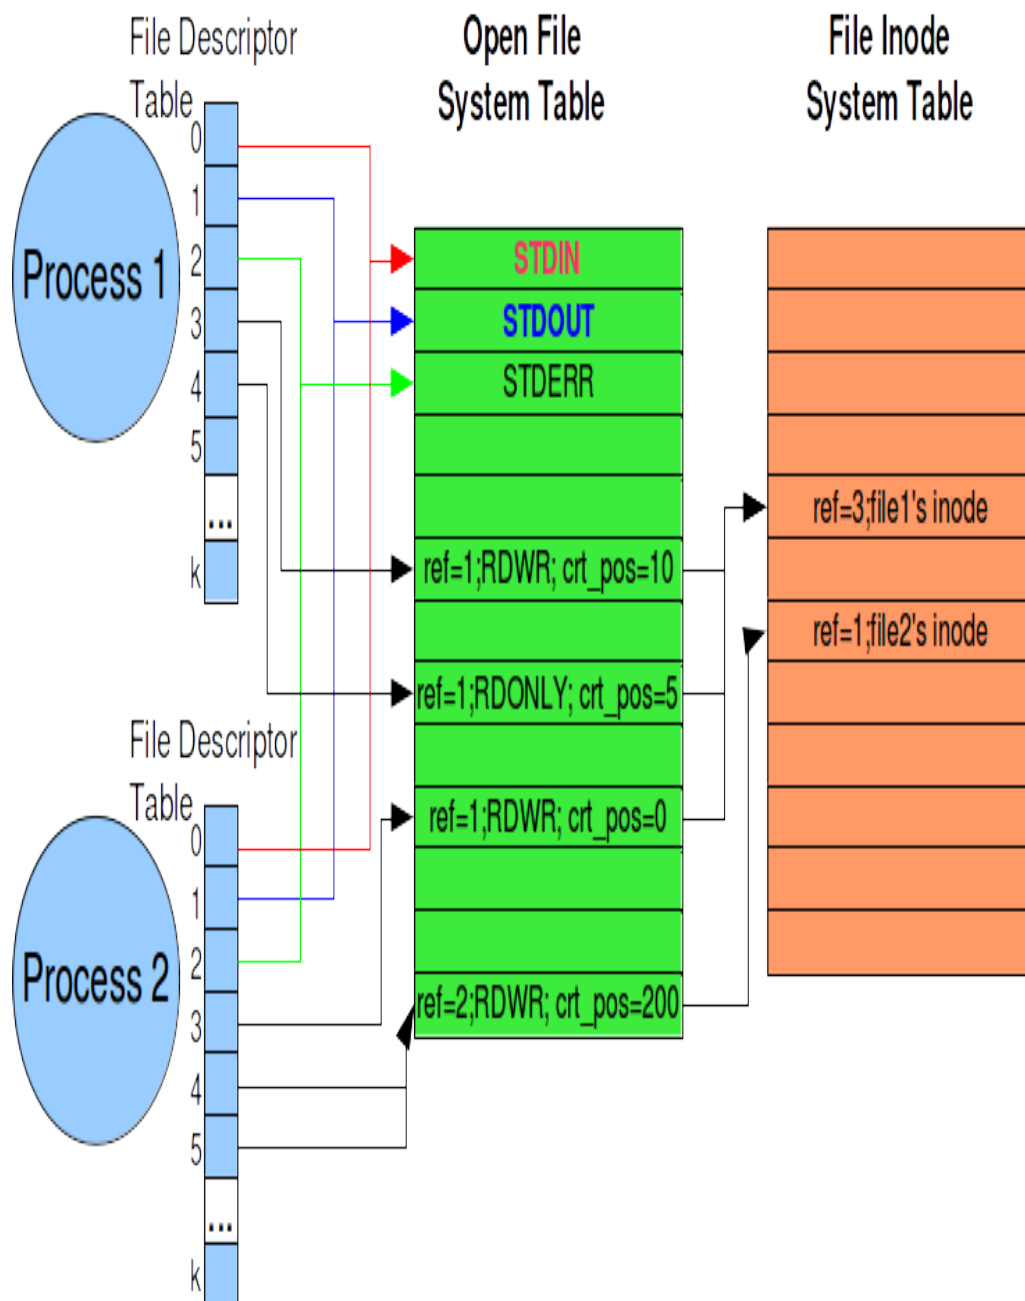


Figure 5: The Effect of More Open and Read Operations

File Format: Text Files

- each byte has its own interpretation
 - contains the code of a printable character
- special characters: new line (0x0A)
- OS knows nothing about (ANY) file's format
 - \Rightarrow read has no notion of reading until some special char
 - i.e. read cannot read a text line
 - **end of line has to be detected by application**

4.2.36

Example: Read the First Text Line

```
#define MAX_LINE 1024

int fd, i;
char c;
char line[MAX_LINE+1];

fd = open("file.txt", O_RDONLY);
if (fd < 0) {
    perror("Cannot open the file");
    exit(1);
}

i=0;
while ( (i < MAX_LINE) && (read(fd, &c, 1) > 0) && (c != '\n')) {
    line[i] = c;
    i++;
}

line[i] = '\0';
printf("The read line is: %s\n", line);
```

4.2.37

File Format: Binary Files

- any non-text file is a binary file
- actually, **any file is a binary file**
 - just a stream of bytes
- \Rightarrow **applications have to know**
 - formats of files they work with (e.g. pdf, docx etc.)
 - i.e. **the way the bytes must be grouped** for good interpretation
 - i.e. the offsets real information (interpretable groups of bytes) is placed

4.2.38

Example: Write / Read into / from Binary Files

- process 1

```
int fd;
int number = 10;
char c = 'A';

if ((fd = creat("intfile.bin", 0644)) < 0) {
    perror("Cannot create the file");
    exit(1);
}

// write a char on the first byte
write(fd, &c, sizeof(c));

// write an integer's representation on the next four bytes
write(fd, &number, sizeof(number));
```
- process 2

```
int fd;
int number;

if ((fd = open("intfile.bin", O_RDONLY)) < 0) {
    perror("Cannot open the file");
    exit(1);
}
```

```
// position where WE (MUST) KNOW the integer is
// i.e. one byte after beginning of file
lseek(fd, sizeof(char), SEEK_SET);

// read four bytes from crt position
// i.e. an integer's representation
read(fd, &number, sizeof(number));
```

4.2.39

Files With Holes (Gaps). Description

- **lseek is allowed to position after the end of file**
- **when write to that position**
 - **a gap results in file**
 - i.e. unwritten space
- Linux does not allocate physical space for gaps
- **read returns zeros from a gap**
 - there is no difference between reading previously written zeros or reading from a gap
 - it is the application's responsibility to remember where gaps are (manages them)
- usage
 - core dumps for crashed processes
 - virtual HDD files (dynamically allocated)

4.2.40

Files With Holes (Gaps). Example

- Example 1: Command line

```
dd if=/dev/zero of=file bs=1024 count=1 seek=1024
```
- Example 2: C Program

```
int fd;
// create a zero-sized file
fd = creat("file.with.gaps.bin", 0644);
if (fd < 0) {
    perror("Cannot create file");
    exit(1);
}

// position 4KiB after the end of file
lseek(fd, 4096, SEEK_END);
// has no effect: only modify the crt position

// write at crt position
write(fd, "END", 3);
// create a gap (unwritten bytes) of 4KB,
// followed by 3 written bytes
// file's size is now 4096 + 3 = 4099
```

4.2.41

File format (text, binary, holes etc.) is the business of user applications — OS is not involved!

4.2.42

Input/Output Redirection. Example

```
int fd1, fd2, n1, n2;

fd1 = open("input.txt", O_RDONLY);
fd2 = creat("output.txt", 0_WRONLY);

// reads an integer from STDIN
scanf("%d", &n1); // calls read(0, ...);

// redirects STDIN
close(0); // breaks the initial association between 0 and STDIN
dup(fd1); // associates 0 with the same open file like fd1

// reads an integer from STDIN
scanf("%d", &n2); // calls read(0, ...);
// actually reads from "input.txt"

// writes an integer to STDOUT
printf("%d\n", n1); // calls write(1, ...);

// redirects STDOUT
dup2(fd2, 1); // makes 1 a duplicate for fd2

// writes an integer to STDOUT
printf("%d\n", n2); // calls write(1, ...);
// actually writes to "output.txt"
```

4.2.43

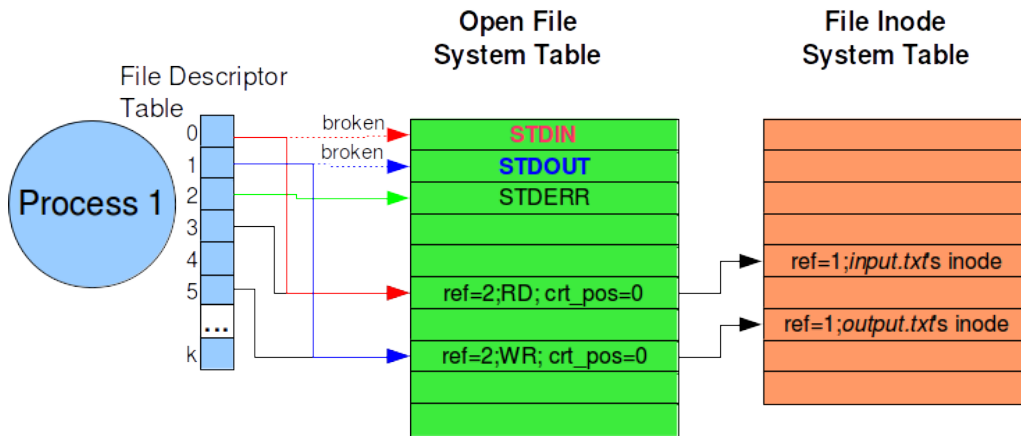


Figure 6: The effect of redirecting STDIN and STDOUT

Input/Output Redirection. Illustration on Open File Tables

4.2.44

Standard channel (STDIN, STDOUT, STDERR) redirection is possible because any resource in the system is modeled as a file!

4.2.45

Changing Permission Rights

- specify permission rights for all three groups of users
 - user: `rwX` → 111 (7)
 - group: `r-X` → 101 (5)
 - others: `--X` → 001 (1)
- example


```
chmod("file", 0751);
```

4.2.46

Getting Information (Metadata) About a File

- Linux metadata = **i-node (information node)**
- **each file and directory has its own, unique i-node**
- all i-nodes
 - have the same fixed size
 - placed together in one HDD area (inode area)
- ⇒ **i-node numbers** used as an index to identify an i-node
- i-node contents
 - file type, size, owner, group, permissions etc.

4.2.47

Example: Get a File's I-Node

```
int res;
struct stat inode;

// gets file's inode
res = lstat("/home/os/input.txt", &inode);
if (res < 0) {
    perror("Cannot get file inode");
}

// identify file's type
if (S_ISREG(inode.st_mode)) {
```



```

    printf("It is a file\n");
    printf("File's size [bytes]: %d\n", inode.st_size);
}

if (S_ISDIR(inode.st_mode))
    printf("It is a directory\n");

if (S_ISLNK(inode.st_mode))
    printf("It is a symbolic link\n");

```

4.2.48

Reading Directory Contents

- directory provided as a **collection of elements**
 - named **directory entries**
- the internal structure of directory (linked-list, B-tree) not visible
 - \Rightarrow the only way to **read a directory** is **entry by entry**
 - i.e. **sequential access**
- a directory entry contains (at least)
 - name (e.g. file or subdirectory)
 - i-node number (not of real interest)
- take care of “.” and “..”
 - “.” points to the current directory
 - “..” points to the parent directory
 - exists as real elements in a directory
 - they could induce cycles in applications that traverse a file tree

4.2.49

Reading Directory Contents

```

DIR* dir;
struct dirent *entry;
char path[MAX_PATH];
struct stat file_info;

// open directory
if ((dir = opendir("/home/os")) == NULL) {
    perror("Cannot open the directory");
    exit(1);
}

// read one-by-one dir entries until NULL returned
while ( (entry = readdir(dir)) != NULL) {
    // avoid "." and ".." as they are not useful
    if (strcmp(entry->d_name, ".") && strcmp(entry->d_name, "..")) {
        // build the complete path = dirpath + dirent's name
        sprintf(path, "%s/%s", "/home/os", entry->d_name);

        // get element's inode
        stat(path, &file_info);

        // identify type
        if (S_ISREG(file_info.st_mode))
            printf("%s is a file\n", path);
        else
            if (S_ISDIR(file_info.st_mode))
                printf("%s is a dir\n", path);
    }
}

```

4.2.50

Searching for an Element in a Directory

- naive, inefficient way (does not benefit from the specialized directory structure)

```

DIR* dir;
struct dirent *entry;
char path[MAX_PATH];
struct stat file_info;

// open directory
if ((dir = opendir("/home/os")) == NULL) {
    perror("Cannot open the directory");
    exit(1);
}

// read one-by-one dir entries until NULL returned
while ( (entry = readdir(dir)) != NULL) {
    if (strcmp(entry->d_name, SEARCHED_NAME) == 0) {
        // build the complete path = dirpath + dirent's name
        sprintf(path, "%s/%s", "/home/os", entry->d_name);
        printf("Found %s\n", path);
        break;
    }
}

```

- efficient way (benefit from the specialized directory structure)

```
char path[MAX_PATH];
struct stat file_info;

// build the complete path = dirpath + SEARCHED_NAME
sprintf(path, "%s/%s", "/home/os", SEARCHED_NAME);

// try getting file info
// succeeds if file exists, fails otherwise
if (stat(path, &file_info) > 0) {
    printf("Found %s\n", path);
} else {
    perror("Cannot get file's inode");
    exit(1);
}
```

4.2.51

A directory is presented and interacted with as a collection of elements! Searching, traversing, creating, changing could be performed only at directory element level.

4.2.52

Security Considerations: File-Path Traversal

- context
 - an application wants to **confine access** (of its users) **to a subdirectory** (subtree)
 - very common to Web applications
- problem
 - when user controls (specifies) parts of the file path
 - “..” **could be used to evade** the restricted directory
- solution
 - check for “..” in user-controlled file paths

4.2.53

File-Path Traversal Illustration

- vulnerable code

```
char filepath[MAX_PATH];

scanf("%s", filename); // <--- provided by the user (a possible attacker)!!!

snprintf(filepath, MAX_PATH, "/home/restricted_user/%s", filename);

fd = open(filepath, O_RDONLY);
... // display the file's contents
```

- malicious value for “filename”

```
"../../../../../../etc/passwd"
```

- secure code

```
char filepath[MAX_PATH];

scanf("%s", filename);
if (strstr(filename, "..") != NULL)
    return;
snprintf(filepath, MAX_PATH, "/home/restricted_user/%s", filename);

fd = open(filepath, O_RDONLY);
```

4.2.54

Never trust the users! Check for “..” in the given file paths, if want avoiding path traversal!

4.2.55

Manipulate Directories

```
BOOL WINAPI CreateDirectory(LPCTSTR lpPathName, LPSECURITY_ATTRIBUTES lpSecurityAttributes);

BOOL WINAPI CreateHardLink(LPCTSTR lpFileName, LPCTSTR lpExistingFileName,
    LPSECURITY_ATTRIBUTES lpSecurityAttributes);

BOOLEAN WINAPI CreateSymbolicLink(LPTSTR lpSymlinkFileName, LPTSTR lpTargetFileName, DWORD dwFlags);

BOOL WINAPI DeleteFile(LPCTSTR lpFileName);

HANDLE WINAPI FindFirstFile(LPCTSTR lpFileName, LPWIN32_FIND_DATA lpFindFileData);

BOOL WINAPI FindNextFile(HANDLE hFindFile, LPWIN32_FIND_DATA lpFindFileData);

BOOL WINAPI RemoveDirectory(LPCTSTR lpPathName);
```

4.2.60

Sparse Files

```
LARGE_INTEGER    FileSize;

FileSize.QuadPart = 8 * 1024 * 1024 * 1024;

FileHandle = CreateFile("file",
    GENERIC_READ | GENERIC_WRITE, FILE_SHARE_READ | FILE_SHARE_WRITE,
    NULL, CREATE_NEW, FILE_FLAG_NO_BUFFERING, NULL);

DeviceIoControl(FileHandle, FSCTL_SET_SPARSE, NULL, 0,
    NULL, 0, &BytesReturned, NULL);

SetFilePointerEx(FileHandle, FileSize, 0, FILE_BEGIN);

SetEndOfFile(FileHandle);
```

4.2.61

Alternate Data Streams. Command Line Examples

- Example 1

```
echo hello > file.txt:alternatestream.txt

more < file.txt:alternatestream.txt

notepad file.txt:alternatestream.txt
```

- Example 2

```
type c:\windows\system32\calc.exe > file.txt:calc.exe

start .\file.txt:calc.exe
```

- Getting alternate streams

```
http://www.microsoft.com/technet/sysinternals/default.mspx

streams [-s] [-d] <file\_name>
```

4.2.62

Alternate Data Streams. C Program Example

```
HANDLE inhandle, outhandle;
char buffer[BUF_SIZE];
int count, s;
DWORD ocnt;

inhandle = CreateFile("sursa.txt", GENERIC_READ, 0,
    NULL, OPEN_EXISTING, 0, NULL);

outhandle = CreateFile("dest.txt:file.txt", GENERIC_WRITE,
    0, NULL, CREATE_ALWAYS,
    FILE_ATTRIBUTE_NORMAL, NULL);

/* copy the file */
do {
    s = ReadFile(inhandle, buffer, BUF_SIZE, &count, NULL);

    if (s && count > 0)
        WriteFile(outhandle, buffer, count, &ocnt, NULL);
} while (s>0 && count>0);
```

4.2.63

3 Conclusions

What We Talked About

- permission rights
 - similar on both Linux and Windows
 - read (r), write (w), execute (x)
 - for both files and directories
- system calls and how they provide access to files and directories
 - creat, unlink, rename, truncate
 - open read, write, lseek, close
 - opendir, readdir, closedir
- security considerations
 - weak permission rights
 - path traversal

4.2.64

Lessons Learned

- the “file” is a sequence of bytes
 - each application should manage its own files’ format
- the “directory” is a collection of elements
 - traverse it element by element
 - search for elements
- do not trust users!
 - set strong permission rights!
 - check for path traversal elements!

4.2.65