

File System's Physical Data Layer

Implementation, Fragmentation, Links and Backup

Adrian Coleșa

Technical University of Cluj-Napoca (UTCN)
Computer Science Department

March 25, 2020

The purpose of today's lecture

- Presents details about the way files and directories are implemented.
- Presents related strategies and problems like: fragmentation, links, backup.

Bibliography

- A. Tanenbaum, *Modern Operating Systems*, 2nd Edition, 2001, Chapter 6, pg. 399 – 421

Outline

- 1 File's Data Allocation
- 2 Directory Implementation
- 3 Hard and Symbolic Links

Outline

- 1 File's Data Allocation
- 2 Directory Implementation
- 3 Hard and Symbolic Links

Context

- files are
 - provided to user application as sequences at bytes
 - a logical view
 - a contiguous area
 - **allocated in terms of blocks**, i.e. group of bytes, on HDD
 - a physical view
 - not necessarily a contiguous area
- we are interested in
 - how blocks of a file are allocated on HDD
 - how does the allocation strategy influence the user applications

Context

- files are
 - **provided to user application as sequences at bytes**
 - a logical view
 - a contiguous area
 - **allocated in terms of blocks**, i.e. group of bytes, on HDD
 - a physical view
 - not necessarily a contiguous area
- we are interested in
 - how blocks of a file are allocated on HDD
 - how does the allocation strategy influence the user applications

Context

- files are
 - **provided to user application as sequences at bytes**
 - a logical view
 - a contiguous area
 - **allocated in terms of blocks**, i.e. group of bytes, on HDD
 - a physical view
 - not necessarily a contiguous area
- we are interested in
 - how blocks of a file are allocated on HDD
 - how does the allocation strategy influence the user applications

Context

- files are
 - **provided to user application as sequences at bytes**
 - a logical view
 - a contiguous area
 - **allocated in terms of blocks**, i.e. group of bytes, on HDD
 - a physical view
 - not necessarily a contiguous area
- we are interested in
 - how blocks of a file are allocated on HDD
 - how does the allocation strategy influence the user applications

Context

- files are
 - **provided to user application as sequences at bytes**
 - a logical view
 - a contiguous area
 - **allocated in terms of blocks**, i.e. group of bytes, on HDD
 - a physical view
 - not necessarily a contiguous area
- we are interested in
 - how blocks of a file are allocated on HDD
 - how does the allocation strategy influence the user applications

Context

- files are
 - **provided to user application as sequences at bytes**
 - a logical view
 - a contiguous area
 - **allocated in terms of blocks**, i.e. group of bytes, on HDD
 - a physical view
 - not necessarily a contiguous area
- we are interested in
 - how blocks of a file are allocated on HDD
 - how does the allocation strategy influence the user applications

Context

- files are
 - **provided to user application as sequences at bytes**
 - a logical view
 - a contiguous area
 - **allocated in terms of blocks**, i.e. group of bytes, on HDD
 - a physical view
 - not necessarily a contiguous area
- we are interested in
 - how blocks of a file are allocated on HDD
 - how does the allocation strategy influence the user applications

Context

- files are
 - **provided to user application as sequences at bytes**
 - a logical view
 - a contiguous area
 - **allocated in terms of blocks**, i.e. group of bytes, on HDD
 - a physical view
 - not necessarily a contiguous area
- we are interested in
 - how blocks of a file are allocated on HDD
 - how does the allocation strategy influence the user applications

Context

- files are
 - **provided to user application as sequences at bytes**
 - a logical view
 - a contiguous area
 - **allocated in terms of blocks**, i.e. group of bytes, on HDD
 - a physical view
 - not necessarily a contiguous area
- we are interested in
 - how blocks of a file are allocated on HDD
 - how does the allocation strategy influence the user applications

Context

- files are
 - **provided to user application as sequences at bytes**
 - a logical view
 - a contiguous area
 - **allocated in terms of blocks**, i.e. group of bytes, on HDD
 - a physical view
 - not necessarily a contiguous area
- we are interested in
 - how blocks of a file are allocated on HDD
 - how does the allocation strategy influence the user applications

Contiguous Allocation

- Files are allocated in **just one contiguous area**




Figure: HDD Partition Structure

Contiguous Allocation

- Files are allocated in **just one contiguous area**




Figure: HDD Partition Structure

Contiguous Allocation

- Files are allocated in **just one contiguous area**



Figure: HDD Partition Structure

Contiguous Allocation

- Files are allocated in **just one contiguous area**



Figure: HDD Partition Structure

Contiguous Allocation

- Files are allocated in **just one contiguous area**




Figure: HDD Partition Structure

Contiguous Allocation

- Files are allocated in **just one contiguous area**




Figure: HDD Partition Structure

Contiguous Allocation

- Files are allocated in **just one contiguous area**




Figure: HDD Partition Structure

Contiguous Allocation

- Files are allocated in **just one contiguous area**




Figure: HDD Partition Structure

Contiguous Allocation

- Advantages
 - Reading large areas from the file is very fast
 - Keeping track of allocated blocks (BAT) is very simple: starting block and the number of allocated blocks




Figure: HDD Partition Structure

Contiguous Allocation

Disadvantages

- Difficult to increase the file size: see for example **File 4**
- Leads to **external fragmentation**
- Complex allocation strategies: first fit, best fit etc.



Figure: HDD Partition Structure

Contiguous Allocation

- Disadvantages

- Difficult to increase the file size: see for example **File 4**
- Leads to **external fragmentation**
- Complex allocation strategies: first fit, best fit etc.




Figure: HDD Partition Structure

Contiguous Allocation

- Disadvantages
 - Difficult to increase the file size: see for example **File 4**
 - Leads to **external fragmentation**
 - Complex allocation strategies: first fit, best fit etc.




Figure: HDD Partition Structure

Contiguous Allocation

- Disadvantages
 - Difficult to increase the file size: see for example **File 4**
 - Leads to **external fragmentation**
 - Complex allocation strategies: first fit, best fit etc.




Figure: HDD Partition Structure

Contiguous Allocation

- Disadvantages
 - Difficult to increase the file size: see for example **File 4**
 - Leads to **external fragmentation**
 - Complex allocation strategies: first fit, best fit etc.




Figure: HDD Partition Structure

Contiguous Allocation

- Disadvantages
 - Difficult to increase the file size: see for example **File 4**
 - Leads to **external fragmentation**
 - Complex allocation strategies: first fit, best fit etc.




Figure: HDD Partition Structure

Contiguous Allocation

Disadvantages

- Difficult to increase the file size: see for example **File 4**
- Leads to **external fragmentation**
- Complex allocation strategies: first fit, best fit etc.




Figure: HDD Partition Structure

Contiguous Allocation

- Disadvantages
 - Difficult to increase the file size: see for example **File 4**
 - Leads to **external fragmentation**
 - Complex allocation strategies: **first fit**, best fit etc.




Figure: HDD Partition Structure

Contiguous Allocation

- Disadvantages

- Difficult to increase the file size: see for example **File 4**
- Leads to **external fragmentation**
- Complex allocation strategies: first fit, **best fit** etc.




Figure: HDD Partition Structure

Any-Free-Block Allocation

- The file can be allocated any free block
- Advantages
 - there is no external fragmentation; any free block can be used
 - file size can be easily extended
 - could be combined with contiguous allocation: the file is allocated more contiguous areas (as large as possible)
- Disadvantages
 - data access (e.g. read entire file) not so efficient
 - BAT structure more complicated
 - still suffers from internal fragmentation and data fragmentation

Any-Free-Block Allocation

- The file can be allocated any free block
- Advantages
 - there is no external fragmentation; any free block can be used
 - file size can be easily extended
 - could be combined with contiguous allocation: the file is allocated more contiguous areas (as large as possible)
- Disadvantages
 - data access (e.g. read entire file) not so efficient
 - BAT structure more complicated
 - still suffers from internal fragmentation and data fragmentation

Any-Free-Block Allocation

- The file can be allocated any free block
- Advantages
 - there is no external fragmentation; any free block can be used
 - file size can be easily extended
 - could be combined with contiguous allocation: the file is allocated more contiguous areas (as large as possible)
- Disadvantages
 - data access (e.g. read entire file) not so efficient
 - BAT structure more complicated
 - still suffers from internal fragmentation and data fragmentation

Any-Free-Block Allocation

- The file can be allocated any free block
- Advantages
 - there is no external fragmentation; any free block can be used
 - file size can be easily extended
 - could be combined with contiguous allocation: the file is allocated more contiguous areas (as large as possible)
- Disadvantages
 - data access (e.g. read entire file) not so efficient
 - BAT structure more complicated
 - still suffers from internal fragmentation and data fragmentation

Any-Free-Block Allocation

- The file can be allocated any free block
- Advantages
 - there is no external fragmentation; any free block can be used
 - file size can be easily extended
 - could be combined with contiguous allocation: the file is allocated more contiguous areas (as large as possible)
- Disadvantages
 - data access (e.g. read entire file) not so efficient
 - BAT structure more complicated
 - still suffers from internal fragmentation and data fragmentation

Any-Free-Block Allocation

- The file can be allocated any free block
- Advantages
 - there is no external fragmentation; any free block can be used
 - file size can be easily extended
 - could be combined with contiguous allocation: the file is allocated more contiguous areas (as large as possible)
- Disadvantages
 - data access (e.g. read entire file) not so efficient
 - BAT structure more complicated
 - still suffers from **internal fragmentation** and **data fragmentation**

Any-Free-Block Allocation

- The file can be allocated any free block
- Advantages
 - there is no external fragmentation; any free block can be used
 - file size can be easily extended
 - could be combined with contiguous allocation: the file is allocated more contiguous areas (as large as possible)
- Disadvantages
 - data access (e.g. read entire file) not so efficient
 - BAT structure more complicated
 - still suffers from **internal fragmentation** and **data fragmentation**

Any-Free-Block Allocation

- The file can be allocated any free block
- Advantages
 - there is no external fragmentation; any free block can be used
 - file size can be easily extended
 - could be combined with contiguous allocation: the file is allocated more contiguous areas (as large as possible)
- Disadvantages
 - data access (e.g. read entire file) not so efficient
 - BAT structure more complicated
 - still suffers from **internal fragmentation** and **data fragmentation**

Any-Free-Block Allocation

- The file can be allocated any free block
- Advantages
 - there is no external fragmentation; any free block can be used
 - file size can be easily extended
 - could be combined with contiguous allocation: the file is allocated more contiguous areas (as large as possible)
- Disadvantages
 - data access (e.g. read entire file) not so efficient
 - BAT structure more complicated
 - still suffers from **internal fragmentation** and **data fragmentation**

Fragmentation Types: External Fragmentation

- context
 - free space scattered in small areas over the entire HDD
 - alternating with allocated areas
- problem
 - free space cannot be used (in some situations)
 - example
 - need to allocate a contiguous area of a given size
 - even though free fragments are found, they could be available only in a smaller size
 - even if total free space would be enough
- specific to
 - contiguous allocation strategies
 - where data can be allocated only in a single contiguous area
- solution (inefficient, i.e. time consuming)
 - defragmentation: move all allocated space at one end of the HDD
 - ⇒ free space in a single contiguous area at the other end of the HDD

Fragmentation Types: External Fragmentation

- context
 - **free space scattered in small areas** over the entire HDD
 - alternating with allocated areas
- problem
 - free space cannot be used (in some situations)
 - example
 - need to allocate a contiguous area of a given size
 - there is no free contiguous area that could be available
 - even if total free space would be enough
- specific to
 - contiguous allocation strategies
 - where data can be allocated **only in a single contiguous area**
- solution (inefficient, i.e. time consuming)
 - **defragmentation**: move all allocated space at one end of the HDD
 - \Rightarrow free space in a single contiguous area at the other end of the HDD

Fragmentation Types: External Fragmentation

- context
 - **free space scattered in small areas** over the entire HDD
 - alternating with allocated areas
- problem
 - free space cannot be used (in some situations)
 - example
 - need to allocate a contiguous area of a given size
 - there is free contiguous area could be available
 - however if total free space would be enough
- specific to
 - contiguous allocation strategies
 - where data can be allocated **only in a single contiguous area**
- solution (inefficient, i.e. time consuming)
 - defragmentation: move all allocated space at one end of the HDD
 - \Rightarrow free space in a single contiguous area at the other end of the HDD

Fragmentation Types: External Fragmentation

- context
 - **free space scattered in small areas** over the entire HDD
 - alternating with allocated areas
- problem
 - **free space cannot be used** (in some situations)
 - example
 - need to allocate a contiguous area of a give size
 - but no free contiguous area could be available
 - even if total free space would be enough
- specific to
 - contiguous allocation strategies
 - where data can be allocated **only in a single contiguous area**
- solution (inefficient, i.e. time consuming)
 - **defragmentation**: move all allocated space at one end of the HDD
 - \Rightarrow free space in a single contiguous area at the other end of the HDD

Fragmentation Types: External Fragmentation

- context
 - **free space scattered in small areas** over the entire HDD
 - alternating with allocated areas
- problem
 - **free space cannot be used** (in some situations)
 - example
 - need to allocate a contiguous area of a give size
 - but no free contiguous area could be available
 - even if total free space would be enough
- specific to
 - contiguous allocation strategies
 - where data can be allocated **only in a single contiguous area**
- solution (inefficient, i.e. time consuming)
 - **defragmentation**: move all allocated space at one end of the HDD
 - \Rightarrow free space in a single contiguous area at the other end of the HDD

Fragmentation Types: External Fragmentation

- context
 - **free space scattered in small areas** over the entire HDD
 - alternating with allocated areas
- problem
 - **free space cannot be used** (in some situations)
 - example
 - need to allocate a contiguous area of a give size
 - but no free contiguous area could be available
 - even if total free space would be enough
- specific to
 - contiguous allocation strategies
 - where data can be allocated **only in a single contiguous area**
- solution (inefficient, i.e. time consuming)
 - **defragmentation**: move all allocated space at one end of the HDD
 - \Rightarrow free space in a single contiguous area at the other end of the HDD

Fragmentation Types: External Fragmentation

- context
 - **free space scattered in small areas** over the entire HDD
 - alternating with allocated areas
- problem
 - **free space cannot be used** (in some situations)
 - example
 - need to allocate a contiguous area of a give size
 - but no free contiguous area could be available
 - even if total free space would be enough
- specific to
 - contiguous allocation strategies
 - where data can be allocated **only in a single contiguous area**
- solution (inefficient, i.e. time consuming)
 - **defragmentation**: move all allocated space at one end of the HDD
 - \Rightarrow free space in a single contiguous area at the other end of the HDD

Fragmentation Types: External Fragmentation

- context
 - **free space scattered in small areas** over the entire HDD
 - alternating with allocated areas
- problem
 - **free space cannot be used** (in some situations)
 - example
 - need to allocate a contiguous area of a give size
 - but no free contiguous area could be available
 - even if total free space would be enough
- specific to
 - contiguous allocation strategies
 - where data can be allocated **only in a single contiguous area**
- solution (inefficient, i.e. time consuming)
 - **defragmentation**: move all allocated space at one end of the HDD
 - \Rightarrow free space in a single contiguous area at the other end of the HDD

Fragmentation Types: External Fragmentation

- context
 - **free space scattered in small areas** over the entire HDD
 - alternating with allocated areas
- problem
 - **free space cannot be used** (in some situations)
 - example
 - need to allocate a contiguous area of a give size
 - but no free contiguous area could be available
 - even if total free space would be enough
- specific to
 - contiguous allocation strategies
 - where data can be allocated **only in a single contiguous area**
- solution (inefficient, i.e. time consuming)
 - **defragmentation**: move all allocated space at one end of the HDD
 - \Rightarrow free space in a single contiguous area at the other end of the HDD

Fragmentation Types: External Fragmentation

- context
 - **free space scattered in small areas** over the entire HDD
 - alternating with allocated areas
- problem
 - **free space cannot be used** (in some situations)
 - example
 - need to allocate a contiguous area of a give size
 - but no free contiguous area could be available
 - even if total free space would be enough
- specific to
 - contiguous allocation strategies
 - where data can be allocated **only in a single contiguous area**
- solution (inefficient, i.e. time consuming)
 - defragmentation: move all allocated space at one end of the HDD
 - \Rightarrow free space in a single contiguous area at the other end of the HDD

Fragmentation Types: External Fragmentation

- context
 - **free space scattered in small areas** over the entire HDD
 - alternating with allocated areas
- problem
 - **free space cannot be used** (in some situations)
 - example
 - need to allocate a contiguous area of a give size
 - but no free contiguous area could be available
 - even if total free space would be enough
- specific to
 - contiguous allocation strategies
 - where data can be allocated **only in a single contiguous area**
- solution (inefficient, i.e. time consuming)
 - defragmentation: move all allocated space at one end of the HDD
 - \Rightarrow free space in a single contiguous area at the other end of the HDD

Fragmentation Types: External Fragmentation

- context
 - **free space scattered in small areas** over the entire HDD
 - alternating with allocated areas
- problem
 - **free space cannot be used** (in some situations)
 - example
 - need to allocate a contiguous area of a give size
 - but no free contiguous area could be available
 - even if total free space would be enough
- specific to
 - contiguous allocation strategies
 - where data can be allocated **only in a single contiguous area**
- solution (inefficient, i.e. time consuming)
 - defragmentation: move all allocated space at one end of the HDD
 - \Rightarrow free space in a single contiguous area at the other end of the HDD

Fragmentation Types: External Fragmentation

- context
 - **free space scattered in small areas** over the entire HDD
 - alternating with allocated areas
- problem
 - **free space cannot be used** (in some situations)
 - example
 - need to allocate a contiguous area of a give size
 - but no free contiguous area could be available
 - even if total free space would be enough
- specific to
 - contiguous allocation strategies
 - where data can be allocated **only in a single contiguous area**
- solution (inefficient, i.e. time consuming)
 - **defragmentation**: move all allocated space at one end of the HDD
 - \Rightarrow free space in a single contiguous area at the other end of HDD

Fragmentation Types: External Fragmentation

- context
 - **free space scattered in small areas** over the entire HDD
 - alternating with allocated areas
- problem
 - **free space cannot be used** (in some situations)
 - example
 - need to allocate a contiguous area of a give size
 - but no free contiguous area could be available
 - even if total free space would be enough
- specific to
 - contiguous allocation strategies
 - where data can be allocated **only in a single contiguous area**
- solution (inefficient, i.e. time consuming)
 - **defragmentation**: move all allocated space at one end of the HDD
 - \Rightarrow free space in a single contiguous area at the other end of the HDD

Fragmentation Types: External Fragmentation

- context
 - **free space scattered in small areas** over the entire HDD
 - alternating with allocated areas
- problem
 - **free space cannot be used** (in some situations)
 - example
 - need to allocate a contiguous area of a give size
 - but no free contiguous area could be available
 - even if total free space would be enough
- specific to
 - contiguous allocation strategies
 - where data can be allocated **only in a single contiguous area**
- solution (inefficient, i.e. time consuming)
 - **defragmentation**: move all allocated space at one end of the HDD
 - \Rightarrow free space in a single contiguous area at the other end of HDD

Fragmentation Types: Internal Fragmentation

- context
 - any free block could be allocated when new space needed
 - BUT ... **allocation is done in terms of predefined units**, i.e. blocks
 - AND ... **needed space not a multiple of block size**
- problem
 - some (internal) "free" space cannot be used
 - **unused space in blocks allocated to files is lost**
- specific to
 - allocation strategies that allocate data in terms of fixed-size blocks
- solutions
 - **smaller block size** to reduce internal fragmentation
 - **fragment blocks**, i.e. share the same block for tails of multiple files

Fragmentation Types: Internal Fragmentation

- context
 - any free block could be allocated when new space needed
 - BUT ... **allocation is done in terms of predefined units**, i.e. blocks
 - AND ... **needed space not a multiple of block size**
- problem
 - some (internal) "free" space cannot be used
 - **unused space in blocks allocated to files is lost**
- specific to
 - allocation strategies that allocate data in terms of fixed-size blocks
- solutions
 - **smaller block size** to reduce internal fragmentation
 - **fragment blocks**, i.e. share the same block for tails of multiple files

Fragmentation Types: Internal Fragmentation

- context
 - any free block could be allocated when new space needed
 - BUT ... **allocation is done in terms of predefined units**, i.e. blocks
 - AND ... **needed space not a multiple of block size**
- problem
 - some (internal) "free" space cannot be used
 - unused space in blocks allocated to files is lost
- specific to
 - allocation strategies that allocate data in terms of fixed-size blocks
- solutions
 - smaller block size to reduce internal fragmentation
 - fragment blocks, i.e. share the same block for tails of multiple files

Fragmentation Types: Internal Fragmentation

- context
 - any free block could be allocated when new space needed
 - BUT ... **allocation is done in terms of predefined units**, i.e. blocks
 - AND ... **needed space not a multiple of block size**
- problem
 - some (internal) "free" space cannot be used
 - unused space in blocks allocated to files is lost
- specific to
 - allocation strategies that allocate data in terms of fixed-size blocks
- solutions
 - smaller block size to reduce internal fragmentation
 - fragment blocks, i.e. share the same block for tails of multiple files

Fragmentation Types: Internal Fragmentation

- context
 - any free block could be allocated when new space needed
 - BUT ... **allocation is done in terms of predefined units**, i.e. blocks
 - AND ... **needed space not a multiple of block size**
- problem
 - some (internal) “free” space cannot be used
 - **unused space in blocks allocated to files** is lost
- specific to
 - allocation strategies that allocate data in terms of fixed-size blocks
- solutions
 - smaller block size to reduce internal fragmentation
 - **fragment blocks**, i.e. share the same block for tails of multiple files

Fragmentation Types: Internal Fragmentation

- context
 - any free block could be allocated when new space needed
 - BUT ... **allocation is done in terms of predefined units**, i.e. blocks
 - AND ... **needed space not a multiple of block size**
- problem
 - some (internal) “free” space cannot be used
 - **unused space in blocks allocated to files** is lost
- specific to
 - allocation strategies that allocate data in terms of fixed-size blocks
- solutions
 - smaller block size to reduce internal fragmentation
 - **fragment blocks**, i.e. share the same block for tails of multiple files

Fragmentation Types: Internal Fragmentation

- context
 - any free block could be allocated when new space needed
 - BUT ... **allocation is done in terms of predefined units**, i.e. blocks
 - AND ... **needed space not a multiple of block size**
- problem
 - some (internal) “free” space cannot be used
 - **unused space in blocks allocated to files** is lost
- specific to
 - allocation strategies that allocate data in terms of fixed-size blocks
- solutions
 - smaller block size to reduce internal fragmentation
 - fragment blocks, i.e. share the same block for tails of multiple files

Fragmentation Types: Internal Fragmentation

- context
 - any free block could be allocated when new space needed
 - BUT ... **allocation is done in terms of predefined units**, i.e. blocks
 - AND ... **needed space not a multiple of block size**
- problem
 - some (internal) “free” space cannot be used
 - **unused space in blocks allocated to files** is lost
- specific to
 - allocation strategies that allocate data in terms of fixed-size blocks
- solutions
 - smaller block size to reduce internal fragmentation
 - fragment blocks, i.e. share the same block for tails of multiple files

Fragmentation Types: Internal Fragmentation

- context
 - any free block could be allocated when new space needed
 - BUT ... **allocation is done in terms of predefined units**, i.e. blocks
 - AND ... **needed space not a multiple of block size**
- problem
 - some (internal) “free” space cannot be used
 - **unused space in blocks allocated to files** is lost
- specific to
 - allocation strategies that allocate data in terms of fixed-size blocks
- solutions
 - smaller block size to reduce internal fragmentation
 - fragment blocks, i.e. share the same block for tails of multiple files

Fragmentation Types: Internal Fragmentation

- context
 - any free block could be allocated when new space needed
 - BUT ... **allocation is done in terms of predefined units**, i.e. blocks
 - AND ... **needed space not a multiple of block size**
- problem
 - some (internal) “free” space cannot be used
 - **unused space in blocks allocated to files** is lost
- specific to
 - allocation strategies that allocate data in terms of fixed-size blocks
- solutions
 - **smaller block size** to reduce internal fragmentation
 - **fragment blocks**, i.e. share the same block for tails of multiple files

Fragmentation Types: Internal Fragmentation

- context
 - any free block could be allocated when new space needed
 - BUT ... **allocation is done in terms of predefined units**, i.e. blocks
 - AND ... **needed space not a multiple of block size**
- problem
 - some (internal) “free” space cannot be used
 - **unused space in blocks allocated to files** is lost
- specific to
 - allocation strategies that allocate data in terms of fixed-size blocks
- solutions
 - **smaller block size** to reduce internal fragmentation
 - **fragment blocks**, i.e. share the same block for tails of multiple files

Fragmentation Types: Internal Fragmentation

- context
 - any free block could be allocated when new space needed
 - BUT ... **allocation is done in terms of predefined units**, i.e. blocks
 - AND ... **needed space not a multiple of block size**
- problem
 - some (internal) “free” space cannot be used
 - **unused space in blocks allocated to files** is lost
- specific to
 - allocation strategies that allocate data in terms of fixed-size blocks
- solutions
 - **smaller block size** to reduce internal fragmentation
 - **fragment blocks**, i.e. share the same block for tails of multiple files

Fragmentation Types: Data Fragmentation

- context
 - the allocated space of a file is distributed in different non-consecutive blocks
 - i.e. more contiguous areas
- problem
 - file access could suffer performance penalties
- specific to
 - any-free-block allocation strategy
 - that do not impose the file to be in a single contiguous area
- solution
 - defragmentation: reallocate file's blocks in consecutive ones

Fragmentation Types: Data Fragmentation

- context
 - the allocated space of a file is distributed in different non-consecutive blocks
 - i.e. more contiguous areas
- problem
 - file access could suffer performance penalties
- specific to
 - any-free-block allocation strategy
 - that do not impose the file to be in a single contiguous area
- solution
 - defragmentation: reallocate file's blocks in consecutive ones

Fragmentation Types: Data Fragmentation

- context
 - the allocated space of a file is distributed in different non-consecutive blocks
 - i.e. more contiguous areas
- problem
 - file access could suffer performance penalties
- specific to
 - any-free-block allocation strategy
 - that do not impose the file to be in a single contiguous area
- solution
 - defragmentation: reallocate file's blocks in consecutive ones

Fragmentation Types: Data Fragmentation

- context
 - the allocated space of a file is distributed in different non-consecutive blocks
 - i.e. more contiguous areas
- problem
 - file access could suffer performance penalties
- specific to
 - any-free-block allocation strategy
 - that do not impose the file to be in a single contiguous area
- solution
 - defragmentation: reallocate file's blocks in consecutive ones

Fragmentation Types: Data Fragmentation

- context
 - the allocated space of a file is distributed in different non-consecutive blocks
 - i.e. more contiguous areas
- problem
 - **file access could suffer performance penalties**
- specific to
 - any-free-block allocation strategy
 - that do not impose the file to be in a single contiguous area
- solution
 - defragmentation: reallocate file's blocks in consecutive ones

Fragmentation Types: Data Fragmentation

- context
 - the allocated space of a file is distributed in different non-consecutive blocks
 - i.e. more contiguous areas
- problem
 - **file access could suffer performance penalties**
- specific to
 - any-free-block allocation strategy
 - that do not impose the file to be in a single contiguous area
- solution
 - defragmentation: reallocate file's blocks in consecutive ones

Fragmentation Types: Data Fragmentation

- context
 - the allocated space of a file is distributed in different non-consecutive blocks
 - i.e. more contiguous areas
- problem
 - **file access could suffer performance penalties**
- specific to
 - **any-free-block allocation strategy**
 - that do not impose the file to be in a single contiguous area
- solution
 - **defragmentation:** reallocate file's blocks in consecutive ones

Fragmentation Types: Data Fragmentation

- context
 - the allocated space of a file is distributed in different non-consecutive blocks
 - i.e. more contiguous areas
- problem
 - **file access could suffer performance penalties**
- specific to
 - **any-free-block allocation strategy**
 - that do not impose the file to be in a single contiguous area
- solution
 - defragmentation: reallocate file's blocks in consecutive ones

Fragmentation Types: Data Fragmentation

- context
 - the allocated space of a file is distributed in different non-consecutive blocks
 - i.e. more contiguous areas
- problem
 - **file access could suffer performance penalties**
- specific to
 - **any-free-block allocation strategy**
 - that do not impose the file to be in a single contiguous area
- solution
 - **defragmentation:** reallocate file's blocks in consecutive ones

Fragmentation Types: Data Fragmentation

- context
 - the allocated space of a file is distributed in different non-consecutive blocks
 - i.e. more contiguous areas
- problem
 - **file access could suffer performance penalties**
- specific to
 - **any-free-block allocation strategy**
 - that do not impose the file to be in a single contiguous area
- solution
 - **defragmentation**: reallocate file's blocks in consecutive ones

File System Block Size

- The problem: How large the block should be?
- Alternatives
 - Larger
 - more efficient read of file data
 - less internal fragmentation (waste HDD space)
 - Smaller
 - less internal fragmentation
 - less data fragmentation and data access time (many disk access)
- no good-for-all solution
 - *performance* and *space utilization* are inherently in conflict
 - the block size should be chosen knowing the way and for what the HDD partition will be used
 - a compromise should be chosen in a general usage case

File System Block Size

- The problem: How large the block should be?
- Alternatives
 - Larger
 - efficient read of file data
 - increase internal fragmentation (waste HDD space)
 - Smaller
 - reduce internal fragmentation
 - increase data fragmentation and data access time (many disk accesses)
- **no good-for-all solution**
 - *performance* and *space utilization* are inherently in conflict
 - the block size should be chosen knowing the way and for what the HDD partition will be used
 - a compromise should be chosen in a general usage case

File System Block Size

- The problem: How large the block should be?
- Alternatives
 - Larger
 - efficient read of file data
 - increase internal fragmentation (waste HDD space)
 - Smaller
 - reduce internal fragmentation
 - increase data fragmentation and data access time (many disk accesses)
- **no good-for-all solution**
 - *performance* and *space utilization* are inherently in conflict
 - the block size should be chosen knowing the way and for what the HDD partition will be used
 - a compromise should be chosen in a general usage case

File System Block Size

- The problem: How large the block should be?
- Alternatives
 - Larger
 - efficient read of file data
 - increase internal fragmentation (waste HDD space)
 - Smaller
 - reduce internal fragmentation
 - increase data fragmentation and data access time (many disk accesses)
- **no good-for-all solution**
 - *performance* and *space utilization* are inherently in conflict
 - the block size should be chosen knowing the way and for what the HDD partition will be used
 - a compromise should be chosen in a general usage case

File System Block Size

- The problem: How large the block should be?
- Alternatives
 - Larger
 - efficient read of file data
 - increase internal fragmentation (waste HDD space)
 - Smaller
 - reduce internal fragmentation
 - increase data fragmentation and data access time (many disk accesses)
- **no good-for-all solution**
 - *performance and space utilization* are inherently in conflict
 - the block size should be chosen knowing the way and for what the HDD partition will be used
 - a compromise should be chosen in a general usage case

File System Block Size

- The problem: How large the block should be?
- Alternatives
 - Larger
 - efficient read of file data
 - increase internal fragmentation (waste HDD space)
 - Smaller
 - reduce internal fragmentation
 - increase data fragmentation and data access time (many disk accesses)
- **no good-for-all solution**
 - *performance and space utilization* are inherently in conflict
 - the block size should be chosen knowing the way and for what the HDD partition will be used
 - a compromise should be chosen in a general usage case

File System Block Size

- The problem: How large the block should be?
- Alternatives
 - Larger
 - efficient read of file data
 - increase internal fragmentation (waste HDD space)
 - Smaller
 - reduce internal fragmentation
 - increase data fragmentation and data access time (many disk accesses)
- **no good-for-all solution**
 - *performance and space utilization* are inherently in conflict
 - the block size should be chosen knowing the way and for what the HDD partition will be used
 - a compromise should be chosen in a general usage case

File System Block Size

- The problem: How large the block should be?
- Alternatives
 - Larger
 - efficient read of file data
 - increase internal fragmentation (waste HDD space)
 - Smaller
 - reduce internal fragmentation
 - increase data fragmentation and data access time (many disk accesses)
- **no good-for-all solution**
 - *performance and space utilization* are inherently in conflict
 - the block size should be chosen knowing the way and for what the HDD partition will be used
 - a compromise should be chosen in a general usage case

File System Block Size

- The problem: How large the block should be?
- Alternatives
 - Larger
 - efficient read of file data
 - increase internal fragmentation (waste HDD space)
 - Smaller
 - reduce internal fragmentation
 - increase data fragmentation and data access time (many disk accesses)
- **no good-for-all solution**
 - *performance* and *space utilization* are inherently in conflict
 - the block size should be chosen knowing the way and for what the HDD partition will be used
 - a compromise should be chosen in a general usage case

File System Block Size

- The problem: How large the block should be?
- Alternatives
 - Larger
 - efficient read of file data
 - increase internal fragmentation (waste HDD space)
 - Smaller
 - reduce internal fragmentation
 - increase data fragmentation and data access time (many disk accesses)
- **no good-for-all solution**
 - *performance* and *space utilization* are inherently in conflict
 - the block size should be chosen knowing the way and for what the HDD partition will be used
 - a compromise should be chosen in a general usage case

File System Block Size

- The problem: How large the block should be?
- Alternatives
 - Larger
 - efficient read of file data
 - increase internal fragmentation (waste HDD space)
 - Smaller
 - reduce internal fragmentation
 - increase data fragmentation and data access time (many disk accesses)
- **no good-for-all solution**
 - *performance* and *space utilization* are inherently in conflict
 - the block size should be chosen knowing the way and for what the HDD partition will be used
 - a compromise should be chosen in a general usage case

File System Block Size

- The problem: How large the block should be?
- Alternatives
 - Larger
 - efficient read of file data
 - increase internal fragmentation (waste HDD space)
 - Smaller
 - reduce internal fragmentation
 - increase data fragmentation and data access time (many disk accesses)
- **no good-for-all solution**
 - *performance* and *space utilization* are inherently in conflict
 - the block size should be chosen knowing the way and for what the HDD partition will be used
 - a compromise should be chosen in a general usage case

Outline

- 1 File's Data Allocation
- 2 Directory Implementation
- 3 Hard and Symbolic Links

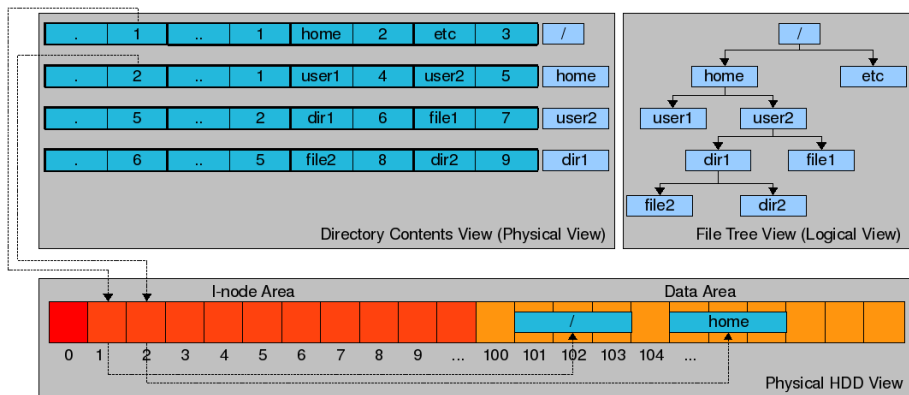
Directory Contents

- a system “file”
- stored as a stream of bytes, but interpreted by the OS
- organized as a collection of records (elements), called **directory entries**
- a directory entry contains
 - the (file, directory etc.) name
 - file's metadata or a reference to them

File “I-node” (Record)

- a physical space (element) and a corresponding data structure used to store information about a FS element (file, directory etc.)
- stores file meta-data, like
 - file type
 - size
 - owner
 - permission rights
 - time stamps
 - the BAT (Block Addresses Table)
 - etc.

The Relationship Between The Directory Entry and The I-node: Illustration



Outline

- 1 File's Data Allocation
- 2 Directory Implementation
- 3 Hard and Symbolic Links**

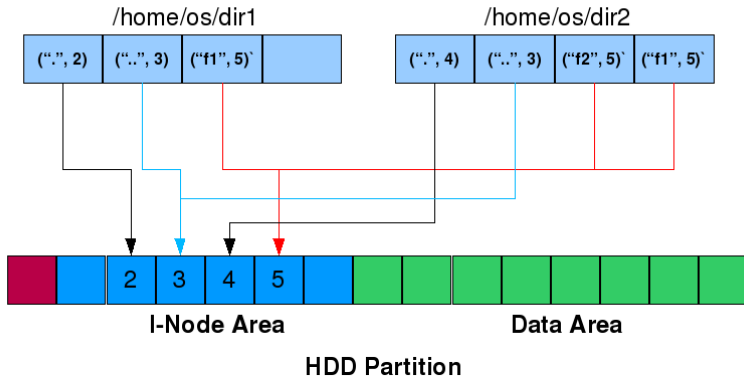
Sharing Data Between Directories

- make a file (directory) appear in different directories
- the operation is called linking files
- two kinds of links
 - hard (physical)
 - soft (symbolic)

Hard Link: Creation and Usage

```
create("/home/os/dir1/f1", 0600);    // only allocate i-node; no data
link("/home/os/dir1/f1", "/home/os/dir2/f2"); // hard link
link("/home/os/dir1/f1", "/home/os/dir2/f1"); // hard link
open("/home/os/dir1/f1", ...); // open file with i-node 5
open("/home/os/dir2/f2", ...); // open file with i-node 5
open("/home/os/dir2/f1", ...); // open file with i-node 5
stat("/home/os/dir1/f1", ...); // read i-node 5 contents
stat("/home/os/dir2/f1", ...); // read i-node 5 contents
stat("/home/os/dir2/f2", ...); // read i-node 5 contents
```

Hard Link: Illustration



Hard Link: Discussion

- Advantages

- a file really belong to the two or more directories, when a path is removed the physical file (space) is not removed until all hard links are removed
- very transparent; there is no difference and distinction between different hard links to the same file

- Disadvantages

- cannot be established between different partitions

Symbolic Link: Creation and Usage

```
create("/home/os/dir1/f1", 0600);    // only allocate i-node; no data
symlink("/home/os/dir1/f1", "/home/os/dir2/f2"); // hard link
open("/home/os/dir1/f1", ...); // open file with i-node 5
open("/home/os/dir2/f2", ...); // open file with i-node 5
stat("/home/os/dir1/f1", ...); // read i-node 5 contents
stat("/home/os/dir2/f2", ...); // read i-node 5 contents
lstat("/home/os/dir2/f2", ...); // read i-node 6 contents
```

Symbolic Link: Illustration

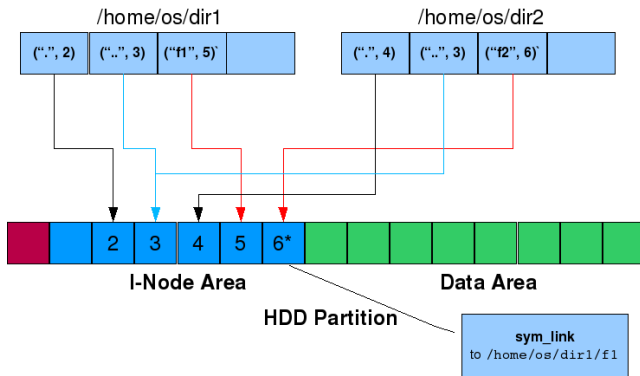


Figure: Symbolic Link Implementation

Symbolic Link: Discussion

- Advantages
 - can be created between different partitions
- Disadvantages
 - once the referenced path is removed, the link become invalid