

Chapter 5.1

Process Management

General Presentation and Linux System Calls

Print Version of Lectures Notes of *Operating Systems*

Technical University of Cluj-Napoca (UTCN)
Computer Science Department

Adrian Coleşa
April 3rd, 2019

5.1.1

Purpose and Contents

The purpose of today's lecture

- Presents general aspects related to process management
- Give examples and details about Linux system calls for processes

5.1.2

Bibliography

- A. Tanenbaum, *Modern Operating Systems*, 2nd Edition, 2001, Chapter 2, Processes, pg. 71 – 100, pg. 132 – 151

5.1.3

Contents

1	General Aspects	1
2	Linux Processes	3
2.1	System Calls	3
2.2	Examples	9
2.3	Relates Issues	10
3	Conclusions	11

5.1.4

1 General Aspects

Process Definition

- Longman dictionary's definition of **process**
 - *a series of actions that are done in order to achieve a particular result*
- a **program in execution** ⇔ an user application
 - a **sequential** stream of **execution** in its own **memory address space**
 - including the current **values of CPU's registers** (e.g. IP)
- **OS abstraction for using the computer**
 - composed by all that is needed to run a program: CPU, memory, I/O devices etc.
 - it is a **virtualization concept** → **virtualizes an entire system** (computer)
 - ⇒ **isolation mechanism**, i.e. isolates one execution (process) by another

5.1.5

Process vs Program

- program's source code $\xrightarrow{\text{compilation}}$ program's executable $\xrightarrow{\text{launching}}$ process
- **program = static (inactive) entity**
- **process = active entity**
- a process
 - is an activity of some kind
 - is created from a program loaded in memory
 - is allocated system resources (memory, file descriptors, CPU etc.)
 - has input, output, and a state
- the two **parts of a process**
 - **sequential execution**: no concurrency inside a process; everything happens sequentially
 - **process state**: everything that process interacts with (registers, memory, files etc)

5.1.6

The Considered Context

- single-processor systems
- multiprogramming and time (processor) sharing
 - pseudo-parallelism
 - switching among processes
 - scheduling algorithm

5.1.7

Process Creation

- automatically by the OS (uncommon case)
 - at system initialization
 - reacting to different events
 - usually run as background processes (vs. foreground processes)
- **by another process** (common case)
 - there is a **system call** provided for process creation
 - situations
 - * when a process needs help doing some computation
 - * when a user action occurs, e.g. interaction with the shell
 - leads to a **process hierarchy** based on the **parent-child** relationship

5.1.8

Process termination

- **voluntarily**, using a special **system call**
 - **normal exit**, i.e. end of program's execution
 - **error detection exit**, like: inexistent files, insufficient or incorrect input etc.
- **involuntarily**, being forcefully terminated
 - initiated by the system due to a "fatal error", like: illegal instructions, division by zero, segmentation fault etc.
 - initiated by another process (i.e. "killed")

5.1.9

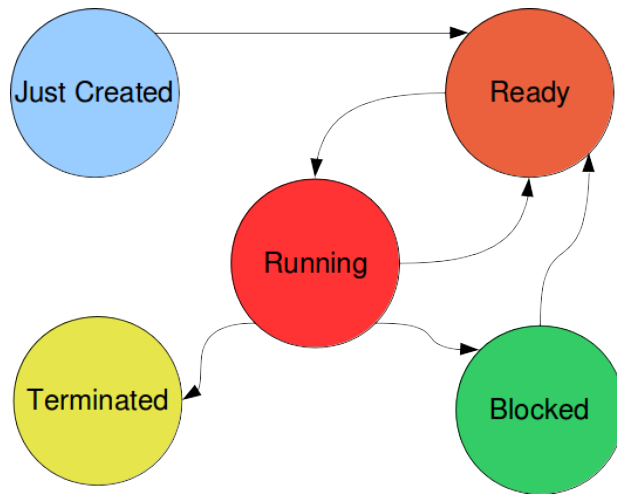


Figure 1: Process States Transition

Process states

- **running**
 - **executed by the CPU**, i.e. using the CPU, at that moment
 - only one process in that state / CPU, actually as many as the number of system's CPUs
- **ready**
 - ready to be executed, but no CPU available
 - so **wait for a CPU to become available**
 - **transparent to the program**
- **blocked**
 - **wait for an event to occur**, a resource to become available
 - **triggered by the application explicitly** through **blocking system calls**
- just created (optional)
 - waiting for some resources to be allocated
- **terminated** (optional)
 - keeping information about the exit state

5.1.10

Process States Transitions

5.1.11

2 Linux Processes

2.1 System Calls

Process Creation: `fork()`

- system call used to **create a new process**
- child process' contents is identical with that of its parent
- *still, two distinct and independent processes*
- the two processes are scheduled independently on the CPU
- parent processes continue its execution returning from `fork`
- child starts its execution returning from `fork`
- `fork` returns
 - a positive value (child's PID) in parent
 - zero in child

5.1.12

fork Usage Example

```
int x;
static int y;
int *px;

int main(int argc, char **argv)
{
    int pid;

    x = 0;
    px = &x;
    y = 0;

    // up to this point only the parent exists
    // now parent calls fork() to create a new process
    pid = fork();
    if (pid < 0) {
        // error case: no child process created
        perror("Cannot create a new process");
        exit(1);
    }
    // from now on there are two processes: parent and child

    // executed by both processes
    printf("x=%d, px=%p, *px=%d, y=%d\n", x, px, *px, y);
    // parent: x=0, px=0x601050, *px=0, y=0
    // child: x=0, px=0x601050, *px=0, y=0

    if (pid == 0) { // executed only by the child
        y = 20;
        *px = 200;
    } else {        // executed only by the parent
        y = 10;
        *px = 100;
    }

    // executed by both processes
    printf("x=%d, px=%p, *px=%d, y=%d\n", x, px, *px, y);
    // parent: x=100, px=0x601050, *px=100, y=10
    // child: x=200, px=0x601050, *px=200, y=20
}
```

5.1.13

fork's Effect Illustration

5.1.14

fork() syscall creates an *independent* child process, which starts as a *copy of its parent*!

5.1.15

Let's practice!

Have you really understood how fork() works?

If you have, try solving the following problems:

1. You are given the following code:

```
fork();
fork();
```

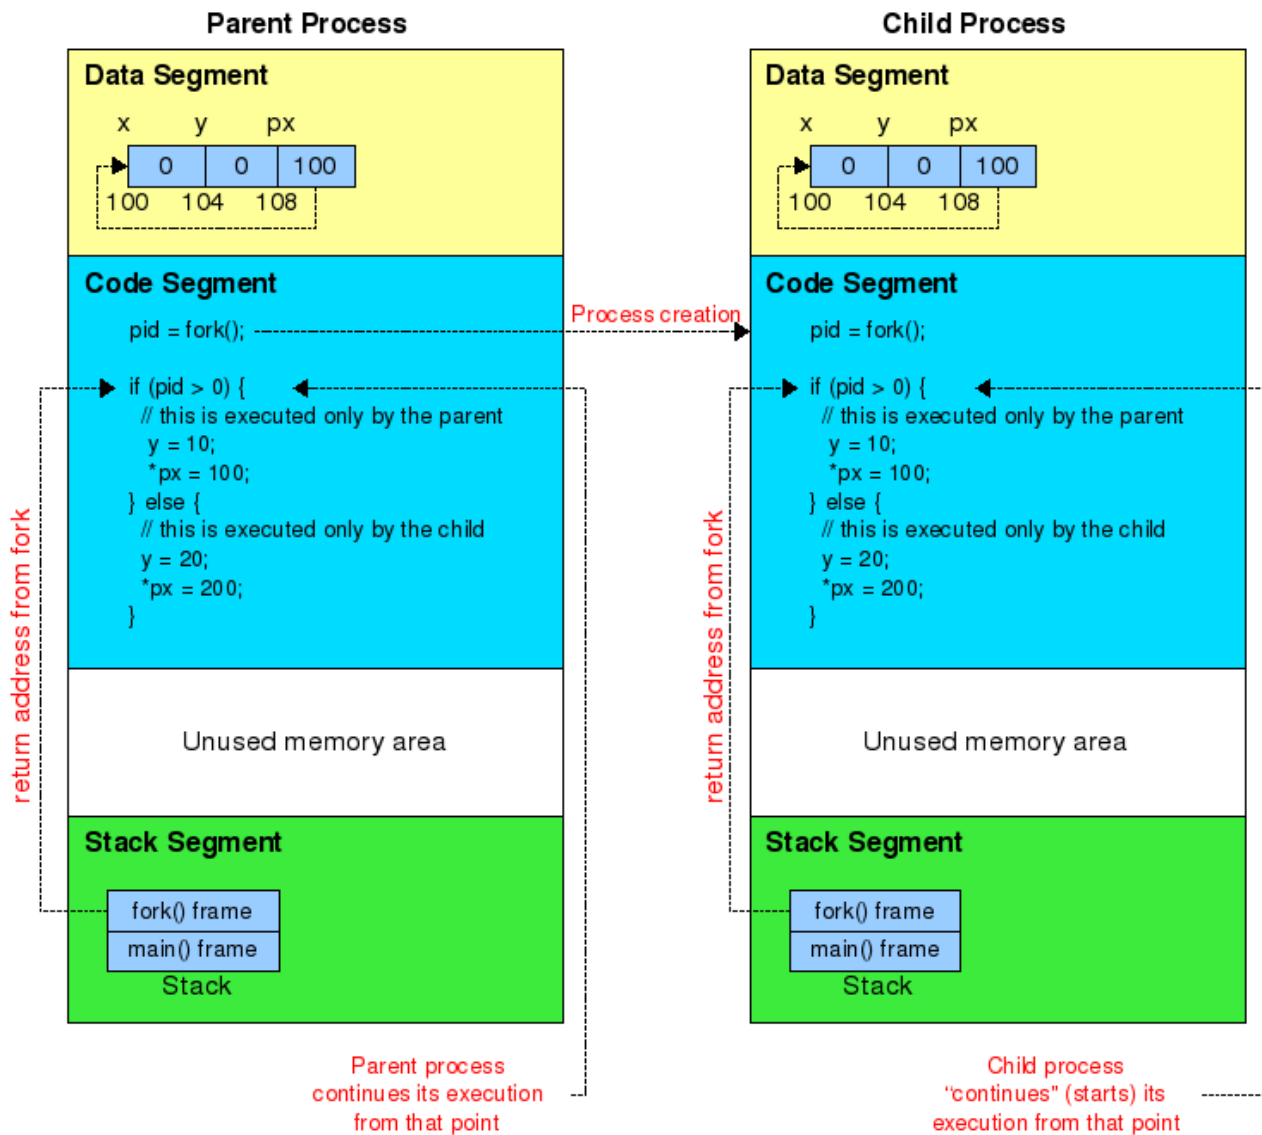
- How many processes does the following code creates?
- Draw the resulted process hierarchy.

2. You are given the following code:

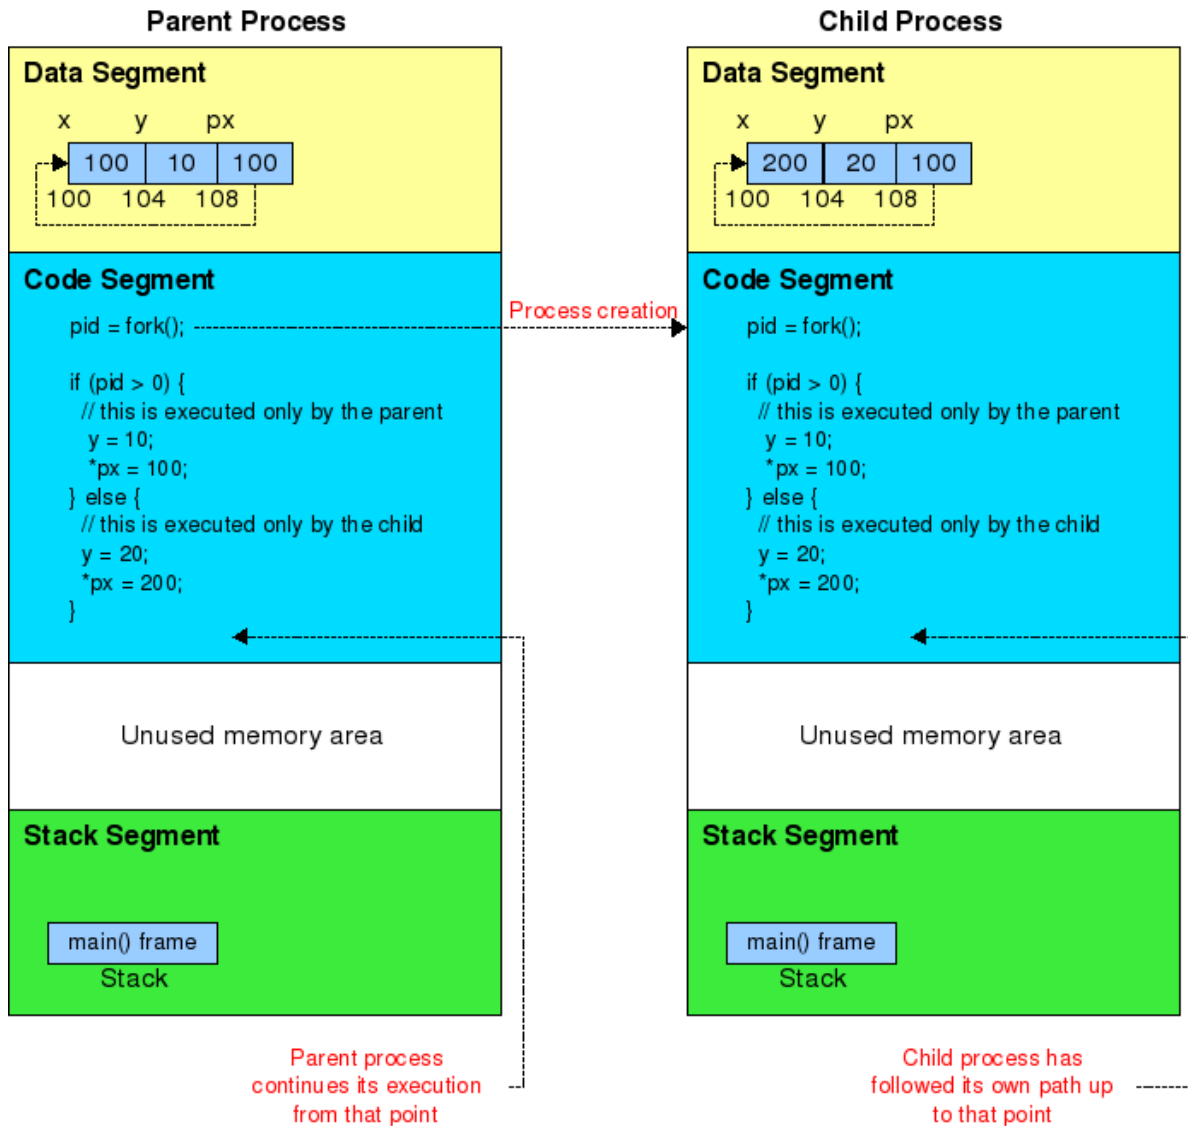
```
for(i=1; i<=100; i++)
    fork();
```

- How many processes does the following code creates?
- Draw the resulted process hierarchy.

5.1.16



**Processes Status Immediately After Fork.
The Child Contents Is Identical With That of Its Parent**



**Processes Status After Some Time.
The Two Processes Evolve Independently**

Code Execution: exec Family

- system call used to **load a new code into the calling process**
 - replace the calling process' contents, but not its identity
- there are more exec system calls
 - `execl`, `execvp`: with variable number of arguments
 - `execv`, `execvp`: with a fixed number of arguments
- the `exec`'s parameters **similar to a command line**
 - the first argument is **always the path** to the executable file
 - the next argument(s) describe the command line, *starting with command name*

5.1.17

`execl` and `execvp` Usage Example

```
// first parameter is the EXPLICIT path to the executable file
execl("/bin/ls", "ls", "-l", NULL);
execl("./myprg.exe", "myprg.exe", "param1", "param2", NULL);

// first parameter is the IMPLICIT path to the executable file
// the path is searched in the directories stored
// in the PATH environment variable
execvp("ls", "ls", "-l", 0);
```

5.1.18

`execv` and `execvp` Usage Example

```
char cmdline[10][100]; // equiv. to char *cmdline[];
                        // equiv. to char **cmdline;

// build the command line
strncpy(cmdline[0], "ls", 99);
strncpy(cmdline[1], "-l", 99);
cmdline[2] = NULL;

// call the exec

// first parameter is the EXPLICIT path to the executable file
execv("/bin/ls", cmdline);

// first parameter is the IMPLICIT path to the executable file
// the path is searched in the directories stored
// in the PATH environment variable
execvp("ls", cmdline);
```

5.1.19

`exec()` syscalls loads a new code in the calling process! There is no return from `exec()` if successfully executed!

5.1.20

Relationship Between `fork` and `exec`

- used to
 - create a child process
 - executing something else than its parent
- **Why there are two separated steps instead of just one?**
 - between them the parent “has control” over its child (see standard input and output redirection below)
 - the **parent is released** by the burden of (i.e. time spent) **loading a new code in child**
 - \Rightarrow **better performance** for the parent

5.1.21

fork and exec Usage Example

- parent code

```
int main()
{
    pid = fork();

    if (pid > 0) {
        // parent doing something
    } else {
        // child loading and executing a new code
        execl("./child.exe", "child.exe", "p1", "10", 0);
        perror("execl has not succeeded");
    }
}
```

- child code

```
int main(int argc, char **argv)
{
    int p;
    for (p=0; p<argc; p++)
        printf("argv[%d]=%s\n", argv[p]);
}
```

5.1.22

Process Termination: exit

- system call used to **terminate voluntarily a process**
- terminate the calling process
- specify an exit code
 - 0 (zero) exit code considered successfully termination
 - anything else considered erroneously termination
- exit code is kept until the parent process asks for it
- example
 - `exit(0);`
 - `exit(1);`

5.1.23

Wait For Termination: wait and waitpid

- system calls used by a process to **wait for the termination of its children**
- return the exit code of the terminated child
- example

```
int child_status;

wait(&child_status);

printf("Child process terminated with code %d\n",
       WEXITSTATUS(status));
```

5.1.24

Relationship Between wait and exit

- a way to **synchronize two processes'** execution (parent and child)
- a simple way to **communicate between processes** (parent and child)
- when a (parent) process terminates
 - all its children get as their new parent a system process
 - * on older systems: the *init* process, having `pid = 1`
 - * on newer systems: a per user *init* process
- when a (child) process terminates before its parent
 - its state is said to be **zombie** and
 - its exit state is maintained by OS until its parent process asks for it or terminates

5.1.25

Let's practice!

You are given the following C code and are required to:

1. Draw the process hierarchy corresponding to the processes created by the code below.
2. Specify the number of times each *printf* is executed, supposing every instruction is executed successfully.

```
1 printf ("[1] Hello world!\n");
2
3 pid = fork();
4
5 printf ("[2] Hello world!\n");
6
7 pid = fork();
8
9 printf ("[3] Hello world!\n");
10
11 if (pid == 0) {
12     execlp("ps", "ps", 0);
13     printf ("[4] Hello world!\n");
14 }
15
16 fork();
17
18 printf ("[5] Hello world!\n");
```

5.1.26

2.2 Examples

Shell Basic Code (Functionality)

```
char **cmdline; // it must be build like argv param of main
while (TRUE) {
    display_prompt_on_screen();
    cmdline = read_cmd_line();

    pid = fork(); // creates a new process
    if (pid < 0) {
        perror("cannot creat a new process");
        continue;
    }

    if (pid == 0)
        execlvp(cmdline[0], cmdline);
    else
        waitpid(pid, NULL, 0);
}
```

5.1.27

Standard Input Redirection

- command line

```
cat < file.txt
```

- STDIN redirection in C program

```
pid = fork();

if (pid > 0) {
    // parent
} else {
    // child
    fd = open("file.txt", O_RDONLY);
    dup2(fd, 0);
    close(fd);

    execlp("cat", "cat", 0);
    perror("execl has not succeded");
}
```

5.1.28

Standard Output Redirection

- command line

```
ls > file.txt
```

- STDIN redirection in C program

```
pid = fork();

if (pid > 0) {
    // parent
} else {
    // child
    fd = creat("file.txt", 0600);
```

```

mc - linux:~/school/os
File Edit View Terminal Tabs Help
acolesa@linux:~/school/os$ ps -l -u acolesa --forest
F S  UID  PID  PPID  C PRI  NI ADDR SZ WCHAN  TTY      TIME CMD
4 S  1000 17351 3772  0 75  0 - 5293 - ?      00:00:00 x-session-manag
1 S  1000 17394 17351  0 75  0 - 1034 429496 ?      00:00:00 \_ ssh-agent
0 S  1000 26952  1  0 75  0 - 8821 - ?      00:00:00 gimp-2.2
0 S  1000 26953 26952  0 78  0 - 3413 429496 ?      00:00:00 \_ script-fu
0 S  1000 24740  1  0 78  0 - 1040 wait ?      00:00:00 soffice
0 S  1000 24754 24740  0 75  0 - 58102 stext ?      00:00:06 \_ soffice.bin
1 S  1000 24621  1  0 82  0 - 7115 429496 ?      00:00:00 kio_userver
1 S  1000 24469  1  0 75  0 - 6674 - ?      00:00:00 kded
1 S  1000 24463  1  0 77  0 - 6187 429496 ?      00:00:00 dcopserver
1 S  1000 24459  1  0 76  0 - 6251 429496 ?      00:00:00 kdeinit
1 S  1000 24467 24459  0 75  0 - 6382 - ?      00:00:00 \_ klauncher

```

```

mc - linux:~/school/os
File Edit View Terminal Tabs Help
acolesa@linux:~/school/os$ ps -l -u acolesa --forest | grep -v '??'
F S  UID  PID  PPID  C PRI  NI ADDR SZ WCHAN  TTY      TIME CMD
0 S  1000 24478 24457  0 75  0 - 1744 - pts/3      00:00:00 \_ bash
0 S  1000 17821 17819  0 75  0 - 1747 wait pts/0      00:00:00 \_ bash
0 S  1000 24111 17821  0 75  0 - 1733 429496 pts/0      00:00:00 \_ mc
0 S  1000 24113 24111  0 75  0 - 1747 wait pts/1      00:00:00 \_ bash
0 R  1000 27214 24113  0 77  0 - 851 - pts/1      00:00:00 \_ ps
0 S  1000 27215 24113  0 75  0 - 712 pipe_w pts/1      00:00:00 \_ grep
acolesa@linux:~/school/os$ 

```

```

dup2(fd, 1);
close(fd);

execlp("ls", "ls", 0);
perror("execl has not succeeded");
}

```

5.1.29

2.3 Relates Issues

ps Command

- displays a snapshot of the active processes in the system
- `ps -l -u acolesa --forest`
- `ps -l -u acolesa --forest | grep -v '??'`
- `ps -l -e --forest | head -n 50`

5.1.30

top and htop Commands

- display a continuously updated list of processes and their on-line scheduling

5.1.31

proc File System

- It is a pseudo file system
- It is mounted in `/proc`
- It is used by the OS to display information about processes
 - each process has a directory named with the process id
 - reading this information is similar to reading any other files and dirs

5.1.32

```

mc - linux:~/school/os
File Edit View Terminal Tabs Help
acolesa@linux:~/school/os$ ps -l -e --forest | head -n 50
F S  UID  PID  PPID  C PRI  NI ADDR SZ WCHAN  TTY      TIME CMD
4 S  0 1 0 0 75 0 - 486 - ?      00:00:01 init
1 S  0 2 1 0 -40 - - 0 migrat ?      00:00:00 migration/0
1 S  0 3 1 0 94 19 - 0 ksofti ?      00:00:00 ksoftirqd/0
1 S  0 4 1 0 -40 - - 0 migrat ?      00:00:00 migration/1
1 S  0 5 1 0 94 19 - 0 ksofti ?      00:00:00 ksoftirqd/1
5 S  0 6 1 0 70 -5 - 0 worker ?      00:00:00 events/0
1 S  0 7 1 0 70 -5 - 0 worker ?      00:00:00 events/1
1 S  0 8 1 0 70 -5 - 0 worker ?      00:00:00 khelper
1 S  0 9 1 0 72 -5 - 0 worker ?      00:00:00 kthread
1 S  0 13 9 0 70 -5 - 0 worker ?      00:00:00 \_ kblockd/0
1 S  0 14 9 0 70 -5 - 0 worker ?      00:00:00 \_ kblockd/1

```

3 Conclusions

What we talked about

- process definition
- process states and state transitions
 - running, ready, blocked, terminated
- system calls to create and terminate a process
- Linux system calls related to processes
 - fork
 - exec
 - exit
 - wait

5.1.33

Lessons Learned

- process is a virtualization and isolation concept
 - virtualize an entire compute for a program's execution
 - isolate one execution by another
- process states
 - running: the desired one
 - ready: exists due to limited no of CPUs; is transparent to processes
 - blocked: triggered explicitly by a process due to a blocking syscall
- create a new process
 - `fork()` called by the parent to create an identical child process
 - `exec()` called by the child to load new code
- every process terminates with `exit()`

5.1.34