

Process Management

Thread Model. Process Scheduling

Adrian Coleșa

Technical University of Cluj-Napoca (UTCN)
Computer Science Department

April 08, 2020

The purpose of this subchapter

- Defines threads and the way they can be used
- Presents scheduling principles and algorithms

The purpose of this subchapter

- Defines threads and the way they can be used
- Presents scheduling principles and algorithms

Bibliography

- A. Tanenbaum, *Modern Operating Systems*, 2nd Edition, 2001, Chapter 2, Processes, pg. 71 – 100, pg. 132 – 151

Outline

- 1 Threads
 - Definition of Concepts
 - Thread Usage
- 2 Scheduling
 - Concepts
 - Scheduling Algorithms
- 3 Conclusions

Outline

1 Threads

- Definition of Concepts
- Thread Usage

2 Scheduling

- Concepts
- Scheduling Algorithms

3 Conclusions

Outline

- 1 Threads
 - Definition of Concepts
 - Thread Usage
- 2 Scheduling
 - Concepts
 - Scheduling Algorithms
- 3 Conclusions

Thread Definition

- describes a **sequential, independent execution** within a process
 - execution = the memory path followed in the code by the IP register
- **there could be more**
 - simultaneous and independent executions in the same process
 - \Rightarrow threads in a process

Thread Definition

- describes a **sequential, independent execution** within a process
 - execution = the memory path followed in the code by the IP register
- **there could be more**
 - simultaneous and independent executions in the same process
 - \Rightarrow threads in a process

Thread Definition

- describes a **sequential, independent execution** within a process
 - execution = the memory path followed in the code by the IP register
- **there could be more**
 - simultaneous and independent executions in the same process
 - \Rightarrow **threads in a process**

Thread Definition

- describes a **sequential, independent execution** within a process
 - execution = the memory path followed in the code by the IP register
- **there could be more**
 - simultaneous and independent executions in the same process
 - \Rightarrow **threads in a process**

Thread Definition

- describes a **sequential, independent execution** within a process
 - execution = the memory path followed in the code by the IP register
- **there could be more**
 - simultaneous and independent executions in the same process
 - \Rightarrow **threads in a process**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- **isolates resources** of one process by that of others

- **thread** model

- models the **processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describing the execution
- if multiple independent executions could be identified in a process
 - there could be many threads of that process
 - one thread is one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- **isolates resources** of one process by that of others

- **thread** model

- models the **processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describing the execution
- if multiple independent executions could be identified in a process
 - there could be many threads of that process
 - one thread is one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- isolates resources of one process by that of others

- **thread** model

- models the **processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describing the execution
- if multiple independent executions could be identified in a process
 - we could be having threads of that process
 - one thread is one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- isolates resources of one process by that of others

- **thread** model

- models the **processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describing the execution
- if multiple independent executions could be identified in a process
 - we could be having threads of that process
 - one thread is one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- isolates resources of one process by that of others

- **thread** model

- models the **processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describing the execution
- if multiple independent executions could be identified in a process
 - we could be having threads of that process
 - one thread is one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- **isolates resources** of one process by that of others

- **thread** model

- models the **processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describing the execution
- if multiple independent executions could be identified in a process
 - each could be using threads of that process
 - one thread is one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- **isolates resources** of one process by that of others

- **thread** model

- models **the processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describes the execution
- if multiple independent executions could be identified in a process
 - there could be more threads of that process
 - one thread \leftrightarrow one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- **isolates resources** of one process by that of others

- **thread** model

- models **the processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describes the execution
- if multiple independent executions could be identified in a process
 - there could be more threads of that process
 - one thread \leftrightarrow one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- **isolates resources** of one process by that of others

- **thread** model

- models **the processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describes the execution
- if multiple independent executions could be identified in a process
 - there could be more threads of that process
 - one thread \leftrightarrow one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- **isolates resources** of one process by that of others

- **thread** model

- models **the processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describes the execution
- if multiple independent executions could be identified in a process
 - there could be more threads of that process
 - one thread \leftrightarrow one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- **isolates resources** of one process by that of others

- **thread** model

- models **the processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describes the execution
- if multiple independent executions could be identified in a process
 - there could be more threads of that process
 - one thread \leftrightarrow one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- **isolates resources** of one process by that of others

- **thread** model

- models **the processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describes the execution
- if multiple independent executions could be identified in a process
 - there could be more threads of that process
 - one thread \leftrightarrow one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- **isolates resources** of one process by that of others

- **thread** model

- models **the processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describes the execution
- if multiple independent executions could be identified in a process
 - there could be more threads of that process
 - one thread \leftrightarrow one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- **isolates resources** of one process by that of others

- **thread** model

- models **the processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describes the execution
- if multiple independent executions could be identified in a process
 - there could be more threads of that process
 - one thread \leftrightarrow one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- **isolates resources** of one process by that of others

- **thread** model

- models **the processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describes the execution
- if multiple independent executions could be identified in a process
 - there could be more threads of that process
 - one thread \leftrightarrow one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Modeling The Execution

- **process** model

- models “**an entire computer**” used for executing some user application
- composed by both
 - process' resources
 - process' execution
- **isolates resources** of one process by that of others

- **thread** model

- models **the processor(s)** given to a process
- makes distinction between the two components of a process
 - the process describes the resources
 - the thread describes the execution
- if multiple independent executions could be identified in a process
 - there could be more threads of that process
 - one thread \leftrightarrow one execution
- threads of a process **share resources** of that process
- threads of a process (normally) **not visible to other processes**

Multiprogramming and Multithreading

- multiprogramming = multiple processes managed simultaneously
- multithreading = multiple threads in the same process managed simultaneously

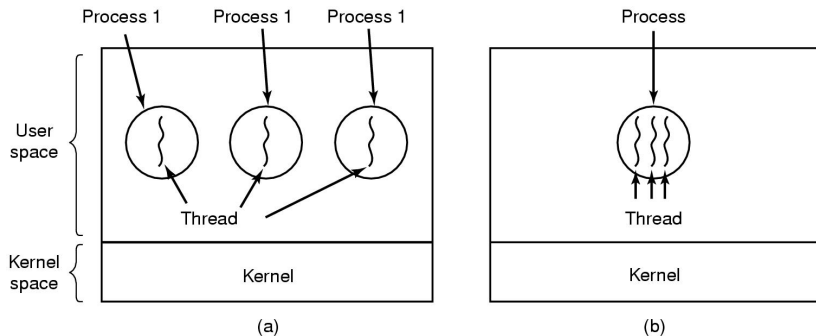


Figure: Taken from Tanenbaum

Multiprogramming and Multithreading

- multiprogramming = multiple processes managed simultaneously
- multithreading = multiple threads in the same process managed simultaneously

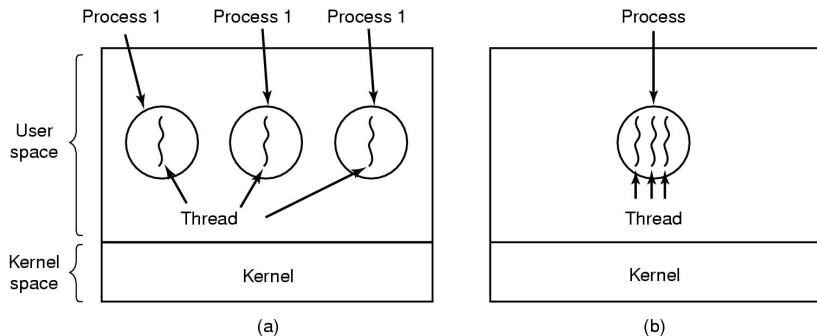


Figure: Taken from Tanenbaum

Thread's Components. Description

- **threads** of the same process
 - **share all the resources of** that **process**: memory, open files etc.
 - \Rightarrow anything (i.e. inside the process) is accessible to all threads (of that process)
- each thread is an independent execution
 - \Rightarrow each thread has also its own resources
 - that describe the corresponding independent execution
- thread's specific resources
 - machine registers (e.g. the IP register)
 - the stack

Thread's Components. Description

- **threads** of the same process
 - **share all the resources of that process:** memory, open files etc.
 - \Rightarrow anything (i.e. inside the process) is accessible to all threads (of that process)
- each thread is an independent execution
 - \Rightarrow each thread has also its own resources
 - that describe the corresponding independent execution
- thread's specific resources
 - machine registers (e.g. the IP register)
 - the stack

Thread's Components. Description

- **threads** of the same process
 - **share all the resources of** that **process**: memory, open files etc.
 - \Rightarrow anything (i.e. inside the process) is accessible to all threads (of that process)
- each thread is an independent execution
 - \Rightarrow each thread has also its own resources
 - that describe the corresponding independent execution
- thread's specific resources
 - machine registers (e.g. the IP register)
 - the stack

Thread's Components. Description

- **threads** of the same process
 - **share all the resources of that process:** memory, open files etc.
 - \Rightarrow anything (i.e. inside the process) is accessible to all threads (of that process)
- each thread is an independent execution
 - \Rightarrow **each thread has also its own resources**
 - that describe the corresponding independent execution
- thread's specific resources
 - machine registers (e.g. the IP register)
 - the stack

Thread's Components. Description

- **threads** of the same process
 - **share all the resources of that process:** memory, open files etc.
 - \Rightarrow anything (i.e. inside the process) is accessible to all threads (of that process)
- each thread is an independent execution
 - \Rightarrow **each thread has also its own resources**
 - that describe the corresponding independent execution
- thread's specific resources
 - machine registers (e.g. the IP register)
 - the stack

Thread's Components. Description

- **threads** of the same process
 - **share all the resources of** that **process**: memory, open files etc.
 - \Rightarrow anything (i.e. inside the process) is accessible to all threads (of that process)
- each thread is an independent execution
 - \Rightarrow **each thread has also its own resources**
 - that describe the corresponding independent execution
- thread's specific resources
 - machine registers (e.g. the IP register)
 - the stack

Thread's Components. Description

- **threads** of the same process
 - **share all the resources of** that **process**: memory, open files etc.
 - \Rightarrow anything (i.e. inside the process) is accessible to all threads (of that process)
- each thread is an independent execution
 - \Rightarrow **each thread has also its own resources**
 - that describe the corresponding independent execution
- thread's specific resources
 - machine registers (e.g. the IP register)
 - the stack

Thread's Components. Description

- **threads** of the same process
 - **share all the resources of** that **process**: memory, open files etc.
 - \Rightarrow anything (i.e. inside the process) is accessible to all threads (of that process)
- each thread is an independent execution
 - \Rightarrow **each thread has also its own resources**
 - that describe the corresponding independent execution
- thread's specific resources
 - **machine registers** (e.g. the IP register)
 - the stack

Thread's Components. Description

- **threads** of the same process
 - **share all the resources of** that **process**: memory, open files etc.
 - \Rightarrow anything (i.e. inside the process) is accessible to all threads (of that process)
- each thread is an independent execution
 - \Rightarrow **each thread has also its own resources**
 - that describe the corresponding independent execution
- thread's specific resources
 - **machine registers** (e.g. the IP register)
 - **the stack**

Thread's Components. Illustration





Figure: Taken from Tanenbaum

Threads Relationship

**Paraphrase a rule from
“Animal Farm” by George Orwell**

Threads Relationship (cont.)




All threads are equal,
but the “main” thread is
*“more equal than the
others!”*

Threads Relationship (cont.)



All threads are equal,
but the “main” thread is
*“more equal than the
others!”*

Threads Relationship (cont.)



**All threads are equal,
but the “main” thread is
*“more equal than the
others!”***

Consequences of “All threads are equal”

- **there is no protection between threads of the same process**
 - there is no need, actually
 - created by the application itself to execute code inside the application
- **there is no restriction on them regarding their scheduling**
 - any one could be scheduled any time, depending on the system's scheduling policy
 - we have no control on thread scheduling
 - if needed, we must synchronize the threads

Consequences of “All threads are equal”

- **there is no protection between threads of the same process**
 - there is no need, actually
 - created by the application itself to execute code inside the application
- **there is no restriction on them regarding their scheduling**
 - any one could be scheduled any time, depending on the system's scheduling policy
 - we have no control on thread scheduling
 - if needed, we must synchronize the threads

Consequences of “All threads are equal”

- **there is no protection between threads of the same process**
 - there is no need, actually
 - created by the application itself to execute code inside the application
- **there is no restriction on them regarding their scheduling**
 - any one could be scheduled any time, depending on the system's scheduling policy
 - we have no control on thread scheduling
 - if needed, we must synchronize the threads

Consequences of “All threads are equal”

- **there is no protection between threads of the same process**
 - there is no need, actually
 - created by the application itself to execute code inside the application
- **there is no restriction on them regarding their scheduling**
 - any one could be scheduled any time, depending on the system's scheduling policy
 - we have no control on thread scheduling
 - if needed, **we must synchronize the threads**

Consequences of “All threads are equal”

- **there is no protection between threads of the same process**
 - there is no need, actually
 - created by the application itself to execute code inside the application
- **there is no restriction on them regarding their scheduling**
 - any one could be scheduled any time, depending on the system's scheduling policy
 - we have no control on thread scheduling
 - if needed, **we must synchronize the threads**

Consequences of “All threads are equal”

- **there is no protection between threads of the same process**
 - there is no need, actually
 - created by the application itself to execute code inside the application
- **there is no restriction on them regarding their scheduling**
 - any one could be scheduled any time, depending on the system's scheduling policy
 - we have no control on thread scheduling
 - if needed, **we must synchronize the threads**

Consequences of “All threads are equal”

- **there is no protection between threads of the same process**
 - there is no need, actually
 - created by the application itself to execute code inside the application
- **there is no restriction on them regarding their scheduling**
 - any one could be scheduled any time, depending on the system's scheduling policy
 - we have no control on thread scheduling
 - if needed, **we must synchronize the threads**

Consequences of “All threads are equal” (cont.)

- ... but the main one is “*more equal than the others*”
 - it is the first one created: **executes the *main()* function**
 - it is the ancestor of all other threads, i.e. initiates creation of the other threads
 - if it terminates returning from the “*main()*” function
 - the entire process (and all threads) terminates
 - ⇒ it make sense for it to wait for the termination of all other threads
 - **though**, if it terminates using the thread termination syscall, the other threads survive
 - see difference between `exit()` and `pthread_exit()` in Linux
 - see a discussion about this at <https://devblogs.microsoft.com/oldnewthing/20100827-00/?p=13023>

Consequences of “All threads are equal” (cont.)

- ... but the main one is “*more equal than the others*”
 - it is the first one created: **executes the *main()* function**
 - it is the ancestor of all other threads, i.e. initiates creation of the other threads
 - if it terminates returning from the “*main()*” function
 - the entire process (and all threads) terminates
 - ⇒ it make sense for it to wait for the termination of all other threads
 - **though**, if it terminates using the thread termination syscall, the other threads survive
 - see difference between `exit()` and `pthread_exit()` in Linux
 - see a discussion about this at <https://devblogs.microsoft.com/oldnewthing/20100827-00/?p=13023>

Consequences of “All threads are equal” (cont.)

- ... but the main one is “*more equal than the others*”
 - it is the first one created: **executes the *main()* function**
 - it is the ancestor of all other threads, i.e. initiates creation of the other threads
 - if it terminates returning from the “*main()*” function
 - the entire process (and all threads) terminates
 - ⇒ it make sense for it to wait for the termination of all other threads
 - **though**, if it terminates using the thread termination syscall, the other threads survive
 - see difference between `exit()` and `pthread_exit()` in Linux
 - see a discussion about this at <https://devblogs.microsoft.com/oldnewthing/20100827-00/?p=13023>

Consequences of “All threads are equal” (cont.)

- ... but the main one is “*more equal than the others*”
 - it is the first one created: **executes the *main()* function**
 - it is the ancestor of all other threads, i.e. initiates creation of the other threads
 - if it terminates returning from the “`main()`” function
 - the entire process (and all threads) terminates
 - \Rightarrow it make sense for it to **wait for the termination of all other threads**
 - **though**, if it terminates using the thread termination syscall, the other threads survive
 - see difference between `exit()` and `pthread_exit()` in Linux
 - see a discussion about this at <https://devblogs.microsoft.com/oldnewthing/20100827-00/?p=13023>

Consequences of “All threads are equal” (cont.)

- ... but the main one is “*more equal than the others*”
 - it is the first one created: **executes the *main()* function**
 - it is the ancestor of all other threads, i.e. initiates creation of the other threads
 - if it terminates returning from the “`main()`” function
 - the entire process (and all threads) terminates
 - \Rightarrow it make sense for it to **wait for the termination of all other threads**
 - **though**, if it terminates using the thread termination syscall, the other threads survive
 - see difference between `exit()` and `pthread_exit()` in Linux
 - see a discussion about this at <https://devblogs.microsoft.com/oldnewthing/20100827-00/?p=13023>

Consequences of “All threads are equal” (cont.)

- ... but the main one is “*more equal than the others*”
 - it is the first one created: **executes the *main()* function**
 - it is the ancestor of all other threads, i.e. initiates creation of the other threads
 - if it terminates returning from the “`main()`” function
 - the entire process (and all threads) terminates
 - \Rightarrow it make sense for it to **wait for the termination of all other threads**
 - **though**, if it terminates using the thread termination syscall, the other threads survive
 - see difference between `exit()` and `pthread_exit()` in Linux
 - see a discussion about this at <https://devblogs.microsoft.com/oldnewthing/20100827-00/?p=13023>

Consequences of “All threads are equal” (cont.)

- ... but the main one is “*more equal than the others*”
 - it is the first one created: **executes the *main()* function**
 - it is the ancestor of all other threads, i.e. initiates creation of the other threads
 - if it terminates returning from the “`main()`” function
 - the entire process (and all threads) terminates
 - \Rightarrow it make sense for it to **wait for the termination of all other threads**
 - **though**, if it terminates using the thread termination syscall, the other threads survive
 - see difference between `exit()` and `pthread_exit()` in Linux
 - see a discussion about this at <https://devblogs.microsoft.com/oldnewthing/20100827-00/?p=13023>

Consequences of “All threads are equal” (cont.)

- ... but the main one is “*more equal than the others*”
 - it is the first one created: **executes the *main()* function**
 - it is the ancestor of all other threads, i.e. initiates creation of the other threads
 - if it terminates returning from the “`main()`” function
 - the entire process (and all threads) terminates
 - \Rightarrow it make sense for it to **wait for the termination of all other threads**
 - **though**, if it terminates using the thread termination syscall, the other threads survive
 - see difference between `exit()` and `pthread_exit()` in Linux
 - see a discussion about this at <https://devblogs.microsoft.com/oldnewthing/20100827-00/?p=13023>

Consequences of “All threads are equal” (cont.)

- ... but the main one is “*more equal than the others*”
 - it is the first one created: **executes the *main()* function**
 - it is the ancestor of all other threads, i.e. initiates creation of the other threads
 - if it terminates returning from the “*main()*” function
 - the entire process (and all threads) terminates
 - \Rightarrow it make sense for it to **wait for the termination of all other threads**
 - **though**, if it terminates using the thread termination syscall, the other threads survive
 - see difference between `exit()` and `pthread_exit()` in Linux
 - see a discussion about this at <https://devblogs.microsoft.com/oldnewthing/20100827-00/?p=13023>

Outline

- 1 Threads
 - Definition of Concepts
 - Thread Usage
- 2 Scheduling
 - Concepts
 - Scheduling Algorithms
- 3 Conclusions

When Do Threads Prove Themselves Useful?

- **useful** in applications

- where **different actions can happen concurrently** (in the same time)
→ **logical parallelism**
- and use **different “independent” system resources** (e.g. CPU, HDD, etc.) → **physical parallelism**

- **NOT useful** in applications

- using only the single CPU (our considered context)
 - though useful, if multiple CPUs are available
- in which possible concurrent executions need all the time (compete for) the same single shared resource

When Do Threads Prove Themselves Useful?

- **useful** in applications
 - where **different actions can happen concurrently** (in the same time)
→ **logical parallelism**
 - and use **different “independent” system resources** (e.g. CPU, HDD, etc.) → **physical parallelism**
- **NOT useful** in applications
 - using only the single CPU (our considered context)
 - though useful, if multiple CPUs are available
 - in which possible concurrent executions need all the time (compete for) the same single shared resource

When Do Threads Prove Themselves Useful?

- **useful** in applications
 - where **different actions can happen concurrently** (in the same time)
→ **logical parallelism**
 - and use **different “independent” system resources** (e.g. CPU, HDD, etc.) → **physical parallelism**
- **NOT useful** in applications
 - using only the single CPU (our considered context)
→ threads are not useful if multiple CPUs are available
 - in which possible concurrent executions need all the time (compete for) the same single shared resource

When Do Threads Prove Themselves Useful?

- **useful** in applications
 - where **different actions can happen concurrently** (in the same time)
→ **logical parallelism**
 - and use **different “independent” system resources** (e.g. CPU, HDD, etc.) → **physical parallelism**
- **NOT useful** in applications
 - using only the single CPU (our considered context)
 - though useful, if multiple CPUs are available
 - in which possible concurrent executions need all the time (compete for) the same single shared resource

When Do Threads Prove Themselves Useful?

- **useful** in applications
 - where **different actions can happen concurrently** (in the same time)
→ **logical parallelism**
 - and use **different “independent” system resources** (e.g. CPU, HDD, etc.) → **physical parallelism**
- **NOT useful** in applications
 - using only the single CPU (our considered context)
 - though useful, if multiple CPUs are available
 - in which possible concurrent executions need all the time (compete for) the same single shared resource

When Do Threads Prove Themselves Useful?

- **useful** in applications
 - where **different actions can happen concurrently** (in the same time)
→ **logical parallelism**
 - and use **different “independent” system resources** (e.g. CPU, HDD, etc.) → **physical parallelism**
- **NOT useful** in applications
 - using only the single CPU (our considered context)
 - though useful, if multiple CPUs are available
 - in which possible concurrent executions need all the time (compete for) the same single shared resource

When Do Threads Prove Themselves Useful?

- **useful** in applications
 - where **different actions can happen concurrently** (in the same time)
→ **logical parallelism**
 - and use **different “independent” system resources** (e.g. CPU, HDD, etc.) → **physical parallelism**
- **NOT useful** in applications
 - using only the single CPU (our considered context)
 - though useful, if multiple CPUs are available
 - in which possible concurrent executions need all the time (compete for) the same single shared resource

Examples of Multi-threaded Apps: A Word Processor




Figure: Taken from Tanenbaum

Examples of Multi-threaded Apps: A Concurrent Server

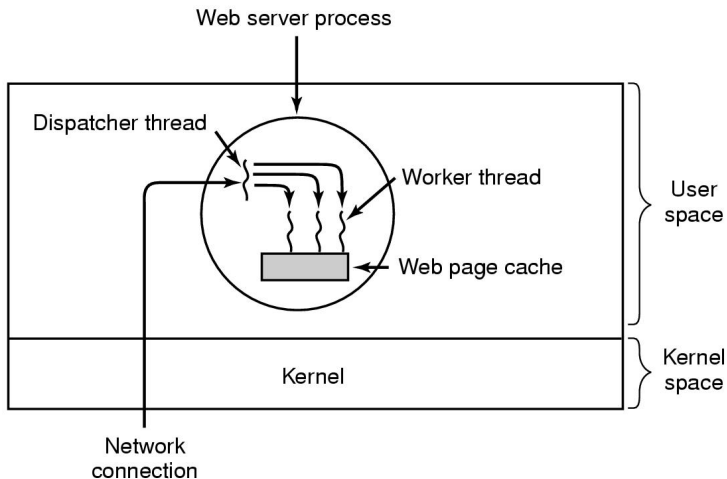


Figure: Taken from Tanenbaum

Threads vs. Processes

- which are better and for which kind of applications
- **processes** → **isolation**, i.e. totally separated by one another → protection
 - ⇒ use them when isolation and protection needed
 - e.g. tabs in an Internet browser
 - e.g. document tabs in a PDF reader
 - e.g. Web server handling client sessions
- **threads** (light-wight processes) → **resource sharing** (of the same process)
 - ⇒ use them when working on the same process' data is needed
 - e.g. parallel, cooperative computations (see examples above)

Threads vs. Processes

- which are better and for which kind of applications
- **processes** → **isolation**, i.e. totally separated by one another → protection
 - ⇒ use them when isolation and protection needed
 - e.g. tabs in an Internet browser
 - e.g. document tabs in a PDF reader
 - e.g. Web server handling client sessions
- **threads** (light-wight processes) → **resource sharing** (of the same process)
 - ⇒ use them when working on the same process' data is needed
 - e.g. parallel, cooperative computations (see examples above)

Threads vs. Processes

- which are better and for which kind of applications
- **processes** → **isolation**, i.e. totally separated by one another → protection
 - ⇒ use them when isolation and protection needed
 - e.g. tabs in an Internet browser
 - e.g. document tabs in a PDF reader
 - e.g. Web server handling client sessions
- **threads** (light-wight processes) → **resource sharing** (of the same process)
 - ⇒ use them when working on the same process' data is needed
 - e.g. parallel, cooperative computations (see examples above)

Threads vs. Processes

- which are better and for which kind of applications
- **processes** → **isolation**, i.e. totally separated by one another → protection
 - ⇒ use them when isolation and protection needed
 - e.g. tabs in an Internet browser
 - e.g. document tabs in a PDF reader
 - e.g. Web server handling client sessions
- **threads** (light-wight processes) → **resource sharing** (of the same process)
 - ⇒ use them when working on the same process' data is needed
 - e.g. parallel, cooperative computations (see examples above)

Threads vs. Processes

- which are better and for which kind of applications
- **processes** → **isolation**, i.e. totally separated by one another → protection
 - ⇒ use them when isolation and protection needed
 - e.g. tabs in an Internet browser
 - e.g. document tabs in a PDF reader
 - e.g. Web server handling client sessions
- **threads** (light-wight processes) → **resource sharing** (of the same process)
 - ⇒ use them when working on the same process' data is needed
 - e.g. parallel, cooperative computations (see examples above)

Threads vs. Processes

- which are better and for which kind of applications
- **processes** → **isolation**, i.e. totally separated by one another → protection
 - ⇒ use them when isolation and protection needed
 - e.g. tabs in an Internet browser
 - e.g. document tabs in a PDF reader
 - e.g. Web server handling client sessions
- **threads** (light-wight processes) → **resource sharing** (of the same process)
 - ⇒ use them when working on the same process' data is needed
 - e.g. parallel, cooperative computations (see examples above)

Threads vs. Processes

- which are better and for which kind of applications
- **processes** → **isolation**, i.e. totally separated by one another → protection
 - ⇒ use them when isolation and protection needed
 - e.g. tabs in an Internet browser
 - e.g. document tabs in a PDF reader
 - e.g. Web server handling client sessions
- **threads** (light-wight processes) → **resource sharing** (of the same process)
 - ⇒ use them when working on the same process' data is needed
 - e.g. parallel, cooperative computations (see examples above)

Threads vs. Processes

- which are better and for which kind of applications
- **processes** → **isolation**, i.e. totally separated by one another → protection
 - ⇒ use them when isolation and protection needed
 - e.g. tabs in an Internet browser
 - e.g. document tabs in a PDF reader
 - e.g. Web server handling client sessions
- **threads** (light-wight processes) → **resource sharing** (of the same process)
 - ⇒ use them when working on the same process' data is needed
 - e.g. parallel, cooperative computations (see examples above)

Threads vs. Processes

- which are better and for which kind of applications
- **processes** → **isolation**, i.e. totally separated by one another → protection
 - ⇒ use them when isolation and protection needed
 - e.g. tabs in an Internet browser
 - e.g. document tabs in a PDF reader
 - e.g. Web server handling client sessions
- **threads** (light-wight processes) → **resource sharing** (of the same process)
 - ⇒ use them when working on the same process' data is needed
 - e.g. parallel, cooperative computations (see examples above)

Questions (1)

On a computing system with **3 logical processors**, specify the **number of threads** needed to **get the best performance** for:

- ① checking if N given numbers are prime or not, using an algorithm that divide the N numbers in equal subsets among the threads.
- ② calculating the N -th element of the array generated by the rule: $x(n) = x(n-1) + x(n-2)$, $x(0)=0$, $x(1)=1$.

Questions (2)

On a computing system with **1 logical processor**, a **HDD** and a **network card** that can act independently of the system's processors, specify the **number of threads** needed to **get the best performance** for:

- 1 copying a file from the local HDD to another file on the same HDD;
- 2 encrypting a file from the local HDD into another file on the same HDD;
- 3 uploading a file from the local HDD on a remote Web site.

Outline

- 1 Threads
 - Definition of Concepts
 - Thread Usage
- 2 Scheduling
 - Concepts
 - Scheduling Algorithms
- 3 Conclusions

Outline

- 1 Threads
 - Definition of Concepts
 - Thread Usage
- 2 Scheduling
 - Concepts
 - Scheduling Algorithms
- 3 Conclusions

Definitions

- **scheduler**: the OS component that **decides**
 - what process / thread will be run next and
 - for how long
- scheduling policy / algorithm
 - the policy (rules, requirements) / algorithm used by the scheduler
- mainly related to CPU(s)
 - but also consider other system resources
- scheduler's role (motivation)
 - mediates between the more competitors of a limited set of resources
- general classes of processes
 - compute-bounded (CPU-bounded), i.e. CPU intensive
 - I/O-bounded, i.e. I/O intensive

Definitions

- **scheduler**: the OS component that **decides**
 - **what process / thread will be run next** and
 - **for how long**
- scheduling policy / algorithm
 - the policy (rules, requirements) / algorithm used by the scheduler
- mainly related to CPU(s)
 - but also consider other system resources
- scheduler's role (motivation)
 - mediates between the more competitors of a limited set of resources
- general classes of processes
 - compute-bounded (CPU-bounded), i.e. CPU intensive
 - I/O-bounded, i.e. I/O intensive

Definitions

- **scheduler**: the OS component that **decides**
 - **what process / thread will be run next** and
 - **for how long**
- scheduling policy / algorithm
 - the policy (rules, requirements) / algorithm used by the scheduler
- mainly related to CPU(s)
 - but also consider other system resources
- scheduler's role (motivation)
 - mediates between the more competitors of a limited set of resources
- general classes of processes
 - compute-bounded (CPU-bounded), i.e. CPU intensive
 - I/O-bounded, i.e. I/O intensive

Definitions

- **scheduler**: the OS component that **decides**
 - **what process / thread will be run next** and
 - **for how long**
- scheduling policy / algorithm
 - the policy (rules, requirements) / algorithm used by the scheduler
- mainly related to CPU(s)
 - but also consider other system resources
- scheduler's role (motivation)
 - mediates between the more competitors of a limited set of resources
- general classes of processes
 - compute-bounded (CPU-bounded), i.e. CPU intensive
 - I/O-bounded, i.e. I/O intensive

Definitions

- **scheduler**: the OS component that **decides**
 - **what process / thread will be run next** and
 - **for how long**
- scheduling policy / algorithm
 - the policy (rules, requirements) / algorithm used by the scheduler
- mainly related to CPU(s)
 - but also consider other system resources
- scheduler's role (motivation)
 - mediates between the more competitors of a limited set of resources
- general classes of processes
 - compute-bounded (CPU-bounded), i.e. CPU intensive
 - I/O-bounded, i.e. I/O intensive

Definitions

- **scheduler**: the OS component that **decides**
 - **what process / thread will be run next** and
 - **for how long**
- scheduling policy / algorithm
 - the policy (rules, requirements) / algorithm used by the scheduler
- mainly related to CPU(s)
 - but also consider other system resources
- scheduler's role (motivation)
 - mediates between the more competitors of a limited set of resources
- general classes of processes
 - compute-bounded (CPU-bounded), i.e. CPU intensive
 - I/O-bounded, i.e. I/O intensive

Definitions

- **scheduler**: the OS component that **decides**
 - **what process / thread will be run next** and
 - **for how long**
- scheduling policy / algorithm
 - the policy (rules, requirements) / algorithm used by the scheduler
- mainly related to CPU(s)
 - but also consider other system resources
- scheduler's role (motivation)
 - mediates between the more competitors of a limited set of resources
- general classes of processes
 - compute-bounded (CPU-bounded), i.e. CPU intensive
 - I/O-bounded, i.e. I/O intensive

Definitions

- **scheduler**: the OS component that **decides**
 - **what process / thread will be run next** and
 - **for how long**
- scheduling policy / algorithm
 - the policy (rules, requirements) / algorithm used by the scheduler
- mainly related to CPU(s)
 - but also consider other system resources
- scheduler's role (motivation)
 - mediates between the more competitors of a limited set of resources
- general classes of processes
 - compute-bounded (CPU-bounded), i.e. CPU intensive
 - I/O-bounded, i.e. I/O intensive

Definitions

- **scheduler**: the OS component that **decides**
 - **what process / thread will be run next** and
 - **for how long**
- scheduling policy / algorithm
 - the policy (rules, requirements) / algorithm used by the scheduler
- mainly related to CPU(s)
 - but also consider other system resources
- scheduler's role (motivation)
 - mediates between the more competitors of a limited set of resources
- general classes of processes
 - compute-bounded (CPU-bounded), i.e. CPU intensive
 - I/O-bounded, i.e. I/O intensive

Definitions

- **scheduler**: the OS component that **decides**
 - **what process / thread will be run next** and
 - **for how long**
- scheduling policy / algorithm
 - the policy (rules, requirements) / algorithm used by the scheduler
- mainly related to CPU(s)
 - but also consider other system resources
- scheduler's role (motivation)
 - mediates between the more competitors of a limited set of resources
- general classes of processes
 - 1 compute-bounded (CPU-bounded), i.e. CPU intensive
 - 2 I/O-bounded, i.e. I/O intensive

Definitions

- **scheduler**: the OS component that **decides**
 - **what process / thread will be run next** and
 - **for how long**
- scheduling policy / algorithm
 - the policy (rules, requirements) / algorithm used by the scheduler
- mainly related to CPU(s)
 - but also consider other system resources
- scheduler's role (motivation)
 - mediates between the more competitors of a limited set of resources
- general classes of processes
 - 1 compute-bounded (CPU-bounded), i.e. CPU intensive
 - 2 I/O-bounded, i.e. I/O intensive

Definitions

- **scheduler**: the OS component that **decides**
 - **what process / thread will be run next** and
 - **for how long**
- scheduling policy / algorithm
 - the policy (rules, requirements) / algorithm used by the scheduler
- mainly related to CPU(s)
 - but also consider other system resources
- scheduler's role (motivation)
 - mediates between the more competitors of a limited set of resources
- general classes of processes
 - 1 compute-bounded (CPU-bounded), i.e. CPU intensive
 - 2 I/O-bounded, i.e. I/O intensive

Classes of Processes




Figure: Taken from Tanenbaum

DEMO: CPU-bound vs. I/O-bound Processes/Threads

- create few I/O bound threads, just to see they let the CPU idle a lot (see the percentage they use from CPU)
 - e.g. threads that display very often something on the screen
 - e.g. threads that mostly all the time read something from a file
- create just one CPU-bound thread, just to see that the CPU will be in use all the time
 - e.g. a simple, unrealistic `while(1);` thread

DEMO: CPU-bound vs. I/O-bound Processes/Threads

- create few I/O bound threads, just to see they let the CPU idle a lot (see the percentage they use from CPU)
 - e.g. threads that display very often something on the screen
 - e.g. threads that mostly all the time read something from a file
- create just one CPU-bound thread, just to see that the CPU will be in use all the time
 - e.g. a simple, unrealistic `while(1);` thread

DEMO: CPU-bound vs. I/O-bound Processes/Threads

- create few I/O bound threads, just to see they let the CPU idle a lot (see the percentage they use from CPU)
 - e.g. threads that display very often something on the screen
 - e.g. threads that mostly all the time read something from a file
- create just one CPU-bound thread, just to see that the CPU will be in use all the time
 - e.g. a simple, unrealistic `while(1);` thread

DEMO: CPU-bound vs. I/O-bound Processes/Threads

- create few I/O bound threads, just to see they let the CPU idle a lot (see the percentage they use from CPU)
 - e.g. threads that display very often something on the screen
 - e.g. threads that mostly all the time read something from a file
- create just one CPU-bound thread, just to see that the CPU will be in use all the time
 - e.g. a simple, unrealistic `while(1);` thread

DEMO: CPU-bound vs. I/O-bound Processes/Threads

- create few I/O bound threads, just to see they let the CPU idle a lot (see the percentage they use from CPU)
 - e.g. threads that display very often something on the screen
 - e.g. threads that mostly all the time read something from a file
- create just one CPU-bound thread, just to see that the CPU will be in use all the time
 - e.g. a simple, unrealistic `“while(1);”` thread

Questions (3)

Which threads in the following code are CPU-bound and IO-bound respectively?

```
thread_1 (int n)
{
    int i, x;
    for (x = 1, i = 1; i <= n; i++)
        x = x * i;
}

thread_2 ()
{
    char c;
    while (1) {
        read (0, &c, 1);
        write (1, &c, 1);
    }
}
```

Scheduling Situations

- any time something is changed regarding the CPU's competitors
 - process creation
 - process termination
 - process blocking, e.g. sleeping for a while, waiting for an event etc.
- external events
 - generally, interrupt occurrences
 - particularly, the clock interrupt occurrence

These events need to implement CPU (time) sharing between processes

Scheduling Situations

- any time something is changed regarding the CPU's competitors
 - **process creation**
 - **process termination**
 - **process blocking**, e.g. sleeping for a while, waiting for an event etc.
- external events
 - generally, **interrupt occurrences**
 - particularly, the **clock interrupt occurrence**

→ OS must react to implement CPU (time) sharing between processes

Scheduling Situations

- any time something is changed regarding the CPU's competitors
 - **process creation**
 - **process termination**
 - **process blocking**, e.g. sleeping for a while, waiting for an event etc.
- external events
 - generally, **interrupt occurrences**
 - particularly, the **clock interrupt occurrence**

→ OS must be able to implement CPU (time) sharing between processes

Scheduling Situations

- any time something is changed regarding the CPU's competitors
 - **process creation**
 - **process termination**
 - **process blocking**, e.g. sleeping for a while, waiting for an event etc.
 - external events
 - generally, **interrupt occurrences**
 - particularly, the **clock interrupt occurrence**
- OS must react to these events by changing CPU timing during scheduled process

Scheduling Situations

- any time something is changed regarding the CPU's competitors
 - **process creation**
 - **process termination**
 - **process blocking**, e.g. sleeping for a while, waiting for an event etc.
- external events
 - generally, **interrupt occurrences**
 - particularly, the **clock interrupt occurrence**
 - used to implement CPU (time) sharing between processes

Scheduling Situations

- any time something is changed regarding the CPU's competitors
 - **process creation**
 - **process termination**
 - **process blocking**, e.g. sleeping for a while, waiting for an event etc.
- external events
 - generally, **interrupt occurrences**
 - particularly, the **clock interrupt occurrence**
 - used to implement CPU (time) sharing between processes

Scheduling Situations

- any time something is changed regarding the CPU's competitors
 - **process creation**
 - **process termination**
 - **process blocking**, e.g. sleeping for a while, waiting for an event etc.
- external events
 - generally, **interrupt occurrences**
 - particularly, the **clock interrupt occurrence**
 - used to implement CPU (time) sharing between processes

Scheduling Situations

- any time something is changed regarding the CPU's competitors
 - **process creation**
 - **process termination**
 - **process blocking**, e.g. sleeping for a while, waiting for an event etc.
- external events
 - generally, **interrupt occurrences**
 - particularly, the **clock interrupt occurrence**
 - used to implement CPU (time) sharing between processes

Preemptive vs. Non-preemptive Scheduler

- depends on the way the clock interrupt occurrence is used by scheduler
- clock interrupt helps OS keeping track of time passage
- from this perspective a scheduler could be

- **non-preemptive**

• a process continues its execution until completion / finishing work
• CPU is given to it until it finishes its work
• i.e. a process could keep CPU a pending process

- **preemptive**

• interrupts a process if it is still in process the CPU is switched to another process i.e. takes the CPU "by force" (not by request) from the current process even if it will finish its work
• resumes later than previously suspended process

Preemptive vs. Non-preemptive Scheduler

- depends on the way the clock interrupt occurrence is used by scheduler
- clock interrupt helps OS keeping track of time passage
- from this perspective a scheduler could be

- **non-preemptive**

• a process continues its execution until completion / finishing
 • CPU is given to it
 • when a process ends then CPU is released

- **preemptive**

• a process is given a time slice, when the CPU is available
 • when the time slice expires the CPU is taken from the process and given to the next process even if it will finish its execution later than previously scheduled process

Preemptive vs. Non-preemptive Scheduler

- depends on the way the clock interrupt occurrence is used by scheduler
- clock interrupt helps OS keeping track of time passage
- from this perspective a scheduler could be
 - **non-preemptive**
 - lets a process continue its execution until completion/blocking once CPU is given to it
 - i.e. a process could keep CPU if needed
 - **preemptive**
 - suspends a process after a while to switch the CPU to another process
 - i.e. takes the CPU "by force" (yet transparently) from the current process even if it still needs it
 - resumes later the previously suspended process

Preemptive vs. Non-preemptive Scheduler

- depends on the way the clock interrupt occurrence is used by scheduler
- clock interrupt helps OS keeping track of time passage
- from this perspective a scheduler could be
 - **non-preemptive**
 - lets a process continue its execution until completion/blocking once CPU is given to it
 - i.e. a process could keep CPU if needed
 - **preemptive**
 - suspends a process after a while to switch the CPU to another process
 - i.e. takes the CPU "by force" (yet transparently) from the current process even if it still needs it
 - resumes later the previously suspended process

Preemptive vs. Non-preemptive Scheduler

- depends on the way the clock interrupt occurrence is used by scheduler
- clock interrupt helps OS keeping track of time passage
- from this perspective a scheduler could be
 - **non-preemptive**
 - lets a process continue its execution until completion/blocking once CPU is given to it
 - i.e. a process could keep CPU if needed
 - **preemptive**
 - suspends a process after a while to switch the CPU to another process
 - i.e. takes the CPU "by force" (yet transparently) from the current process even if it still needs it
 - resumes later the previously suspended process

Preemptive vs. Non-preemptive Scheduler

- depends on the way the clock interrupt occurrence is used by scheduler
- clock interrupt helps OS keeping track of time passage
- from this perspective a scheduler could be
 - **non-preemptive**
 - lets a process continue its execution until completion/blocking once CPU is given to it
 - i.e. a process could keep CPU if needed
 - **preemptive**
 - suspends a process after a while to switch the CPU to another process
 - i.e. takes the CPU "by force" (yet transparently) from the current process even if it still needs it
 - resumes later the previously suspended process

Preemptive vs. Non-preemptive Scheduler

- depends on the way the clock interrupt occurrence is used by scheduler
- clock interrupt helps OS keeping track of time passage
- from this perspective a scheduler could be
 - **non-preemptive**
 - lets a process continue its execution until completion/blocking once CPU is given to it
 - i.e. a process could keep CPU if needed
 - **preemptive**
 - suspends a process after a while to switch the CPU to another process
 - i.e. takes the CPU “by force” (yet transparently) from the current process even if it still needs it
 - resumes later the previously suspended process

Preemptive vs. Non-preemptive Scheduler

- depends on the way the clock interrupt occurrence is used by scheduler
- clock interrupt helps OS keeping track of time passage
- from this perspective a scheduler could be
 - **non-preemptive**
 - lets a process continue its execution until completion/blocking once CPU is given to it
 - i.e. a process could keep CPU if needed
 - **preemptive**
 - suspends a process after a while to switch the CPU to another process
 - i.e. takes the CPU “by force” (yet transparently) from the current process even if it still needs it
 - resumes later the previously suspended process

Preemptive vs. Non-preemptive Scheduler

- depends on the way the clock interrupt occurrence is used by scheduler
- clock interrupt helps OS keeping track of time passage
- from this perspective a scheduler could be
 - **non-preemptive**
 - lets a process continue its execution until completion/blocking once CPU is given to it
 - i.e. a process could keep CPU if needed
 - **preemptive**
 - suspends a process after a while to switch the CPU to another process
 - i.e. takes the CPU “by force” (yet transparently) from the current process even if it still needs it
 - resumes later the previously suspended process

Preemptive vs. Non-preemptive Scheduler

- depends on the way the clock interrupt occurrence is used by scheduler
- clock interrupt helps OS keeping track of time passage
- from this perspective a scheduler could be
 - **non-preemptive**
 - lets a process continue its execution until completion/blocking once CPU is given to it
 - i.e. a process could keep CPU if needed
 - **preemptive**
 - suspends a process after a while to switch the CPU to another process
 - i.e. takes the CPU “by force” (yet transparently) from the current process even if it still needs it
 - resumes later the previously suspended process

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)

- do not interact with the users, usually performing long running jobs
- their users normally expect **reasonable termination time**
- **fit better non-preemptive scheduling algorithms**
- or preemptive algorithms with long time quanta
- reduce the no of process switches \Rightarrow increase overall performance

- **interactive systems** (processes)

- interact with users
- their users expect **good response time**
- **fit only preemptive algorithms**

- **real-time systems** (processes)

- must perform certain actions in limited time
- i.e. they must meet their deadlines
- e.g. react to an external event with a maximum delay
- **need specific scheduling algorithms and system characteristics**

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)

- do not interact with the users, usually performing long running jobs
- their users normally expect **reasonable termination time**
- **fit better non-preemptive scheduling algorithms**
- or preemptive algorithms with long time quanta
- reduce the no of process switches \Rightarrow increase overall performance

- **interactive systems** (processes)

- interact with users
- their users expect **good response time**
- **fit only preemptive algorithms**

- **real-time systems** (processes)

- must perform certain actions in limited time
- i.e. they must meet their deadlines
- e.g. react to an external event with a maximum delay
- **need specific scheduling algorithms and system characteristics**

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)

- do not interact with the users, usually performing long running jobs
- their users normally expect **reasonable termination time**
- **fit better non-preemptive scheduling algorithms**
- or preemptive algorithms with long time quanta
- reduce the no of process switches \Rightarrow increase overall performance

- **interactive systems** (processes)

- interact with users
- their users expect **good response time**
- **fit only preemptive algorithms**

- **real-time systems** (processes)

- must perform certain actions in limited time
- i.e. they must meet their deadlines
- e.g. react to an external event with a maximum delay
- **need specific scheduling algorithms and system characteristics**

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)

- do not interact with the users, usually performing long running jobs
- their users normally expect **reasonable termination time**
- **fit better non-preemptive scheduling algorithms**
- or preemptive algorithms with long time quanta
- reduce the no of process switches \Rightarrow increase overall performance

- **interactive systems** (processes)

- interact with users
- their users expect **good response time**
- **fit only preemptive algorithms**

- **real-time systems** (processes)

- must perform certain actions in limited time
- i.e. they must meet their deadlines
- e.g. react to an external event with a maximum delay
- **need specific scheduling algorithms and system characteristics**

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)

- do not interact with the users, usually performing long running jobs
- their users normally expect **reasonable termination time**
- **fit better non-preemptive scheduling algorithms**
- or preemptive algorithms with long time quanta
- reduce the no of process switches \Rightarrow increase overall performance

- **interactive systems** (processes)

- interact with users
- their users expect **good response time**
- **fit only preemptive algorithms**

- **real-time systems** (processes)

- must perform certain actions in limited time
- i.e. they must meet their deadlines
- e.g. react to an external event with a maximum delay
- **need specific scheduling algorithms and system characteristics**

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)

- do not interact with the users, usually performing long running jobs
- their users normally expect **reasonable termination time**
- **fit better non-preemptive scheduling algorithms**
- or preemptive algorithms with long time quanta
- reduce the no of process switches \Rightarrow increase overall performance

- **interactive systems** (processes)

- interact with users
- their users expect **good response time**
- **fit only preemptive algorithms**

- **real-time systems** (processes)

- must perform certain actions in limited time
- i.e. they must meet their deadlines
- e.g. react to an external event with a maximum delay
- **need specific scheduling algorithms and system characteristics**

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)

- do not interact with the users, usually performing long running jobs
- their users normally expect **reasonable termination time**
- **fit better non-preemptive scheduling algorithms**
- or preemptive algorithms with long time quanta
- reduce the no of process switches \Rightarrow increase overall performance

- **interactive systems** (processes)

- interact with users
- their users expect **good response time**
- **fit only preemptive algorithms**

- **real-time systems** (processes)

- must perform certain actions in limited time
- i.e. they must meet their deadlines
- e.g. react to an external event with a maximum delay
- **need specific scheduling algorithms and system characteristics**

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)

- do not interact with the users, usually performing long running jobs
- their users normally expect **reasonable termination time**
- **fit better non-preemptive scheduling algorithms**
- or preemptive algorithms with long time quanta
- reduce the no of process switches \Rightarrow increase overall performance

- **interactive systems** (processes)

- interact with users
- their users expect **good response time**
- **fit only preemptive algorithms**

- **real-time systems** (processes)

- must perform certain actions in limited time
- i.e. they must meet their deadlines
- e.g. react to an external event with a maximum delay
- **need specific scheduling algorithms and system characteristics**

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)

- do not interact with the users, usually performing long running jobs
- their users normally expect **reasonable termination time**
- **fit better non-preemptive scheduling algorithms**
- or preemptive algorithms with long time quanta
- reduce the no of process switches \Rightarrow increase overall performance

- **interactive systems** (processes)

- interact with users
- their users expect **good response time**
- **fit only preemptive algorithms**

- **real-time systems** (processes)

- must perform certain actions in limited time
- i.e. they must meet their deadlines
- e.g. react to an external event with a maximum delay
- **need specific scheduling algorithms and system characteristics**

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)

- do not interact with the users, usually performing long running jobs
- their users normally expect **reasonable termination time**
- **fit better non-preemptive scheduling algorithms**
- or preemptive algorithms with long time quanta
- reduce the no of process switches \Rightarrow increase overall performance

- **interactive systems** (processes)

- interact with users
- their users expect **good response time**
- **fit only preemptive algorithms**

- **real-time systems** (processes)

- must perform certain actions in limited time
- i.e. they must meet their deadlines
- e.g. react to an external event with a maximum delay
- need specific scheduling algorithms and system characteristics

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)
 - do not interact with the users, usually performing long running jobs
 - their users normally expect **reasonable termination time**
 - **fit better non-preemptive scheduling algorithms**
 - or preemptive algorithms with long time quanta
 - reduce the no of process switches \Rightarrow increase overall performance
- **interactive systems** (processes)
 - interact with users
 - their users expect **good response time**
 - **fit only preemptive algorithms**
- **real-time systems** (processes)
 - must perform certain actions in limited time
 - i.e. they must meet their deadlines
 - e.g. react to an external event with a maximum delay
 - **need specific scheduling algorithms** and system characteristics

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)
 - do not interact with the users, usually performing long running jobs
 - their users normally expect **reasonable termination time**
 - **fit better non-preemptive scheduling algorithms**
 - or preemptive algorithms with long time quanta
 - reduce the no of process switches \Rightarrow increase overall performance
- **interactive systems** (processes)
 - interact with users
 - their users expect **good response time**
 - **fit only preemptive algorithms**
- **real-time systems** (processes)
 - must perform certain actions in limited time
 - i.e. they must meet their deadlines
 - e.g. react to an external event with a maximum delay
 - **need specific scheduling algorithms** and system characteristics

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)
 - do not interact with the users, usually performing long running jobs
 - their users normally expect **reasonable termination time**
 - **fit better non-preemptive scheduling algorithms**
 - or preemptive algorithms with long time quanta
 - reduce the no of process switches \Rightarrow increase overall performance
- **interactive systems** (processes)
 - interact with users
 - their users expect **good response time**
 - **fit only preemptive algorithms**
- **real-time systems** (processes)
 - must perform certain actions in limited time
 - i.e. they must meet their deadlines
 - e.g. react to an external event with a maximum delay
 - **need specific scheduling algorithms** and system characteristics

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)
 - do not interact with the users, usually performing long running jobs
 - their users normally expect **reasonable termination time**
 - **fit better non-preemptive scheduling algorithms**
 - or preemptive algorithms with long time quanta
 - reduce the no of process switches \Rightarrow increase overall performance
- **interactive systems** (processes)
 - interact with users
 - their users expect **good response time**
 - **fit only preemptive algorithms**
- **real-time systems** (processes)
 - must perform certain actions in limited time
 - i.e. they must meet their deadlines
 - e.g. react to an external event with a maximum delay
 - **need specific scheduling algorithms and system characteristics**

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)
 - do not interact with the users, usually performing long running jobs
 - their users normally expect **reasonable termination time**
 - **fit better non-preemptive scheduling algorithms**
 - or preemptive algorithms with long time quanta
 - reduce the no of process switches \Rightarrow increase overall performance
- **interactive systems** (processes)
 - interact with users
 - their users expect **good response time**
 - **fit only preemptive algorithms**
- **real-time systems** (processes)
 - must perform certain actions in limited time
 - i.e. they must meet their deadlines
 - e.g. react to an external event with a maximum delay
 - **need specific scheduling algorithms** and system characteristics

Outline

- 1 Threads
 - Definition of Concepts
 - Thread Usage
- 2 Scheduling
 - Concepts
 - Scheduling Algorithms
- 3 Conclusions

First-Come First-Served (FCFS)

- it is a non-preemptive algorithm
- simple to understand and implement
- the average waiting time (a.w.t.) depends on the thread order and execution time
 - see Gantt diagram, i.e. graphical illustration of chronological time intervals each thread was scheduled

First-Come First-Served (FCFS)

- it is a non-preemptive algorithm
- simple to understand and implement
- the average waiting time (a.w.t.) depends on the thread order and execution time
 - see Gantt diagram, i.e. graphical illustration of chronological time intervals each thread was scheduled

First-Come First-Served (FCFS)

- it is a non-preemptive algorithm
- simple to understand and implement
- the average waiting time (a.w.t.) depends on the thread order and execution time
 - see Gantt diagram, i.e. graphical illustration of chronological time intervals each thread was scheduled

First-Come First-Served (FCFS)

- it is a non-preemptive algorithm
- simple to understand and implement
- the average waiting time (a.w.t.) depends on the thread order and execution time
 - see Gantt diagram, i.e. graphical illustration of chronological time intervals each thread was scheduled

First-Come First-Served (FCFS). Example 1

- T1:23, T2:2, T3:3 (thread : execution time)
- waiting time for T1 is 0 ms, for T2 is 23 ms, and for T3 is 25 ms
- $a.w.t. = \frac{0+23+25}{3} = 16ms$




Figure: Gantt diagram of process scheduling

First-Come First-Served (FCFS). Example 1

- T1:23, T2:2, T3:3 (thread : execution time)
- waiting time for T1 is 0 ms, for T2 is 23 ms, and for T3 is 25 ms
- $a.w.t. = \frac{0+23+25}{3} = 16ms$




Figure: Gantt diagram of process scheduling

First-Come First-Served (FCFS). Example 1

- T1:23, T2:2, T3:3 (thread : execution time)
- waiting time for T1 is 0 ms, for T2 is 23 ms, and for T3 is 25 ms
- $a.w.t. = \frac{0+23+25}{3} = 16ms$




Figure: Gantt diagram of process scheduling

First-Come First-Served (FCFS). Example 2

- T2:2, T3:3, T1:23 (thread : execution time)
- waiting time for T1 is 5 ms, for T2 is 0 ms, and for T3 is 2 ms
- $a.w.t. = \frac{5+0+2}{3} = 2.33ms$



Figure: Gantt diagram of process scheduling

First-Come First-Served (FCFS). Example 2

- T2:2, T3:3, T1:23 (thread : execution time)
- waiting time for T1 is 5 ms, for T2 is 0 ms, and for T3 is 2 ms
- $a.w.t. = \frac{5+0+2}{3} = 2.33ms$



Figure: Gantt diagram of process scheduling

First-Come First-Served (FCFS). Example 2

- T2:2, T3:3, T1:23 (thread : execution time)
- waiting time for T1 is 5 ms, for T2 is 0 ms, and for T3 is 2 ms
- $a.w.t. = \frac{5+0+2}{3} = 2.33ms$



Figure: Gantt diagram of process scheduling

Shortest Job First (SJF)

- runtime should be known in advance
- is the optimal algorithm when all the ready threads are available simultaneously
- non-preemptive
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5 (submit time : process : run time)
 - a.w.t. = $(0 + (8-1) + (17-3) + (12-2))/4 = 7.75$
- preemptive (*shortest remaining time next*)
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5
 - a.w.t. = $((10-1)+(1-1)+(17-2)+(5-3))/4 = 6,5$

Shortest Job First (SJF)

- runtime should be known in advance
- is the optimal algorithm when all the ready threads are available simultaneously
- non-preemptive
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5 (submit time : process : run time)
 - a.w.t. = $(0 + (8-1) + (17-3) + (12-2))/4 = 7.75$
- preemptive (*shortest remaining time next*)
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5
 - a.w.t. = $((10-1)+(1-1)+(17-2)+(5-3))/4 = 6,5$

Shortest Job First (SJF)

- runtime should be known in advance
- is the optimal algorithm when all the ready threads are available simultaneously
- non-preemptive
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5 (submit time : process : run time)
 - a.w.t. = $(0 + (8-1) + (17-3) + (12-2))/4 = 7.75$
- preemptive (*shortest remaining time next*)
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5
 - a.w.t. = $((10-1)+(1-1)+(17-2)+(5-3))/4 = 6,5$

Shortest Job First (SJF)

- runtime should be known in advance
- is the optimal algorithm when all the ready threads are available simultaneously
- non-preemptive
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5 (submit time : process : run time)
 - a.w.t. = $(0 + (8-1) + (17-3) + (12-2))/4 = 7.75$
- preemptive (*shortest remaining time next*)
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5
 - a.w.t. = $((10-1)+(1-1)+(17-2)+(5-3))/4 = 6,5$

Shortest Job First (SJF)

- runtime should be known in advance
- is the optimal algorithm when all the ready threads are available simultaneously
- non-preemptive
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5 (submit time : process : run time)
 - a.w.t. = $(0 + (8-1) + (17-3) + (12-2))/4 = 7.75$
- preemptive (*shortest remaining time next*)
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5
 - a.w.t. = $((10-1)+(1-1)+(17-2)+(5-3))/4 = 6,5$

Shortest Job First (SJF)

- runtime should be known in advance
- is the optimal algorithm when all the ready threads are available simultaneously
- non-preemptive
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5 (submit time : process : run time)
 - a.w.t. = $(0 + (8-1) + (17-3) + (12-2))/4 = 7.75$
- preemptive (*shortest remaining time next*)
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5
 - a.w.t. = $((10-1)+(1-1)+(17-2)+(5-3))/4 = 6,5$

Shortest Job First (SJF)

- runtime should be known in advance
- is the optimal algorithm when all the ready threads are available simultaneously
- non-preemptive
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5 (submit time : process : run time)
 - a.w.t. = $(0 + (8-1) + (17-3) + (12-2))/4 = 7.75$
- preemptive (*shortest remaining time next*)
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5
 - a.w.t. = $((10-1)+(1-1)+(17-2)+(5-3))/4 = 6,5$

Shortest Job First (SJF)

- runtime should be known in advance
- is the optimal algorithm when all the ready threads are available simultaneously
- non-preemptive
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5 (submit time : process : run time)
 - a.w.t. = $(0 + (8-1) + (17-3) + (12-2))/4 = 7.75$
- preemptive (*shortest remaining time next*)
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5
 - a.w.t. = $((10-1)+(1-1)+(17-2)+(5-3))/4 = 6,5$

Round-Robin (RR)

- each thread is assigned a time interval
 - time quantum or slice
- it is a preemptive algorithm
- maintain a FIFO list for ready threads
 - like FCFS, but preemptive, based on time quanta
- the length of the time quantum
 - too short \Rightarrow many CPU switches \Rightarrow lower the CPU efficiency, i.e. a lot of CPU is wasted for running the context switching code
 - too large \Rightarrow longer wait times in ready queue \Rightarrow poor response time for interactive threads
 - 20–100 msec is a reasonable compromise

Round-Robin (RR)

- each thread is assigned a time interval
 - time quantum or slice
- it is a preemptive algorithm
- maintain a FIFO list for ready threads
 - like FCFS, but preemptive, based on time quanta
- the length of the time quantum
 - too short \Rightarrow many CPU switches \Rightarrow lower the CPU efficiency, i.e. a lot of CPU is wasted for running the context switching code
 - too large \Rightarrow longer wait times in ready queue \Rightarrow poor response time for interactive threads
 - 20–100 msec is a reasonable compromise

Round-Robin (RR)

- each thread is assigned a time interval
 - time quantum or slice
- it is a preemptive algorithm
- maintain a FIFO list for ready threads
 - like FCFS, but preemptive, based on time quanta
- the length of the time quantum
 - too short \Rightarrow many CPU switches \Rightarrow lower the CPU efficiency, i.e. a lot of CPU is wasted for running the context switching code
 - too large \Rightarrow longer wait times in ready queue \Rightarrow poor response time for interactive threads
 - 20–100 msec is a reasonable compromise

Round-Robin (RR)

- each thread is assigned a time interval
 - time quantum or slice
- it is a preemptive algorithm
- maintain a FIFO list for ready threads
 - like FCFS, but preemptive, based on time quanta
- the length of the time quantum
 - too short \Rightarrow many CPU switches \Rightarrow lower the CPU efficiency, i.e. a lot of CPU is wasted for running the context switching code
 - too large \Rightarrow longer wait times in ready queue \Rightarrow poor response time for interactive threads
 - 20–100 msec is a reasonable compromise

Round-Robin (RR)

- each thread is assigned a time interval
 - time quantum or slice
- it is a preemptive algorithm
- maintain a FIFO list for ready threads
 - like FCFS, but preemptive, based on time quanta
- the length of the time quantum
 - too short \Rightarrow many CPU switches \Rightarrow lower the CPU efficiency, i.e. a lot of CPU is wasted for running the context switching code
 - too large \Rightarrow longer wait times in ready queue \Rightarrow poor response time for interactive threads
 - 20–100 msec is a reasonable compromise

Round-Robin (RR)

- each thread is assigned a time interval
 - time quantum or slice
- it is a preemptive algorithm
- maintain a FIFO list for ready threads
 - like FCFS, but preemptive, based on time quanta
- the length of the time quantum
 - too short \Rightarrow many CPU switches \Rightarrow lower the CPU efficiency, i.e. a lot of CPU is wasted for running the context switching code
 - too large \Rightarrow longer wait times in ready queue \Rightarrow poor response time for interactive threads
 - 20–100 msec is a reasonable compromise

Round-Robin (RR)

- each thread is assigned a time interval
 - time quantum or slice
- it is a preemptive algorithm
- maintain a FIFO list for ready threads
 - like FCFS, but preemptive, based on time quanta
- the length of the time quantum
 - too short \Rightarrow many CPU switches \Rightarrow lower the CPU efficiency, i.e. a lot of CPU is wasted for running the context switching code
 - too large \Rightarrow longer wait times in ready queue \Rightarrow poor response time for interactive threads
 - 20–100 msec is a reasonable compromise

Round-Robin (RR)

- each thread is assigned a time interval
 - time quantum or slice
- it is a preemptive algorithm
- maintain a FIFO list for ready threads
 - like FCFS, but preemptive, based on time quanta
- the length of the time quantum
 - too short \Rightarrow many CPU switches \Rightarrow lower the CPU efficiency, i.e. a lot of CPU is wasted for running the context switching code
 - too large \Rightarrow longer wait times in ready queue \Rightarrow poor response time for interactive threads
 - 20–100 msec is a reasonable compromise

Round-Robin (RR)

- each thread is assigned a time interval
 - time quantum or slice
- it is a preemptive algorithm
- maintain a FIFO list for ready threads
 - like FCFS, but preemptive, based on time quanta
- the length of the time quantum
 - too short \Rightarrow many CPU switches \Rightarrow lower the CPU efficiency, i.e. a lot of CPU is wasted for running the context switching code
 - too large \Rightarrow longer wait times in ready queue \Rightarrow poor response time for interactive threads
 - 20–100 msec is a reasonable compromise

Priority Based

- each thread has a priority assigned
 - usually a number that describes the thread “importance” (from some perspective)
 - e.g. running time could be such a number: the smaller the higher thread’s priority \Rightarrow SJF
- rule: **the ready thread with the greatest priority is always run**
- give a chance also to smaller priority threads
 - i.e. **avoid starvation**
 - by changing periodically the priority of processes
- priority-based policies
 - fixed priorities
 - dynamically adjusted priorities
- priority classes
 - there could be more processes with the same priority
 - use priority scheduling between classes
 - use another scheduling algorithm with each class, e.g. round-robin

Priority Based

- each thread has a priority assigned
 - usually a number that describes the thread “importance” (from some perspective)
 - e.g. running time could be such a number: the smaller the higher thread's priority \Rightarrow SJF
- rule: **the ready thread with the greatest priority is always run**
- give a chance also to smaller priority threads
 - i.e. **avoid starvation**
 - by changing periodically the priority of processes
- priority-based policies
 - fixed priorities
 - dynamically adjusted priorities
- priority classes
 - there could be more processes with the same priority
 - use priority scheduling between classes
 - use another scheduling algorithm with each class, e.g. round-robin

Priority Based

- each thread has a priority assigned
 - usually a number that describes the thread “importance” (from some perspective)
 - e.g. running time could be such a number: the smaller the higher thread’s priority \Rightarrow SJF
- rule: **the ready thread with the greatest priority is always run**
- give a chance also to smaller priority threads
 - i.e. **avoid starvation**
 - by changing periodically the priority of processes
- priority-based policies
 - fixed priorities
 - dynamically adjusted priorities
- priority classes
 - there could be more processes with the same priority
 - use priority scheduling between classes
 - use another scheduling algorithm with each class, e.g. round-robin

Priority Based

- each thread has a priority assigned
 - usually a number that describes the thread “importance” (from some perspective)
 - e.g. running time could be such a number: the smaller the higher thread’s priority \Rightarrow SJF
- rule: **the ready thread with the greatest priority is always run**
- give a chance also to smaller priority threads
 - i.e. avoid starvation
 - by changing periodically the priority of processes
- priority-based policies
 - fixed priorities
 - dynamically adjusted priorities
- priority classes
 - there could be more processes with the same priority
 - use priority scheduling between classes
 - use another scheduling algorithm with each class, e.g. round-robin

Priority Based

- each thread has a priority assigned
 - usually a number that describes the thread “importance” (from some perspective)
 - e.g. running time could be such a number: the smaller the higher thread’s priority \Rightarrow SJF
- rule: **the ready thread with the greatest priority is always run**
- give a chance also to smaller priority threads
 - i.e. **avoid starvation**
 - by changing periodically the priority of processes
- priority-based policies
 - fixed priorities
 - dynamically adjusted priorities
- priority classes
 - there could be more processes with the same priority
 - use priority scheduling between classes
 - use another scheduling algorithm with each class, e.g. round-robin

Priority Based

- each thread has a priority assigned
 - usually a number that describes the thread “importance” (from some perspective)
 - e.g. running time could be such a number: the smaller the higher thread’s priority \Rightarrow SJF
- rule: **the ready thread with the greatest priority is always run**
- give a chance also to smaller priority threads
 - i.e. **avoid starvation**
 - by changing periodically the priority of processes
- priority-based policies
 - fixed priorities
 - dynamically adjusted priorities
- priority classes
 - there could be more processes with the same priority
 - use priority scheduling between classes
 - use another scheduling algorithm with each class, e.g. round-robin

Priority Based

- each thread has a priority assigned
 - usually a number that describes the thread “importance” (from some perspective)
 - e.g. running time could be such a number: the smaller the higher thread’s priority \Rightarrow SJF
- rule: **the ready thread with the greatest priority is always run**
- give a chance also to smaller priority threads
 - i.e. **avoid starvation**
 - by changing periodically the priority of processes
- priority-based policies
 - fixed priorities
 - dynamically adjusted priorities
- priority classes
 - there could be more processes with the same priority
 - use priority scheduling between classes
 - use another scheduling algorithm with each class, e.g. round-robin

Priority Based

- each thread has a priority assigned
 - usually a number that describes the thread “importance” (from some perspective)
 - e.g. running time could be such a number: the smaller the higher thread’s priority \Rightarrow SJF
- rule: **the ready thread with the greatest priority is always run**
- give a chance also to smaller priority threads
 - i.e. **avoid starvation**
 - by changing periodically the priority of processes
- priority-based policies
 - fixed priorities
 - dynamically adjusted priorities
- priority classes
 - there could be more processes with the same priority
 - use priority scheduling between classes
 - use another scheduling algorithm with each class, e.g. round-robin

Priority Based

- each thread has a priority assigned
 - usually a number that describes the thread “importance” (from some perspective)
 - e.g. running time could be such a number: the smaller the higher thread’s priority \Rightarrow SJF
- rule: **the ready thread with the greatest priority is always run**
- give a chance also to smaller priority threads
 - i.e. **avoid starvation**
 - by changing periodically the priority of processes
- priority-based policies
 - fixed priorities
 - dynamically adjusted priorities
- priority classes
 - there could be more processes with the same priority
 - use priority scheduling between classes
 - use another scheduling algorithm with each class, e.g. round-robin

Priority Based

- each thread has a priority assigned
 - usually a number that describes the thread “importance” (from some perspective)
 - e.g. running time could be such a number: the smaller the higher thread’s priority \Rightarrow SJF
- rule: **the ready thread with the greatest priority is always run**
- give a chance also to smaller priority threads
 - i.e. **avoid starvation**
 - by changing periodically the priority of processes
- priority-based policies
 - fixed priorities
 - dynamically adjusted priorities
- priority classes
 - there could be more processes with the same priority
 - use priority scheduling between classes
 - use another scheduling algorithm with each class, e.g. round-robin

Priority Based

- each thread has a priority assigned
 - usually a number that describes the thread “importance” (from some perspective)
 - e.g. running time could be such a number: the smaller the higher thread’s priority \Rightarrow SJF
- rule: **the ready thread with the greatest priority is always run**
- give a chance also to smaller priority threads
 - i.e. **avoid starvation**
 - by changing periodically the priority of processes
- priority-based policies
 - fixed priorities
 - dynamically adjusted priorities
- priority classes
 - there could be more processes with the same priority
 - use priority scheduling between classes
 - use another scheduling algorithm with each class, e.g. round-robin

Priority Based

- each thread has a priority assigned
 - usually a number that describes the thread “importance” (from some perspective)
 - e.g. running time could be such a number: the smaller the higher thread’s priority \Rightarrow SJF
- rule: **the ready thread with the greatest priority is always run**
- give a chance also to smaller priority threads
 - i.e. **avoid starvation**
 - by changing periodically the priority of processes
- priority-based policies
 - fixed priorities
 - dynamically adjusted priorities
- priority classes
 - there could be more processes with the same priority
 - use priority scheduling between classes
 - use another scheduling algorithm with each class, e.g. round-robin

Priority Based

- each thread has a priority assigned
 - usually a number that describes the thread “importance” (from some perspective)
 - e.g. running time could be such a number: the smaller the higher thread’s priority \Rightarrow SJF
- rule: **the ready thread with the greatest priority is always run**
- give a chance also to smaller priority threads
 - i.e. **avoid starvation**
 - by changing periodically the priority of processes
- priority-based policies
 - fixed priorities
 - dynamically adjusted priorities
- priority classes
 - there could be more processes with the same priority
 - use priority scheduling between classes
 - use another scheduling algorithm with each class, e.g. round-robin

Priority Based

- each thread has a priority assigned
 - usually a number that describes the thread “importance” (from some perspective)
 - e.g. running time could be such a number: the smaller the higher thread’s priority \Rightarrow SJF
- rule: **the ready thread with the greatest priority is always run**
- give a chance also to smaller priority threads
 - i.e. **avoid starvation**
 - by changing periodically the priority of processes
- priority-based policies
 - fixed priorities
 - dynamically adjusted priorities
- priority classes
 - there could be more processes with the same priority
 - use priority scheduling between classes
 - use another scheduling algorithm with each class, e.g. round-robin

Questions (4)

Let us suppose that an OS' scheduler uses the **Round-Robin** algorithm with **200 ms** time-quantum. Illustrate on a **Gantt diagram** for the time interval **0 – 1 seconds** the executions of the threads whose code is given below, for the following scenarios.

1

```
// Thread1          // Thread2
while(1){}          while(1){}
```

2

```
// Thread1          // Thread2
usleep(300000); // 300 ms
while(1) {}        while(1) {}
```

Outline

- 1 Threads
 - Definition of Concepts
 - Thread Usage
- 2 Scheduling
 - Concepts
 - Scheduling Algorithms
- 3 Conclusions

What we talked about

- threads

- describe multiple, independent executions in the same process
- share all resources of the process they belong to

- multi-threading applications

- scheduling

- needed when there are more competitors for a limited set of resources (e.g. CPUs)
- first-come first-served, shortest job first, round-robin, priority-based

What we talked about

- threads
 - describe multiple, independent executions in the same process
 - share all resources of the process they belong to
- multi-threading applications
- scheduling
 - needed when there are more competitors for a limited set of resources (e.g. CPUs)
 - first-come first-served, shortest job first, round-robin, priority-based

What we talked about

- threads
 - describe multiple, independent executions in the same process
 - share all resources of the process they belong to
- multi-threading applications
- scheduling
 - needed when there are more competitors for a limited set of resources (e.g. CPUs)
 - first-come first-served, shortest job first, round-robin, priority-based

What we talked about

- threads
 - describe multiple, independent executions in the same process
 - share all resources of the process they belong to
- multi-threading applications
- scheduling
 - needed when there are more competitors for a limited set of resources (e.g. CPUs)
 - first-come first-served, shortest job first, round-robin, priority-based

What we talked about

- threads
 - describe multiple, independent executions in the same process
 - share all resources of the process they belong to
- multi-threading applications
- scheduling
 - needed when there are more competitors for a limited set of resources (e.g. CPUs)
 - first-come first-served, shortest job first, round-robin, priority-based

What we talked about

- threads
 - describe multiple, independent executions in the same process
 - share all resources of the process they belong to
- multi-threading applications
- scheduling
 - needed when there are more competitors for a limited set of resources (e.g. CPUs)
 - first-come first-served, shortest job first, round-robin, priority-based

What we talked about

- threads
 - describe multiple, independent executions in the same process
 - share all resources of the process they belong to
- multi-threading applications
- scheduling
 - needed when there are more competitors for a limited set of resources (e.g. CPUs)
 - first-come first-served, shortest job first, round-robin, priority-based

Lessons Learned

Multiple threads in a process are **useful** and **effective** only in particular cases:

- 1 **logical application parallelism** exists
- 2 needed **physical resources** available

Lessons Learned

Multiple threads in a process are **useful** and **effective** only in particular cases:

- 1 **logical application parallelism** exists
- 2 needed **physical resources** available