

Chapter 5.2

Process Management

Thread Model. Process Scheduling

Print Version of Lectures Notes of *Operating Systems*

Technical University of Cluj-Napoca (UTCN)
Computer Science Department

Adrian Coleşa

April 08, 2020

5.2.1

Purpose and Contents

The purpose of this subchapter

- Defines threads and the way they can be used
- Presents scheduling principles and algorithms

5.2.2

Bibliography

- A. Tanenbaum, *Modern Operating Systems*, 2nd Edition, 2001, Chapter 2, Processes, pg. 71 – 100, pg. 132 – 151

5.2.3

Contents

1	Threads	1
1.1	Definition of Concepts	1
1.2	Thread Usage	3
2	Scheduling	6
2.1	Concepts	6
2.2	Scheduling Algorithms	8
3	Conclusions	10

5.2.4

1 Threads

1.1 Definition of Concepts

Thread Definition

- describes a **sequential, independent execution** within a process
 - execution = the memory path followed in the code by the IP register
- **there could be more**
 - simultaneous and independent executions in the same process
 - \Rightarrow **threads in a process**

5.2.5

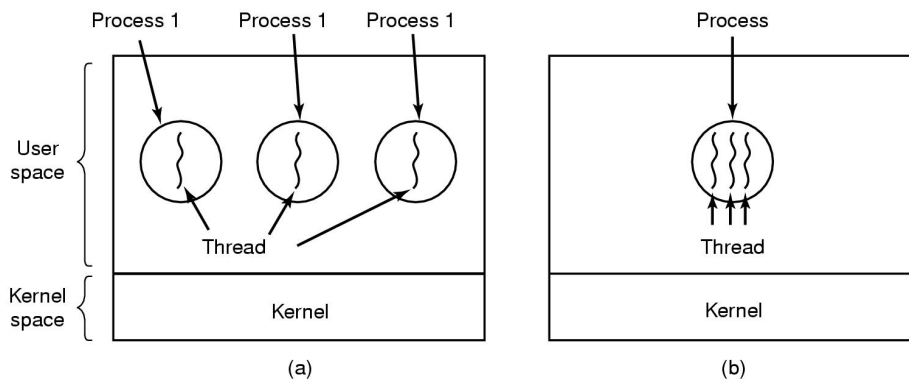


Figure 1: Taken from Tanenbaum

Modeling The Execution

- **process model**
 - models “**an entire computer**” used for executing some user application
 - composed by both
 - * process’ resources
 - * process’ execution
 - **isolates resources** of one process by that of others
- **thread model**
 - models **the processor(s)** given to a process
 - makes distinction between the two components of a process
 - * the process describes the resources
 - * the thread describes the execution
 - if multiple independent executions could be identified in a process
 - * there could be more threads of that process
 - * one thread \leftrightarrow one execution
 - threads of a process **share resources** of that process
 - threads of a process (normally) **not visible to other processes**

5.2.6

Multiprogramming and Multithreading

- multiprogramming = multiple processes managed simultaneously
- multithreading = multiple threads in the same process managed simultaneously

5.2.7

Thread’s Components. Description

- **threads** of the same process
 - **share all the resources of that process**: memory, open files etc.
 - \Rightarrow anything (i.e. inside the process) is accessible to all threads (of that process)
- each thread is an independent execution
 - \Rightarrow **each thread has also its own resources**
 - that describe the corresponding independent execution
- thread’s specific resources
 - **machine registers** (e.g. the IP register)
 - **the stack**

5.2.8

Thread’s Components. Illustration

5.2.9




Figure 2: Taken from Tanenbaum

Threads Relationship

Paraphrase a rule from “Animal Farm” by George Orwell

5.2.10

Threads Relationship (cont.)

All threads are equal, but the “main” thread is “*more equal than the others!*”

5.2.11

Consequences of “All threads are equal”


- **there is no protection between threads of the same process**
 - there is no need, actually
 - created by the application itself to execute code inside the application
- **there is no restriction on them regarding their scheduling**
 - any one could be scheduled any time, depending on the system’s scheduling policy
 - we have no control on thread scheduling
 - if needed, **we must synchronize the threads**

5.2.12

Consequences of “All threads are equal” (cont.)

- **... but the main one is “*more equal than the others!*”**
 - it is the first one created: **executes the *main()* function**
 - it is the ancestor of all other threads, i.e. initiates creation of the other threads
 - if it terminates returning from the “*main()*” function
 - * the entire process (and all threads) terminates
 - * \Rightarrow it make sense for it to **wait for the termination of all other threads**
 - **though**, if it terminates using the thread termination syscall, the other threads survive
 - * see difference between `exit()` and `pthread_exit()` in Linux
 - * see a discussion about this at <https://devblogs.microsoft.com/oldnewthing/20100827-00/?p=13023>

5.2.13






Figure 3: Taken from Tanenbaum

1.2 Thread Usage

When Do Threads Prove Themselves Useful?

- **useful** in applications
 - where **different actions can happen concurrently** (in the same time) → **logical parallelism**
 - and use **different “independent” system resources** (e.g. CPU, HDD, etc.) → **physical parallelism**
- **NOT useful** in applications
 - using only the single CPU (our considered context)
 - * though useful, if multiple CPUs are available
 - in which possible concurrent executions need all the time (compete for) the same single shared resource

5.2.14

Examples of Multi-threaded Apps: A Word Processor

5.2.15

Examples of Multi-threaded Apps: A Concurrent Server

5.2.16

Threads vs. Processes

- which are better and for which kind of applications
- **processes** → **isolation**, i.e. totally separated by one another → protection
 - ⇒ use them when isolation and protection needed
 - e.g. tabs in an Internet browser
 - e.g. document tabs in a PDF reader
 - e.g. Web server handling client sessions
- **threads** (light-wight processes) → **resource sharing** (of the same process)
 - ⇒ use them when working on the same process' data is needed
 - e.g. parallel, cooperative computations (see examples above)

5.2.17

Questions (1)

On a computing system with **3 logical processors**, specify the **number of threads** needed to **get the best performance** for:

1. checking if N given numbers are prime or not, using an algorithm that divide the N numbers in equal subsets among the threads.
2. calculating the N -th element of the array generated by the rule: $x(n) = x(n-1) + x(n-2)$, $x(0)=0$, $x(1)=1$.

5.2.18

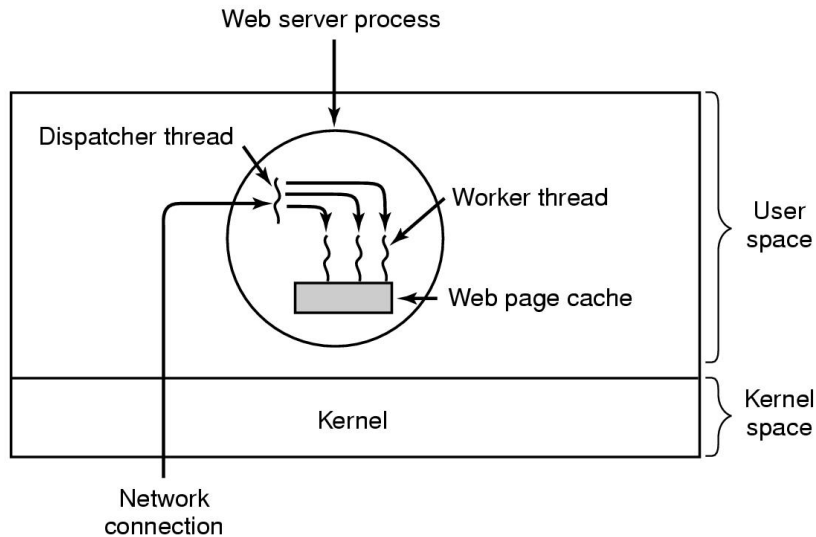


Figure 4: Taken from Tanenbaum

Questions (2)

On a computing system with **1 logical processor**, a **HDD** and a **network card** that can act independently of the system's processors, specify the **number of threads** needed to **get the best performance** for:

1. copying a file from the local HDD to another file on the same HDD;
2. encrypting a file from the local HDD into another file on the same HDD;
3. uploading a file from the local HDD on a remote Web site.

5.2.19

2 Scheduling

2.1 Concepts

Definitions

- **scheduler**: the OS component that **decides**
 - **what process / thread will be run next** and
 - **for how long**
- scheduling policy / algorithm
 - the policy (rules, requirements) / algorithm used by the scheduler
- mainly related to CPU(s)
 - but also consider other system resources
- scheduler's role (motivation)
 - mediates between the more competitors of a limited set of resources
- general classes of processes
 1. compute-bounded (CPU-bounded), i.e. CPU intensive
 2. I/O-bounded, i.e. I/O intensive

5.2.20

Classes of Processes

5.2.21




Figure 5: Taken from Tanenbaum

DEMO: CPU-bound vs. I/O-bound Processes/Threads

- create few I/O bound threads, just to see they let the CPU idle a lot (see the percentage they use from CPU)
 - e.g. threads that display very often something on the screen
 - e.g. threads that mostly all the time read something from a file
- create just one CPU-bound thread, just to see that the CPU will be in use all the time
 - e.g. a simple, unrealistic “while(1);” thread

5.2.22

Questions (3)

Which threads in the following code are CPU-bound and IO-bound respectively?

```
thread_1 (int n)
{
    int i, x;
    for (x = 1, i = 1; i <= n; i++)
        x = x * i;
}

thread_2 ()
{
    char c;
    while (1) {
        read (0, &c, 1);
        write (1, &c, 1);
    }
}
```

5.2.23

Scheduling Situations

- any time something is changed regarding the CPU’s competitors
 - **process creation**
 - **process termination**
 - **process blocking**, e.g. sleeping for a while, waiting for an event etc.
- external events
 - generally, **interrupt occurrences**
 - particularly, the **clock interrupt occurrence**
 - * used to implement CPU (time) sharing between processes

5.2.24

Preemptive vs. Non-preemptive Scheduler

- depends on the way the clock interrupt occurrence is used by scheduler
- clock interrupt helps OS keeping track of time passage
- from this perspective a scheduler could be
 - **non-preemptive**
 - * lets a process continue its execution until completion/blocking once CPU is given to it
 - * i.e. a process could keep CPU if needed
 - **preemptive**
 - * suspends a process after a while to switch the CPU to another process
 - * i.e. takes the CPU “by force” (yet transparently) from the current process even if it still needs it
 - * resumes later the previously suspended process

5.2.25

Categories of Scheduling Algorithms (Examples)

- **batch systems** (processes)
 - do not interact with the users, usually performing long running jobs
 - their users normally expect **reasonable termination time**
 - **fit better non-preemptive scheduling algorithms**
 - or preemptive algorithms with long time quanta
 - reduce the no of process switches \Rightarrow increase overall performance
- **interactive systems** (processes)
 - interact with users
 - their users expect **good response time**
 - **fit only preemptive algorithms**
- **real-time systems** (processes)
 - must perform certain actions in limited time
 - i.e. they must meet their deadlines
 - e.g. react to an external event with a maximum delay
 - **need specific scheduling algorithms** and system characteristics

5.2.26

2.2 Scheduling Algorithms

First-Come First-Served (FCFS)

- it is a non-preemptive algorithm
- simple to understand and implement
- the average waiting time (a.w.t.) depends on the thread order and execution time
 - see Gantt diagram, i.e. graphical illustration of chronological time intervals each thread was scheduled

5.2.27

First-Come First-Served (FCFS). Example 1

- T1:23, T2:2, T3:3 (thread: execution time)
- waiting time for T1 is 0 ms, for T2 is 23 ms, and for T3 is 25 ms
- $a.w.t. = \frac{0+23+25}{3} = 16ms$

5.2.28

First-Come First-Served (FCFS). Example 2

- T2:2, T3:3, T1:23 (thread: execution time)
- waiting time for T1 is 5 ms, for T2 is 0 ms, and for T3 is 2 ms
- $a.w.t. = \frac{5+0+2}{3} = 2.33ms$

5.2.29




Figure 6: Gantt diagram of process scheduling



Figure 7: Gantt diagram of process scheduling

Shortest Job First (SJF)

- runtime should be known in advance
- is the optimal algorithm when all the ready threads are available simultaneously
- non-preemptive
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5 (submit time : process : run time)
 - a.w.t. = $(0 + (8-1) + (17-3) + (12-2))/4 = 7.75$
- preemptive (*shortest remaining time next*)
 - 0:P1:8, 1:P2:4, 2:P3:9, 3:P4:5
 - a.w.t. = $((10-1)+(1-1)+(17-2)+(5-3))/4 = 6,5$

5.2.30

Round-Robin (RR)

- each thread is assigned a time interval
 - time quantum or slice
- it is a preemptive algorithm
- maintain a FIFO list for ready threads
 - like FCFS, but preemptive, based on time quanta
- the length of the time quantum
 - too short \Rightarrow many CPU switches \Rightarrow lower the CPU efficiency, i.e. a lot of CPU is wasted for running the context switching code
 - too large \Rightarrow longer wait times in ready queue \Rightarrow poor response time for interactive threads
 - 20–100 msec is a reasonable compromise

5.2.31

Priority Based

- each thread has a priority assigned
 - usually a number that describes the thread “importance” (from some perspective)
 - e.g. running time could be such a number: the smaller the higher thread’s priority \Rightarrow SJF
- rule: **the ready thread with the greatest priority is always run**
- give a chance also to smaller priority threads
 - i.e. **avoid starvation**
 - by changing periodically the priority of processes
- priority-based policies
 - fixed priorities

- dynamically adjusted priorities
- priority classes
 - there could be more processes with the same priority
 - use priority scheduling between classes
 - use another scheduling algorithm with each class, e.g. round-robin

5.2.32

Questions (4)

Let us suppose that an OS' scheduler uses the **Round-Robin** algorithm with **200 ms** time-quantum. Illustrate on a **Gantt diagram** for the time interval **0 – 1 seconds** the executions of the threads whose code is given below, for the following scenarios.

- | | |
|-------------------------|-------------------------|
| <code>// Thread1</code> | <code>// Thread2</code> |
| <code>while(1){}</code> | <code>while(1){}</code> |
- | | |
|--|--------------------------|
| <code>// Thread1</code> | <code>// Thread2</code> |
| <code>usleep(300000); // 300 ms</code> | |
| <code>while(1) {}</code> | <code>while(1) {}</code> |

5.2.33

3 Conclusions

What we talked about

- threads
 - describe multiple, independent executions in the same process
 - share all resources of the process they belong to
- multi-threading applications
- scheduling
 - needed when there are more competitors for a limited set of resources (e.g. CPUs)
 - first-come first-served, shortest job first, round-robin, priority-based

5.2.34

Lessons Learned

Multiple threads in a process are **useful** and **effective** only in particular cases:

1. **logical application parallelism** exists
2. needed **physical resources available**

5.2.35