

# Chapter 6.2

## Process and Thread Synchronization

### Semaphores and Condition Variables

Print Version of Lectures Notes of *Operating Systems*

Technical University of Cluj-Napoca (UTCN)  
Computer Science Department

Adrian Coleșa

April 29, 2020

---

6.2.1

### Purpose and Contents

#### The purpose of this chapter

- Present different synchronization mechanisms: semaphores and condition variables
- Present some classical synchronization patterns (problems): producer / consumer

---

6.2.2

#### Bibliography

- A. Tanenbaum, *Modern Operating Systems*, 2nd Edition, 2001, Chapter 2, Processes, p. 100 – 132
- A. Downey, *The Little Book of Semaphores*, 2nd Edition, 2008, p. 1 – 106

---

6.2.3

### Contents

<b>1 Semaphores</b>	<b>1</b>
<b>2 The “Producer/Consumer” Problem</b>	<b>3</b>
<b>3 Condition Variables</b>	<b>5</b>
<b>4 Conclusions</b>	<b>10</b>

---

6.2.4

## 1 Semaphores

### Characteristics

- **generalization of locks**
  - more threads could be given permission
  - to enter simultaneously in the critical regions protected by the semaphore
- consists of
  - **a value**
    - \* number of available permissions to pass the checkpoint
    - \* only positive values are allowed
  - **two primitives**

- \* **P (sem\_down)**: blocks the current process if the value is 0, otherwise decrements the value
- \* **V (sema\_up)**: increments the value and wakes up a sleeping process
- main difference from locks
  - any thread is allowed to increase a semaphore value (no of permissions), not only those that previously got one
  - $\Rightarrow$  could be used both as a (generalized) lock and an event counter

## Possible Implementation (uni-processor systems)

```
P() // sema_down()
{
    disable_interrupts();
    while (value == 0) {
        insert(crt_thread, w_queue);
        block_current_thread();
    }
    value = value - 1;
    enable_interrupts();
}
```

```
V() // sema_up()
{
    disable_interrupts();
    value = value + 1;
    if (! empty(waiting_queue)) {
        th = remove_first(w_queue);
        insert(ready_queue, th);
    }
    enable_interrupts();
}
```

- **interrupts are disabled**, to provide **code atomicity**
  - after critical code is executed, interrupts are enabled back
- **“sleep / wake-up technique”** is used instead of the *“busy waiting”*
  - **“sending to sleep” (blocking)** a thread
    - \* append it to a waiting queue and
    - \* take the CPU from it
  - **“waking up” (unblocking)** a thread
    - \* remove it from waiting queue
    - \* append it to the ready queue
    - \* eventually being given the CPU again

## Usage Example

```
#include <pthread.h>

// restrict the number of "persons" in a room
// i.e. number of threads in their critical region
sem_t sem;

void three_in_a_room();

main()
{
    // initialize the semaphore!
    // do that before having threads working with it!
    sem_init(&sem, 1, 3);

    // create the competing threads
    for (int i=0; i < 100; i++)
        pthread_create(&t[i], NULL, three_in_a_room, NULL);

    // wait for all threads' termination
    for (int i=0; i < 100; i++)
        pthread_join(t[i], NULL);

    // remove the semaphore
    sem_destroy(&sem);
}

void three_in_a_room()
{
    // ask for a permission, i.e. try decrementing the semaphore
    // the thread will be blocked, if no permission is available,
    // i.e. semaphore's value is 0
    sem_wait(&sem);

    // stay in room (critical region) for a while

    // release the obtained permission
    // i.e. increment the semaphore's value
    // this operation NEVER blocks the thread
    sem_post(&sem);
}
```

## Practice (1)

- Using semaphores, write the (pseudo)code for the functions in the code below such that to not allow more than 22 players enter simultaneously on the football field. A player is represented by a thread executing the *football\_player()* function.

```
1 void football_player()
2 {
3     enter_football_field();
4     play_football();
5     exit_the_football_field();
6 }
```

6.2.8

## Practice (2)

- Using semaphores, write the (pseudo)code that limits to 5 the number of threads executing simultaneously the critical section in the function below.

```
1 #define FREE 1
2 #define BUSY 0
3 #define MAX_TH 5
4
5 int available_unit[MAX_TH] = {FREE, FREE, FREE, FREE, FREE};
6 int unit[MAX_TH] = {0, 0, 0, 0, 0};
7
8 void limited_area()
9 {
10     int pos = -1;
11
12     for (pos = 0; pos < MAX_TH; pos++)
13         if (available_unit[pos] == FREE) {
14             available_unit[pos] = BUSY;
15             break;
16         }
17
18     unit[pos]++;
19     available_unit[pos] = FREE;
20 }
21
```

6.2.9

## 2 The “Producer/Consumer” Problem

### Problem Description

- two types of processes
  - producers*: produce messages
  - consumers*: consume messages
- use a bounded buffer
- waiting conditions
  - for producer: “buffer is full”
  - for consumer: “buffer is empty”
- wakeup conditions
  - for producer: “there is space in buffer”
  - for consumer: “there are messages in buffer”

6.2.10

### Basic Solution With Race Conditions

```
// Global variables
const N = 100; // number of slots in buffer
int count = 0; // number of messages in the buffer

void producer(int item)
{
    // wait for available space
    while (count == N) {} // full buffer

    insert_item(item);

    // wakeup a consumer
    count = count + 1;
}
```

```

int consumer(int *item)
{
    // wait for available message
    while (count == 0) {} // empty buffer

    *item = remove_item();

    // wakeup a producer
    count = count - 1;
}

```

- **race conditions**
  - uncontrolled concurrent access to variable count
  - same message could be consumed more times
  - messages could be overwritten
- **busy waiting**
  - suggests us that **synchronization is needed**
  - **should be avoided** in practice
  - **replace it** with a sleep / wakeup **synchronization mechanism**

---

6.2.11

## Producers/Consumers Problem's Implementation With Semaphores

```

// Global variables and synchronization mechanisms
const N = 100; // number of slots in buffer
Semaphore mutex(1); // provides mutual exclusion to buffer
Semaphore can_produce_msg(N); // controls producers' access to buffer
Semaphore can_consume_msg(0); // controls consumers' access to buffer

void producer(int item)
{
    // check if buffer is full
    can_produce_msg.P();

    // gets mutual exclusion to buffer
    mutex.P();

    insert_item(item);

    // releases the buffer
    mutex.V();

    // new permission for consumers
    can_consume_msg.V();
}

int consumer(int *item)
{
    // check if buffer is empty
    can_consume_msg.P();

    // gets mutual exclusion to buffer
    mutex.P();

    *item = remove_item();

    // releases the buffer
    mutex.V();

    // new permission for producers
    can_produce_msg.V();
}

```

---

6.2.12

## Remarks On Producers/Consumers Implementation With Semaphores

- there are two ways of using semaphores
  - as a (generalized) lock to provide limited access (mutual exclusion in the example with mutex semaphore)
  - as a condition checking and event counter (semaphores can\_consume\_msg and can\_produce\_msg)
- the order of getting the semaphores is important and it can lead to deadlock
  - see the solution with mutex semaphore get first

---

6.2.13

## Deadlock In The Producers/Consumers Problem's Implementation With Semaphores

```

// Global variables and synchronization mechanisms
const N = 100; // number of slots in buffer
Semaphore mutex(1); // provides mutual exclusion to buffer
Semaphore can_produce_msg(N); // controls producers' access to buffer
Semaphore can_consume_msg(0); // controls consumers' access to buffer

```

```

void producer(int item)
{
    // gets mutual exclusion to buffer
    mutex.P();

    // check if buffer is full
    can_produce_msg.P();

    insert_item(item);

    // new permission for consumers
    can_consume_msg.V();

    // releases the buffer
    mutex.V();
}

int consumer(int *item)
{
    // gets mutual exclusion to buffer
    mutex.P();

    // check if buffer is empty
    can_consume_msg.P();

    *item = insert_item();

    // new permission for producers
    can_produce_msg.V();

    // releases the buffer
    mutex.V();
}

```

---

6.2.14

## 3 Condition Variables

### Condition Variables: Context and Problem

- **used in a mutual exclusion context**
- **a thread** having a lock **must wait for a specific condition** to be fulfilled before going further
  - waited condition reflect the shared resource state
  - lock must be taken before testing / changing the condition
- **should not wait (sleep) keeping the lock**
  - could block the other threads that could change condition to true
- ⇒ waiting thread should **release the lock while waiting** (sleeping)
- ⇒ another thread can take the lock in the meantime and change condition
- that thread **wakes up the sleeping (waiting) thread**

---

6.2.15

### Condition Variables: Possible Deadlock Scenario

```

// Thread waiting
// for a condition
// to be fulfilled
thread_1()
{
    mutex.lock();

    while (cond == FALSE) {
        // e.g. (no_msg == 0)
        // for a consumer
        sleep(1);
    }

    mutex.unlock();
}

// Thread changing
// the waited condition
// to TRUE
thread_2()
{
    mutex.lock();

    cond = TRUE;
    // e.g. no_msg++
    // for a producer
}

```

```
mutex.unlock();
}
```

---

6.2.16

## Condition Variables: Functionality

- what is it?
  - **a specialized waiting mechanism**
  - transparently release the lock while blocking the lock holder
  - at wakeup transparently reacquire the lock before resuming the waiting thread
- consists of
  - a waiting queue
  - two (three) primitives
    - \* **wait**: the calling process is inserted in the waiting queue and suspended
    - \* **signal**: one sleeping process is removed from the waiting queue and woken up
    - \* **broadcast**: all sleeping processes are woken up

---

6.2.17

## Condition Variables. Basic Pattern Usage

```
// Thread waiting for
// an event (condition)
// to occur
thread_1()
{
    mutex.lock();

    while (! cond)
        c.wait(mutex);

    mutex.unlock();
}

// Thread generating
// the event
//
thread_2()
{
    mutex.lock();

    cond = TRUE;
    c.signal();

    mutex.unlock();
}
```

---

6.2.18

## Condition Variables: Implementation

```
wait(Lock mutex)
{
    disable_interrupts();

    mutex.unlock();

    insert(crt_th, w_queue);
    block_current_thread();

    mutex.lock();

    enable_interrupts();
}
```

```

signal()
{
    disable_interrupts();

    if (! empty(waiting_queue)) {
        th = remove_first(w_queue);
        insert(ready_queue, th);
    }

    enable_interrupts();
}

```

---

6.2.19

## Condition Variables. General Pattern Usage (Var 1)

```

// Synchronization mechanisms: one condition variable
Lock mutex;
Condition c;

// Thread 1's Function
thread_1()
{
    // critical region entrance
    mutex.lock();
    while (! cond_1)
        c.wait(mutex);
    mutex.unlock();

    // thread T1 in its critical region
    ....

    // critical region exit
    mutex.lock();
    cond_2 = TRUE;
    c.broadcast();
    mutex.unlock();
}

// Thread 2's Function
thread_2()
{
    // critical region entrance
    mutex.lock();
    while (! cond_2)
        c.wait(mutex);
    mutex.unlock();

    // thread T2 in its critical region
    ...

    // critical region exit
    mutex.lock();
    cond_1 = TRUE;
    c.broadcast();
    mutex.unlock();
}

```

---

6.2.20

## Condition Variables. General Pattern Usage (Var 2)

```

// Synchronization mechanisms: two condition variables
Lock mutex;
Condition c1, c2;

// Thread 1's Function
thread_1()
{
    // critical region entrance
    mutex.lock();
    while (! cond_1)
        c1.wait(mutex);
    mutex.unlock();

    // thread T1 in its critical region
    ....

    // critical region exit
    mutex.lock();
    cond_2 = TRUE;
    c2.signal(); // c2.broadcast();
    mutex.unlock();
}

// Thread 2's Function
thread_2()
{
    // critical region entrance
    mutex.lock();
    while (! cond_2)
        c2.wait(mutex);
    mutex.unlock();

    // thread T2 in its critical region
    ...

    // critical region exit
    mutex.lock();
    cond_1 = TRUE;
    c1.signal(); // c1.broadcast();
    mutex.unlock();
}

```

---

6.2.21

## Condition Variables. Multiple Conditions (Var 1)

```
// Synchronization mechanisms: multiple conditions & one condition variable
Lock mutex;
Condition c;

// Thread 1's Function
thread_1()
{
    // critical region entrance
    mutex.lock();
    while (! cond_1 || ! cond_2)
        c.wait(mutex);
    mutex.unlock();

    // T1's critical region
    ....

    // critical region exit
    mutex.lock();
    .....
    mutex.unlock();
}

// Thread 2's Function
thread_2()
{
    ....
    // critical region exit
    mutex.lock();
    switch (x) {
        case 1:
            cond_1 = TRUE;
            c.broadcast(); // c.signal() ?
            break;
        case 2:
            cond_2 = TRUE;
            c.broadcast(); // c.signal() ?
            break;
    }
    mutex.unlock();
}
```

6.2.22

## Condition Variables. Multiple Conditions (Var 2 - BAD)

```
// Synchronization mechanisms: multiple conditions & multiple condition variables & one lock
Lock mutex;
Condition c1, c2;

// Thread 1's Function
thread_1()
{
    // critical region entrance
    mutex.lock();
    while (! cond_1)
        c1.wait(mutex);

    while (! cond_2)
        c2.wait(mutex);
    mutex.unlock();

    // T1's critical region
    ....

    // critical region exit
    mutex.lock();
    .....
    mutex.unlock();
}

// Thread 2's Function
thread_2()
{
    ....
    // critical region exit
    mutex.lock();
    switch (x) {
        case 1:
            cond_1 = TRUE;
            c1.signal(); // c1.broadcast();
            break;
        case 2:
            cond_2 = TRUE;
            c2.signal(); // c2.broadcast();
            break;
    }
    mutex.unlock();
}
```

6.2.23

## Condition Variables. Multiple Conditions (Var 2 - OK)

```
// Synchronization mechanisms: multiple conditions & multiple condition variables & multiple locks
Lock mutex1, mutex2;
Condition c1, c2;

// Thread 1's Function
thread_1()
{
    // critical region entrance
    mutex1.lock();
    while (! cond_1)
        c1.wait(mutex1);

    mutex2.lock();
    while (! cond_2)
```



```

        c2.wait(mutex2);
mutex2.unlock();
mutex1.unlock();

// T1's critical region
....

// critical region exit
.....
}

// Thread 2's Function
thread_2()
{
    ....
    // critical region exit

    switch (x) {
    case 1:
        mutex1.lock();
        cond_1 = TRUE;
        c1.signal(); // c1.broadcast();
        mutex1.unlock();
        break;
    case 2:
        mutex2.lock();
        cond_2 = TRUE;
        c2.signal(); // c2.broadcast();
        mutex2.unlock();
        break;
    }
}
}

```

---

6.2.24

## Condition Variables: Comparison With Semaphores

- *wait* always suspends the calling process
  - while *P* does this only in case the semaphore's value is zero
- a signal can be lost (i.e. not seen)
  - while a *V* is not, because it increments the semaphore's value
- safe to use inside a mutual exclusion area
  - while semaphore could lead to deadlock

---

6.2.25

## Rules to Remember About Using Condition Variables

1. **always use them** (wait and signal) **inside a mutual exclusion area** (protected by a lock)
2. the lock does (normally) not protect the shared resource
  - is would simply provide mutual exclusion, but maybe not needed
  - does not support deadlock-free waiting
3. the lock protects the entrance in and exit from critical region
  - where some condition checks or changes are done
  - in a mutual exclusion manner
4. it **could be needed to recheck the condition** after returning from wait
  - use *while* instead of *if*
5. **do not use wait without checking a condition**
  - it is not sure the thread must wait (we do not control thread scheduling)
  - signal has no history, so can be lost

---

6.2.26

## Rules to Remember About Using Condition Variables (review)

1. **eyes** for seeing
2. **ears** for hearing
3. **mind** for thinking (understanding, learning, remembering)
4. see <https://www.youtube.com/watch?v=RZxAc3Grkck>
5. review previous slide and make use of the “elements” mentioned above

---

6.2.27

## Producers/Consumers Problem's Implementation With Locks and Condition Variables

```
// Global variables and synchronization mechanisms
const N = 100;           // number of slots in buffer
int count = 0;           // number of messages in buffer
Lock mutex;              // provides mutual exclusion to buffer
Condition producers;     // controls producers' access to buffer
Condition consumers;     // controls consumers' access to buffer

void producer(int item)
{
    // gets mutual exclusion to buffer
    mutex.lock();

    // check if buffer is full
    while (count == N)
        producers.wait(mutex);

    insert_item(item);
    count = count + 1;

    // wakes up a consumer
    consumers.signal();

    // releases the buffer
    mutex.unlock();
}

int consumer(int *item)
{
    // gets mutual exclusion to buffer
    mutex.lock();

    // check if buffer is empty
    while (count == 0)
        consumers.wait(mutex);

    *item = remove_item();
    count = count - 1;

    // wakes up a producer
    producers.signal();

    // releases the buffer
    mutex.unlock();
}
```

6.2.28

### Practice (3)

- Using locks and condition variables, write the (pseudo)code for the functions in the code below such that to not allow more than 22 players enter simultaneously on the football field. A player is represented by a thread executing the *football\_player()* function.

```
void football_player()
{
    enter_football_field();
    play_football();
    exit_the_footbal_field();
}
```

6.2.29

## 4 Conclusions

### What we talked about

- semaphores
  - a generalization of locks  $\Rightarrow$  could allow more threads pass the barrier
  - important to initialize the semaphore
  - “P()” and “V()” primitives
  - could also be used as an event counter, usually initialized with 0
- condition variables
  - wait: provide a specialized way to wait for a condition to be fulfilled
  - signal: provide a way to signal that a condition is fulfilled
- producer-consumer synchronization pattern

6.2.30

### Lessons Learned

- locks could be too restrictive
- semaphores are more flexible, though must be used with care to avoid deadlock
- condition variables must be used with a lock, i.e. inside a mutual exclusion area

6.2.31