

Chapter 10

Review and Conclusions

What have you learned? What else could you still learn?

Print Version of Lectures Notes of *Operating Systems*
Technical University of Cluj-Napoca (UTCN)
Computer Science Department

Adrian Coleșa
June 5, 2019

10.1

Purpose and Contents

The purpose of this chapter

- 1. review subjects presented during the OS course (this semester)
- 2. draw some conclusions
- 3. presents future OS-related courses

10.2

Contents

1	OS Subject Review	1
2	Conclusions	5
2.1	Past	5
2.2	Future	6

10.3

1 OS Subject Review

OS Definition, Role, Architecture

- system software placed between hardware and user software (applications)
- roles
 - virtual machine provider (hide hardware and provide abstractions)
 - resource manager
- every resource access / need must be required from OS
 - by calling SO services, i.e. system calls
- OS protects itself from user applications based on different CPU execution modes
 - privileged: kernel mode
 - non-privileged: user mode
- architectures
 - monolith
 - micro-kernel

10.4

Practice: OS Definition Related Questions

- When is the OS code executed on an uni-processor system? Give examples of at least two situations.
- How is the system call mechanism implemented?
- Which software is run in kernel and user mode respectively in the following two cases?
 - monolithic OS
 - micro-kernel OS

10.5

OS Shell

- provides the user interface to the OS
- usually a simple application
- takes user requests and translated them in system calls
- two types
 - graphical
 - text, named command interpreter
- command interpreter functionality
 - get user's command line
 - split it in tokens, i.e. command line and arguments
 - create a new process to execute the specified command
 - waits for created child process' termination
- command line
 - a string of characters separated by spaces
 - first item: command name, actually an executable path
 - * searched in directories from PATH
 - * security: trust user-established environment (e.g. PATH)
 - other items: command line arguments
 - special characters, like STDIN/OUT redirection, pipe etc.

10.6

Practice: Shell Related Questions

Which is the effect of the following commands?

- `ls > file`
- `read n < file`
- `ls -R / 1>good 2>err`
- `cat dict.txt | sort`

10.7

File System (FS)

- file concept
 - basic unit of data allocation
 - unstructured stream of bytes
 - file contents managed by user applications, not by OS
 - components: data and meta-data
 - security: too much permissions vulnerability
- directory concept
 - used for organizing the file system space
 - impose file system hierarchy
 - paths: absolute and relative
 - security: path traversal vulnerability

- FS system calls
 - file: open, read, write, lseek, close
 - directory: opendir, readdir, stat, unlink, link
- allocation aspects
 - contiguous allocation \Rightarrow external fragmentation
 - any-free-block allocation \Rightarrow internal fragmentation
 - i-nodes, directory entries, links, files with holes

10.8

Practice: FS Related Questions

- How is usually the file provided like to the user applications?
- Which of the following extensions usually correspond to text and binary files respectively: html, pdf, c, zip?
- What is an i-node in Linux?
- What does 0640 means in terms of permission rights in Linux?
- Which is the most probable file descriptor returned (and displayed) by the following Linux program? Explain your answer.

```
main()
{
    int fd = open ("/etc/passwd", O_RDONLY);
    printf("fd = %d\n", fd);
}
```

- Write in one line the C code to read a text line from a file, whose file descriptor is given.
- Write the C code to read an integer from offset 16 from a file, whose file descriptor is given.

10.9

Process and Thread Management

- process
 - models execution: abstractizes the machine (CPU, memory)
 - describe execution and resources needed for that execution
 - isolates (separates) resources / executions
 - states: running, ready, blocked, terminated
- thread
 - models execution in a process
 - more threads = more concurrent executions in the same process
 - threads of a process share all resources of that process
 - threads useful and effective when
 - * logical parallelism exist in the application
 - * enough hardware resources available
- scheduling
 - decides who runs and for how long
 - preemptive vs. non-preemptive
 - * preemptiveness based on timer interrupt
 - first-come first-served (FCFS), shortest job first (SJB), round-robin (RR), priority-based

10.10

Practice: Process Related Questions

- How many times is each message in the code below displayed on the screen? Explain your answer.

```
int fd[2];
pipe(fd);
printf("Step 1\n");
fork();
fork();
fork();
printf("Step 2\n");
```

- How many readers and writers, respectively, will exist in the system for the created pipe after the execution of the given code, supposing no process is terminated at that moment?
- Which is the optimum number of threads that should be created by an application to get the best performance (i.e. execution time) when running on a uni-process system and copying a file from one disk to another disk, by encrypting the file contents during the copy operation?

10.11

Synchronization Mechanisms

- problem: race conditions of concurrent threads sharing resources
- synchronization means imposing access rules
 - leads to waiting (blocking) \Rightarrow reduce parallelism
- synchronization needs OS and hardware support to get atomicity
- synchronization mechanisms
 - lock \Rightarrow mutual exclusion
 - semaphore
 - * generalized lock
 - * event counter
 - condition variable
 - * specialized waiting mechanism in mutual exclusion area
- classical patterns
 - producers-consumers, readers-writers, rendez-vous (barrier)

10.12

Practice: Synchronization Related Questions

- Use semaphores to allow just 10 threads run simultaneously a given function's body.
- Rendezvous: synchronize threads executing functions `boy()` and `girl()` respectively, such that to allow them returning from that functions only in pairs of a "boy" and a "girl".
- Synchronize two concurrent threads executing the two functions below respectively, such that to make them display on the screen the message "*Life is wonderful, isn't it?*"

```
thread_1()                thread_2()
{                          {
    printf("Life ");        printf("is ");
    printf("wonderful, ");  printf("isn't it?");
}                          }
```

10.13

Memory Management

- memory addresses: physical vs. virtual
 - address space
 - virtual address space structure: code, data, heap, stack
- memory binding / translation
 - compile-time (very limited), load-time, run-time (most flexible)
 - need translation tables
- contiguous allocation
 - simple, efficient
 - leads to external fragmentation

- base and limit registers
- segmentation
 - one contiguous area for each segment (area)
 - segment table
- paging
 - allocates memory in fixed-size chunks \Rightarrow internal fragmentation
 - virtual pages and physical frames
 - page tables and page table entries
 - page sharing, memory mapped files

10.14

Practice: Memory Mng Related Questions

- Which is the size in bytes of a process' virtual address space on a system using 64 bits for a memory address?
- How many page table entries must be used by such a system to map a process VAS, supposing the page size to be 4MB?
- Illustrate on such a system the part of a process' VAS and page table (entries) used for mapping the memory required by the following code:

```
unsigned int *p = malloc(4*4*1024*1024);
printf("p=%u\n", p); // displays p = 1000*4*1024*1024
```

- Which page does the following instruction refer to? Could it be executed successfully or not?

```
p[4*1024*1024] = 10;
```

10.15

Security Aspects

- untrusted application environment
 - untrusted PATH variable
- file system
 - too much permissions
 - path traversal
- memory-related: bad / wrong memory accesses
 - NULL-pointer usage
 - use-after free
 - buffer overflow

10.16

2 Conclusions

2.1 Past

What did we talk about?

- OS's place and role (and definition), relative to hardware and other software
- OS structure
- Shell, i.e. command interpreter
- File System
- Process Management
- Memory Management
- Security Aspects

10.17

What have we learned?

- basic concepts like
 - process, thread — execution, resources, isolation, scheduling
 - file: storage, sequence of bytes (no format), meta-data (i-node), links, fragmentation, open files
 - directory: organization, file tree, collection of directory entries
 - synchronization: locks (mutex), semaphores, condition variables, producers/consumers, readers/writers, barrier
 - memory: address space, virtual/physical memory addresses, ELF, paging
 - IPC mechanisms: pipes, shared memory
 - fragmentation: external (contiguous allocation, best-fit, worst-fit), internal
 - basic security issues: buffer overflow, path traversal
- system calls to access various OS (Linux) services
- write C programs to have access to Linux system calls

10.18

What can we do?

- explain OS functionality and concepts
- understand better applications functionality, their relationship with the OS and some of their crash reasons
- write (if needed) C programs to access low-level OS services
- choose the appropriate OS for a particular purpose
- tune better an OS for particular applications/context/purposes

10.19

2.2 Future

Will I need OS knowledge?

- all the time: understand, explain, evaluate, configure

10.20

Will I use OS system calls?

- sometimes, especially when you need a particular (efficient, lower-level) functionality
- all the time, if you work at low-level (in C, asm)

10.21

Operating System Design

- this is (our) next step: the advanced course (UTCN, CS Dept.) regarding OS
- what about: internals of an OS, design/implementation alternatives of
 - scheduling algorithms and synchronization mechanisms
 - user processes and threads
 - system calls for open files, processes and threads
 - memory management, virtual memory, page replacement algorithms
 - file system
- practical aspects: design, implement (in C) and test an OS (*HAL9000* — UTCN; *Pintos* — Stanford, USA); learn how to debug an OS on a remote virtual machine
- practical aspects: work in team (3 members)

10.22

Operating System Design (cont.)

- usefulness
 - an OS is a complex management software; all the management algorithms and techniques work for many other complex management software even at higher levels
 - understanding low-level mechanisms makes you understand better higher-level ones' behavior
- it is an 4th year *optional subject*
- *students' myth*: (highly) difficult
- *reality (trust me)*: just interesting and challenging, not more difficult than other subject-s/projects

10.23

Operating System Design (cont.)

- I would like you to think like this:
 - “I’d choose this subject just because I am really interested in, it sounds great and challenging, I feel I need it in my future career, I want to understand how low-level mechanisms and aspects work, it fits my aptitudes and interests, I have heard about learning useful things etc.”
- I would not like to hear about you saying:
 - “I DID NOT choose that subject just because I have a job and do not have enough time for it, I’ve heard it takes long(er) to solve the assignments, I’ve heard it is (more) difficult etc.”

10.24

Security Master Program (SISC)

- many OS-related courses
 - build low-level OS layers on the 64-bit architecture
 - kernel driver development
 - build virtualization security-oriented OS (hypervisor)
- many security-related aspects
 - secure coding
 - Web security
 - mobile systems (Android) security
 - big-data and security
 - penetration testing
 - risk management
 - cryptography

10.25

Finish

That’s all folks! So long, folks!

10.26